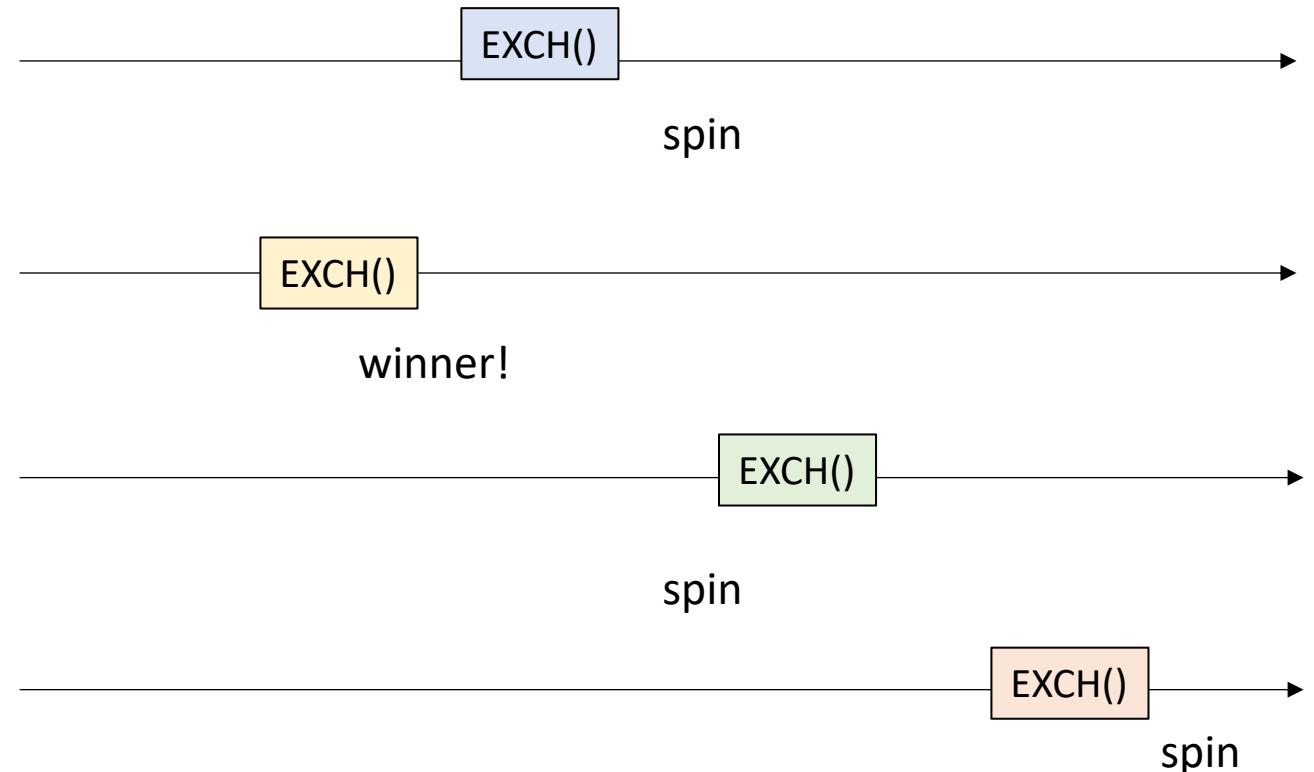


# CSE113: Parallel Programming

April 20, 2021

- **Topic:** Practical Mutual Exclusion

- Atomic RMW locks
- Optimizing locks



# Announcements

- Homework due on Thursday at Midnight
  - Gan set up a submission link on Canvas
  - Questions:
    - Visit a few from Canvas
  - Should Gan do office hours on Thursday instead of Friday?
    - 10 - 11 AM on Thursday
- My office hours are on Wednesday
  - Unfortunately do not have many on Thursday

# Announcements

- Next homework:
  - Assigned by midnight Thursday
- Sign up for Piazza:
  - Lots of good discussions
  - Only 50 of you have
- I trust things are going well on Discord...

# Announcements

- A little bit out of sync with the book
  - We'll do first half of Chapter 7 today
  - Next lecture will be last of Chapter 7 and some of Chapter 8
  - Still on track to start Module 3 (concurrent data structures) in 1 week

# Quiz

- Canvas Quiz

# Quiz

- Canvas Quiz
- Go over answers

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

recall this snippet of code.  
What was the issue?

Tyler's coffee addiction:

```
m.lock();
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    return;
}
m.unlock();
return;
```

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

recall this snippet of code.  
What was the issue?

We need to unlock  
at two places

Tyler's coffee addiction:

```
m.lock();
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    m.unlock();
    return;
}
m.unlock();
return;
```

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

recall this snippet of code.  
What was the issue?

We need to unlock  
at two places

Tyler's coffee addiction:

```
m.lock();
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    m.unlock();
    return;
}
m.unlock();
return;
```

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

Tyler's coffee addiction:

```
lock_guard<mutex> lck(m);
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    return;
}
return;
```

recall this snippet of code.  
What was the issue?

Or we use lock\_guard

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

recall this snippet of code.  
What was the issue?

Or we use lock\_guard

A template on a mutex type (i.e. implements lock and unlock).

Tyler's coffee addiction:

```
lock_guard<mutex> lck(m);
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    return;
}
return;
```

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

recall this snippet of code.  
What was the issue?

Or we use lock\_guard

A template on a mutex type (i.e. implements lock and unlock).

Tyler's coffee addiction:

```
lock_guard<mutex> lck(m);
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    return;
}
return;
```

pass the mutex into the constructor.  
No other methods!

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

recall this snippet of code.  
What was the issue?

Or we use lock\_guard

Constructor locks the mutex,  
destructor unlocks it.

A template on a mutex type (i.e. implements lock and unlock).

Tyler's coffee addiction:

```
lock_guard<mutex> lck(m);
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    return;
}
return;
```

pass the mutex into the constructor.  
No other methods!

# Before we start: C++ lock\_guard

- C++ lock\_guard - Pretty cool!
- Uses C++ constructor and destructor

recall this snippet of code.  
What was the issue?

Or we use lock\_guard

Constructor locks the mutex,  
destructor unlocks it.

Writing to files can throw exceptions, if  
we don't handle the exception, then the system  
could deadlock.

Tyler's coffee addiction:

```
lock_guard<mutex> lck(m);
tylers_account -= 1;
if (tylers_account < -100) {
    printf("warning!\n");
    log_warning_to_file();
    return;
}
return;
```

pass the mutex into the constructor.  
No other methods!

# Lecture Schedule

- Atomic RMW mutexes
  - Exchange
  - CAS
  - Ticket
- Optimizations
  - Relaxed peeking
  - Backoff

# Lecture Schedule

- **Atomic RMW mutexes**

- Exchange
- CAS
- Ticket

- Optimizations

- Relaxed peeking
- Backoff

# From previous lecture: Peterson's mutex

- Peterson's algorithm: 2 threaded mutex implementation
  - 2 flag values
  - 1 victim
- We used primitives: atomic loads and stores
- Proof in book:
  - If you only use loads and stores, synchronizing  $N$  threads requires  $O(N)$  memory

# From previous lecture: Peterson's mutex

- Peterson's algorithm: 2 threaded mutex implementation
  - 2 flag values
  - 1 victim
- We used primitives: atomic loads and stores

```
class Mutex {  
public:  
    Mutex() {  
        victim = -1;  
        flag[0] = flag[1] = 0;  
    }  
  
    void lock();  
    void unlock();  
  
private:  
    atomic_int victim;  
    atomic_bool flag[2];  
};
```

# From previous lecture: Peterson's mutex

- Peterson's algorithm: 2 threaded mutex implementation
  - 2 flag values
  - 1 victim
- We used primitives: atomic loads and stores

```
void unlock() {
    int i = thread_id;
    flag[i].store(0);
}
```

```
void lock() {
    int i = thread_id;
    flag[i].store(1);
    victim.store(thread_id);
    while (victim.load() == thread_id
           && flag[j].load() == 1);
}
```

```
class Mutex {
public:
    Mutex() {
        victim = -1;
        flag[0] = flag[1] = 0;
    }

    void lock();
    void unlock();

private:
    atomic_int victim;
    atomic_bool flag[2];
};
```

# From previous lecture: Peterson's mutex

- Peterson's algorithm
  - 2 threads
  - 2 flag values
- Generalizations:
  - Filter lock - N threaded Peterson's algorithm (uses  $2*N$  memory)
  - Bakery lock - N threaded fair mutex (uses  $2*N$  memory)
  - Implementations in the Book! Chapter 3

# From previous lecture: Peterson's mutex

- Peterson's algorithm
  - 2 threads
  - 2 flag values
- Implementing Peterson's was difficult because of loads/stores interleaving!

# From previous lecture: Peterson's mutex

- Peterson's algorithm
  - 2 threads
  - 2 flag values
- Implementing Peterson's was difficult because of loads/stores interleaving!
- But what if there was another way...

# Buggy Mutex implementation

```
#include <atomic>
using namespace std;

class Mutex {
public:
    Mutex() {
        flag = 0;
    }
    void lock();
    void unlock();
private:
    atomic_bool flag;
};

mutex is initialized to "free"
atomic_bool for our memory location
```

# Buggy Mutex implementation

```
void lock() {  
    while (flag.load() == 1);  
    flag.store(1);  
}
```

While the mutex is not available (i.e. another thread has it)

Once the mutex is available, we will claim it

# Buggy Mutex implementation

```
void unlock() {  
    flag.store(0);  
}
```

To release the mutex, we just set it back to 0 (available)

## Buggy Mutex implementation: Analysis

```
void lock() {  
    while (flag.load() == 1);  
    flag.store(1);  
}
```

```
void unlock() {  
    flag.store(0);  
}
```

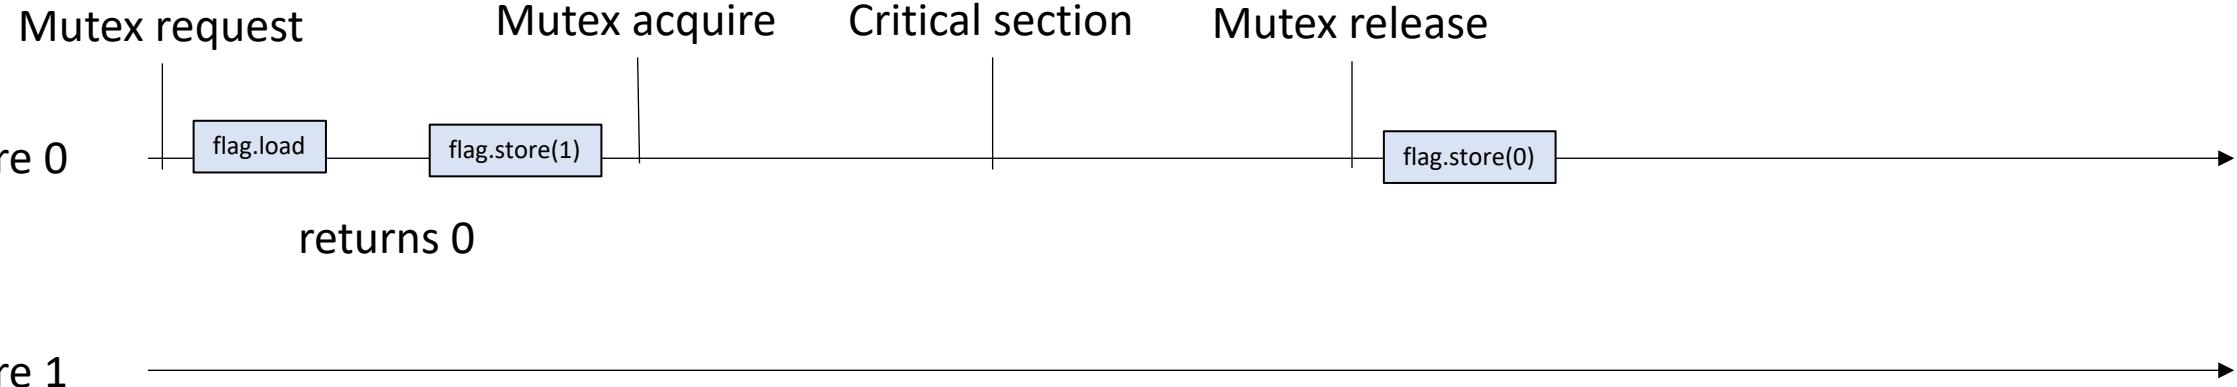
Thread 0:

```
m.lock();  
m.unlock();
```

Thread 1:

```
m.lock();  
m.unlock();
```

*Lets try another interleaving*



## Buggy Mutex implementation: Analysis

```
void lock() {  
    while (flag.load() == 1);  
    flag.store(1);  
}
```

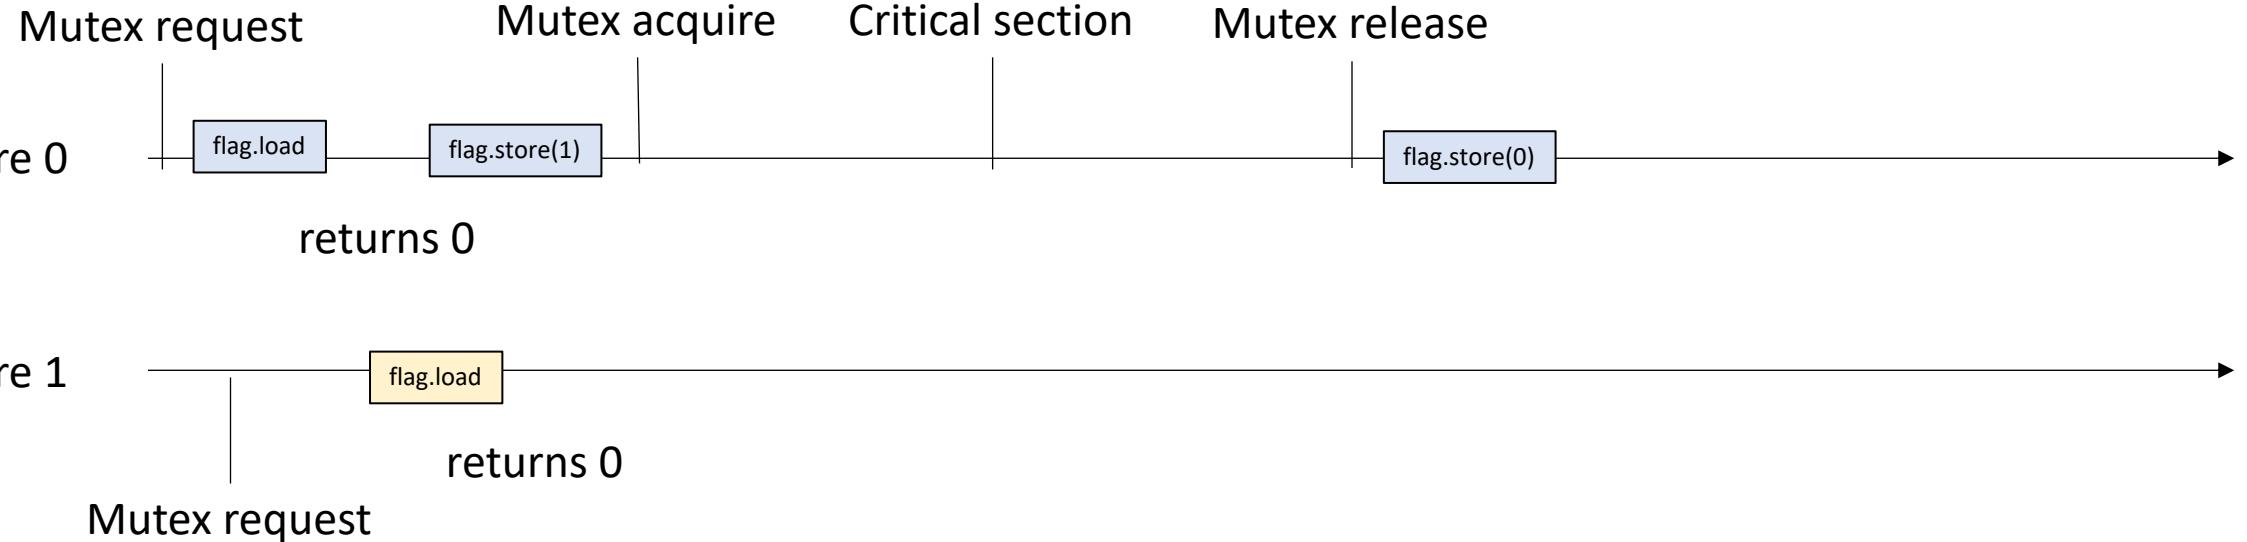
```
void unlock() {  
    flag.store(0);  
}
```

Thread 0:

```
m.lock();  
m.unlock();
```

Thread 1:

```
m.lock();  
m.unlock();
```



## Buggy Mutex implementation: Analysis

```
void lock() {  
    while (flag.load() == 1);  
    flag.store(1);  
}
```

```
void unlock() {  
    flag.store(0);  
}
```

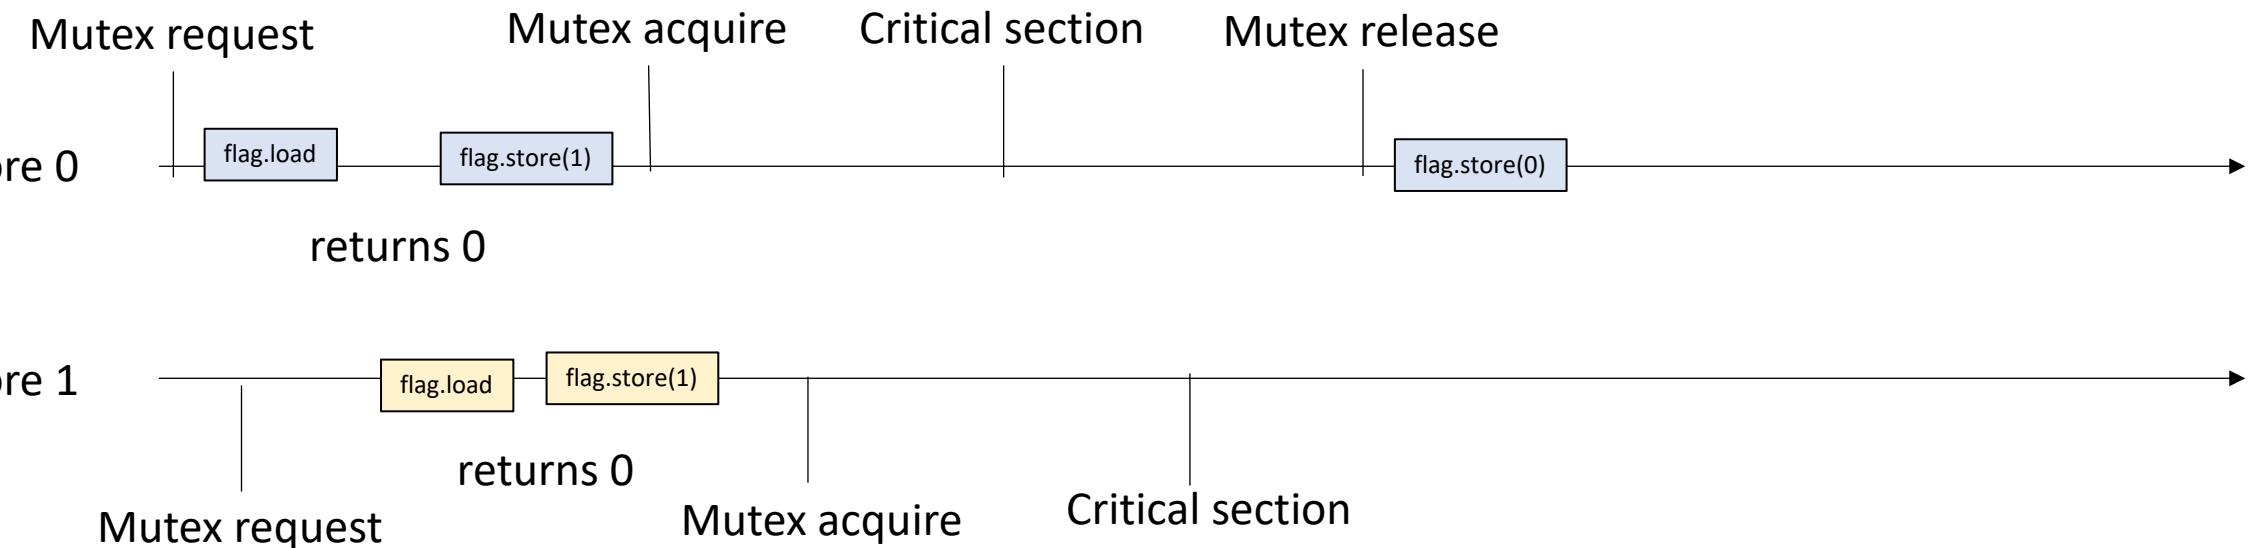
Thread 0:

```
m.lock();  
m.unlock();
```

Thread 1:

```
m.lock();  
m.unlock();
```

*Critical sections overlap! This mutex implementation is not correct!*



# What went wrong?

- The load and stores from two threads interleaved
  - What if there was a way to prevent this?

# What went wrong?

- The load and stores from two threads interleaved
  - What if there was a way to prevent this?
- Atomic RMWs
  - operate on atomic types (we already have atomic types)
  - recall the non-locking bank accounts:  
`atomic_fetch_add(atomic *a, value v);`

# What is a RMW

A read-modify-write consists of:

- *read*
- *modify*
- *write*

done atomically, i.e. they cannot interleave.

Typically returns the value (in some way) from the read.

# atomic\_fetch\_add

Recall the lock free account

Atomic Read-modify-write (RMWs): primitive instructions that implement a read event, modify event, and write event indivisibly, i.e. it cannot be interleaved.

```
atomic_fetch_add(atomic_int * addr, int value) {
    int tmp = *addr; // read
    tmp += value;    // modify
    *addr = tmp;     // write
}
```

# atomic\_fetch\_add

Recall the lock free account

Atomic Read-modify-write (RMWs): primitive instructions that implement a read event, modify event, and write event indivisibly, i.e. it cannot be interleaved.

```
int atomic_fetch_add(atomic_int * addr, int value) {
    int tmp = *addr; // read
    tmp += value;   // modify
    *addr = tmp;     // write
    return tmp;      // return previous value in the memory location
}
```

# lock-free accounts

Tyler's coffee addiction:

```
atomic_fetch_add(&tylers_account, -1);
```

Tyler's employer

```
atomic_fetch_add(&tylers_account, 1);
```



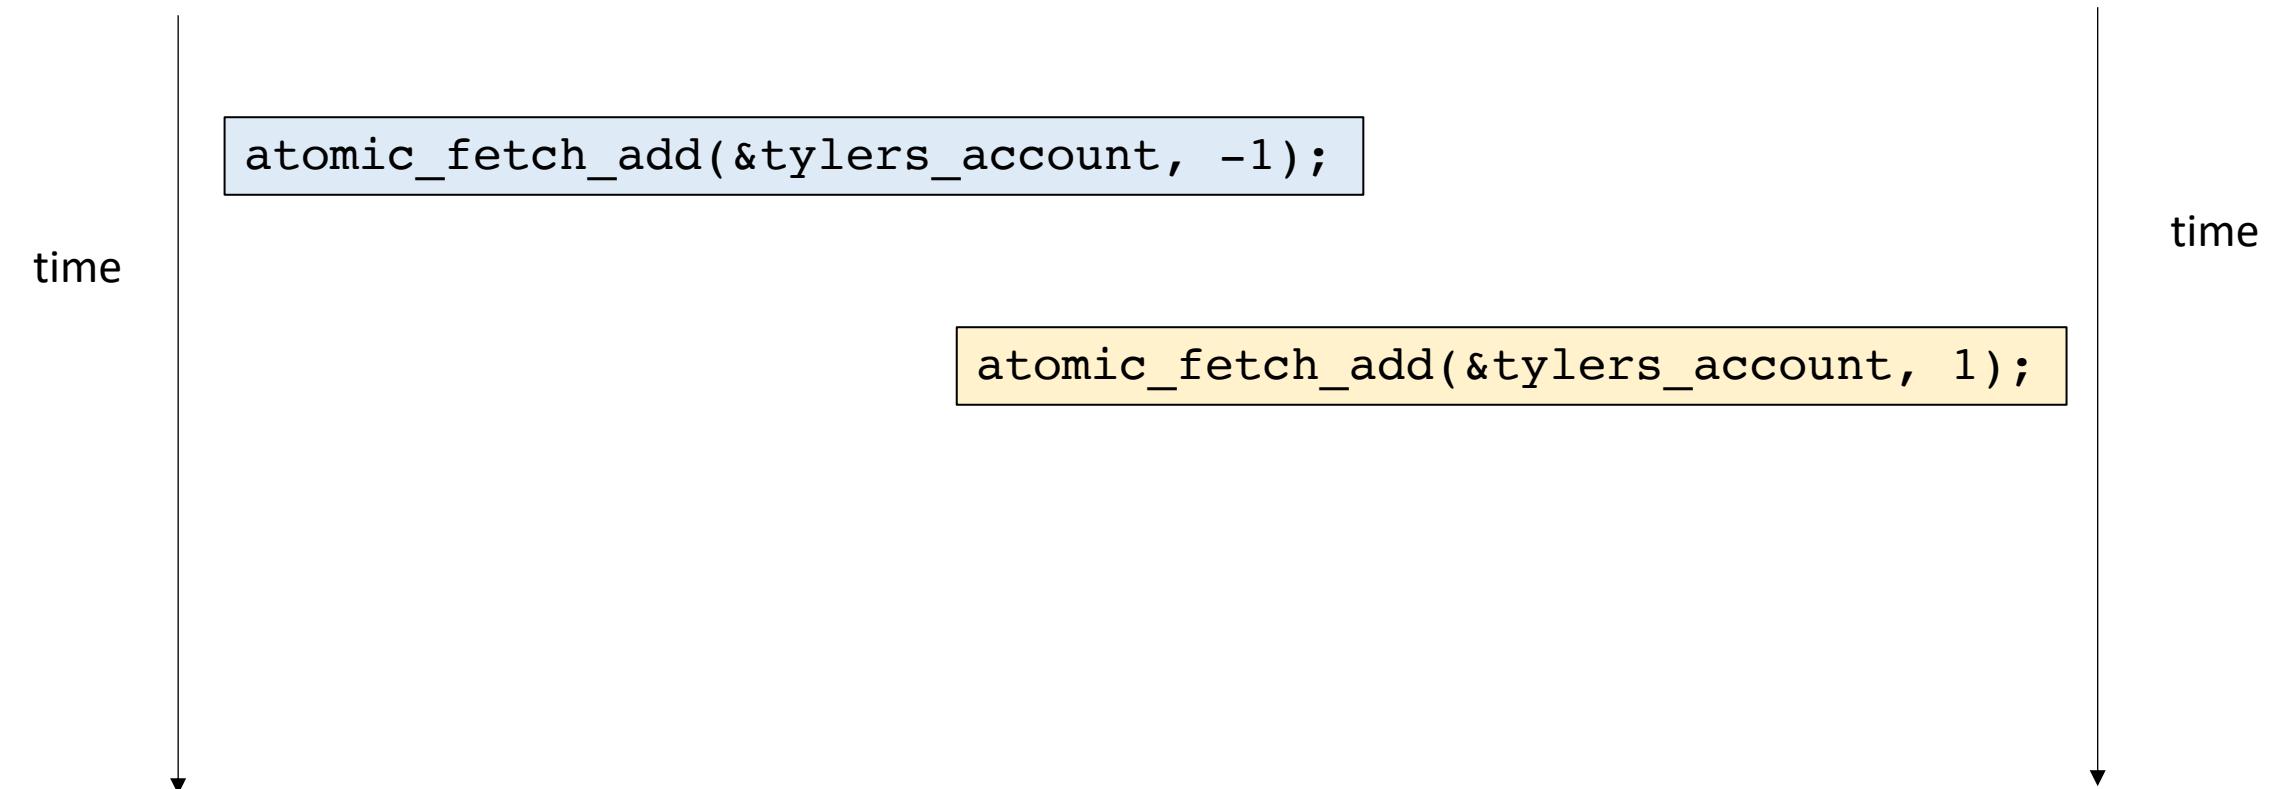
# lock-free accounts

Tyler's coffee addiction:

```
atomic_fetch_add(&tylers_account, -1);
```

Tyler's employer

```
atomic_fetch_add(&tylers_account, 1);
```



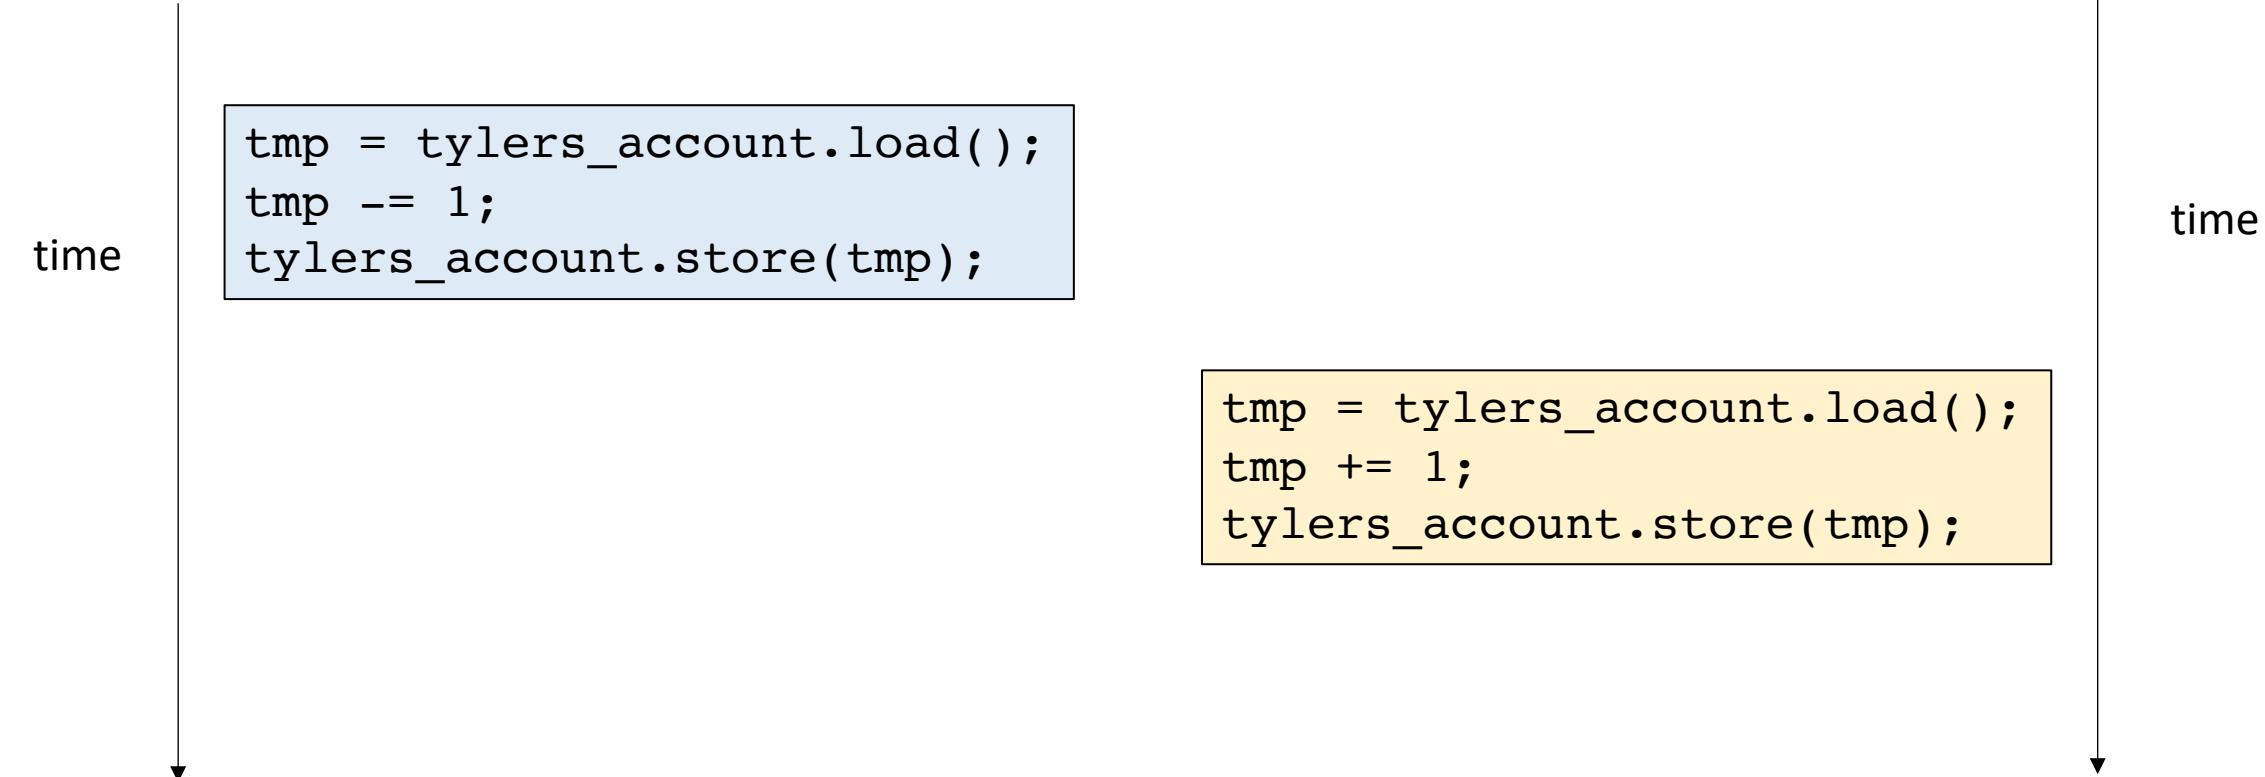
## lock-free accounts

Tyler's coffee addiction:

```
atomic_fetch_add(&tylers_account, -1);
```

Tyler's employer

```
atomic_fetch_add(&tylers_account, 1);
```



## lock-free accounts

Tyler's coffee addiction:

```
atomic_fetch_add(&tylers_account, -1);
```

Tyler's employer

```
atomic_fetch_add(&tylers_account, 1);
```

```
tmp = tylers_account.load();
tmp -= 1;
tylers_account.store(tmp);
```

cannot interleave!

time

```
tmp = tylers_account.load();
tmp += 1;
tylers_account.store(tmp);
```

↓

# lock-free accounts

Tyler's coffee addiction:

```
atomic_fetch_add(&tylers_account, -1);
```

Tyler's employer

```
atomic_fetch_add(&tylers_account, 1);
```

cannot interleave!

```
tmp = tylers_account.load();
tmp += 1;
tylers_account.store(tmp);
```

```
tmp = tylers_account.load();
tmp -= 1;
tylers_account.store(tmp);
```

either way, account breaks even at the end!

time

time

# RMW-based locks

- A few simple RMWs enable lots of interesting mutex implementations
- When we have simpler implementations, we can focus on performance

# Lecture Schedule

- **Atomic RMW mutexes**

- Exchange
  - CAS
  - Ticket

- Optimizations

- Relaxed peeking
- Backoff

# First example: Exchange Lock

- Simplest atomic RMW will allow us to implement an:
- N-threaded mutex with 1 bit!

# First example: Exchange Lock

```
value atomic_exchange(atomic *a, value v);
```

Loads the value at a and stores the value in v at a. Returns the value that was loaded.

# First example: Exchange Lock

```
value atomic_exchange(atomic *a, value v);
```

Loads the value at a and stores the value in v at a. Returns the value that was loaded.

```
value atomic_exchange(atomic *a, value v) {  
    value tmp = a.load();  
    a.store(v);  
    return tmp;  
}
```

# First example: Exchange Lock

```
value atomic_exchange(atomic *a, value v);
```

Loads the value at a and stores the value in v at a. Returns the value that was loaded.

```
value atomic_exchange(atomic *a, value v) {  
    value tmp = a.load();  
    a.store(v);  
    return tmp;  
}
```

*no "modify" step!*

# First example: Exchange Lock

```
#include <atomic>
using namespace std;

class Mutex {
public:
    Mutex() {
        flag = false;
    }

    void lock();
    void unlock();

private:
    atomic_bool flag;
};
```

Lets make a mutex with just one atomic bool!

# First example: Exchange Lock

```
#include <atomic>
using namespace std;

class Mutex {
public:
    Mutex() {
        flag = false;
    }

    void lock();
    void unlock();

private:
    atomic_bool flag;
};
```

Lets make a mutex with just one atomic bool!

initialized to false

one atomic flag

# First example: Exchange Lock

```
#include <atomic>
using namespace std;

class Mutex {
public:
    Mutex() {
        flag = false;
    }

    void lock();
    void unlock();

private:
    atomic_bool flag;
};
```

Lets make a mutex with just one atomic bool!

initialized to `false`

**main idea:**

The flag is false when the mutex is free.

one atomic flag

The flag is true when some thread has the mutex.

# First example: Exchange Lock

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

# First example: Exchange Lock

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

So what's going on?

# First example: Exchange Lock

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Two cases:

So what's going on?

**mutex is free:** the value loaded is false. We store true. The value returned is False, so we don't spin

**mutex is taken:** the value loaded is true, we put the SAME value back (true). The returned value is true, so we spin.

# First example: Exchange Lock

```
void unlock() {  
    flag.store(false);  
}
```

Unlock is simple: just store false to the flag,  
marking the mutex as available.

# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
m.lock();  
m.unlock();

Thread 1:  
m.lock();  
m.unlock();

```
void unlock() {  
    flag.store(false);  
}
```

core 0



core 1



# Analysis

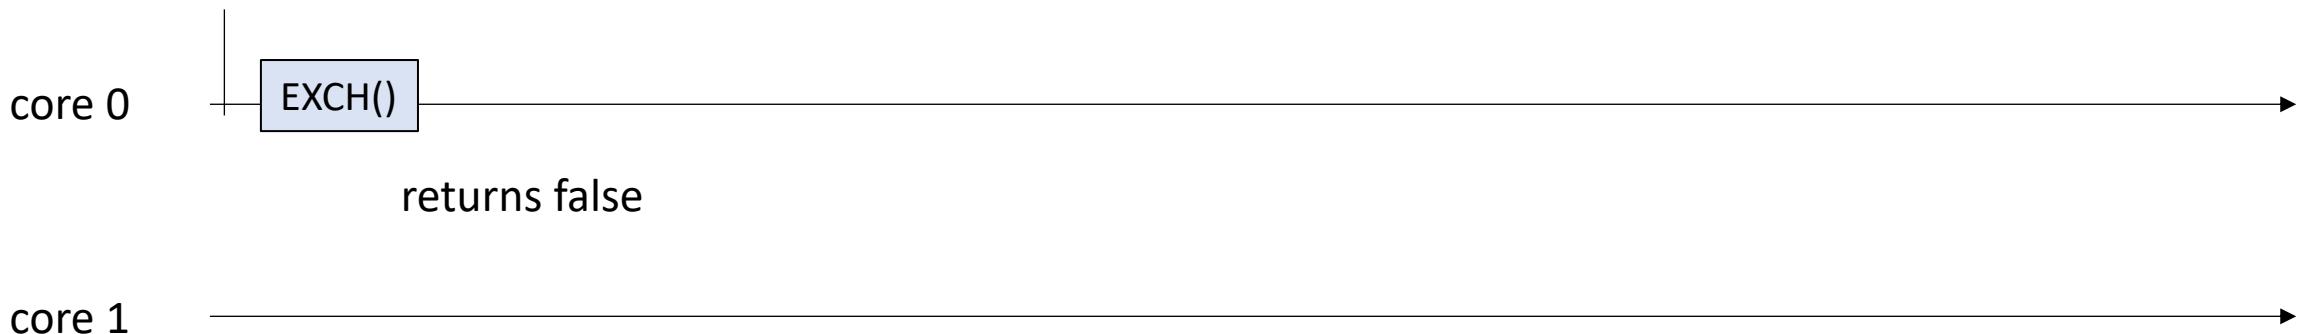
```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```

Mutex request



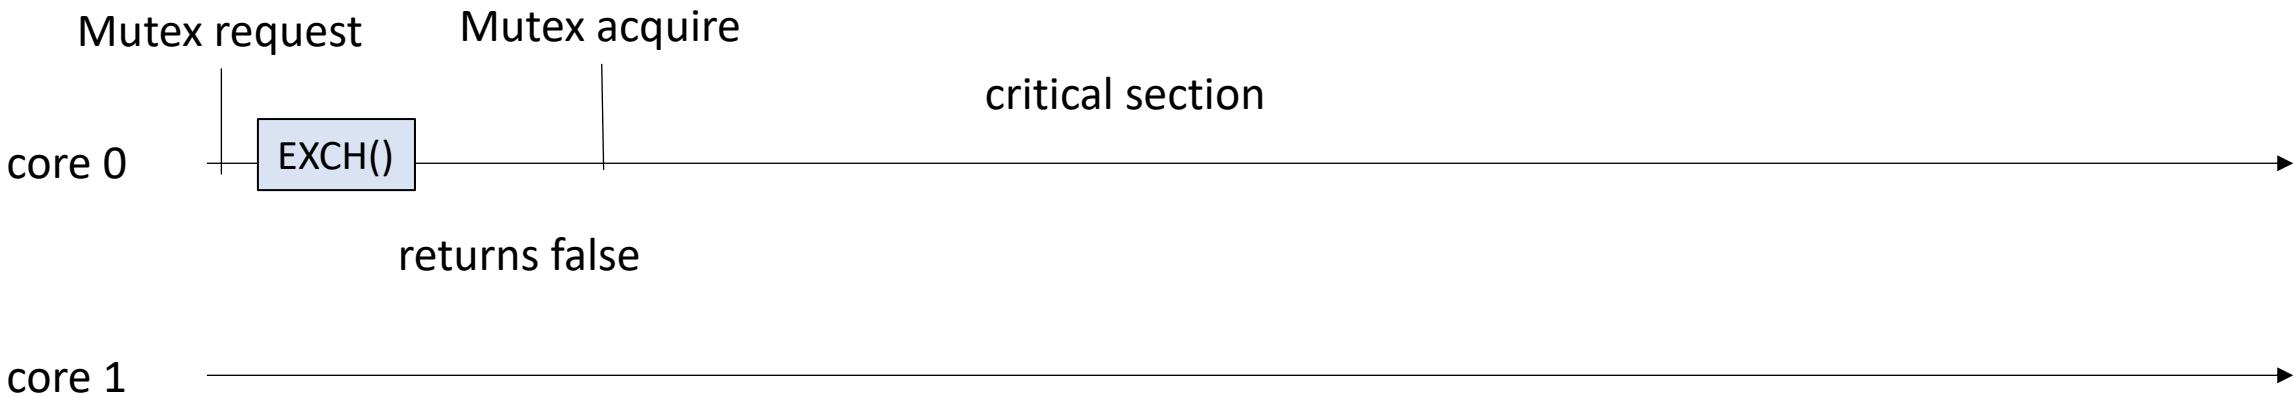
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



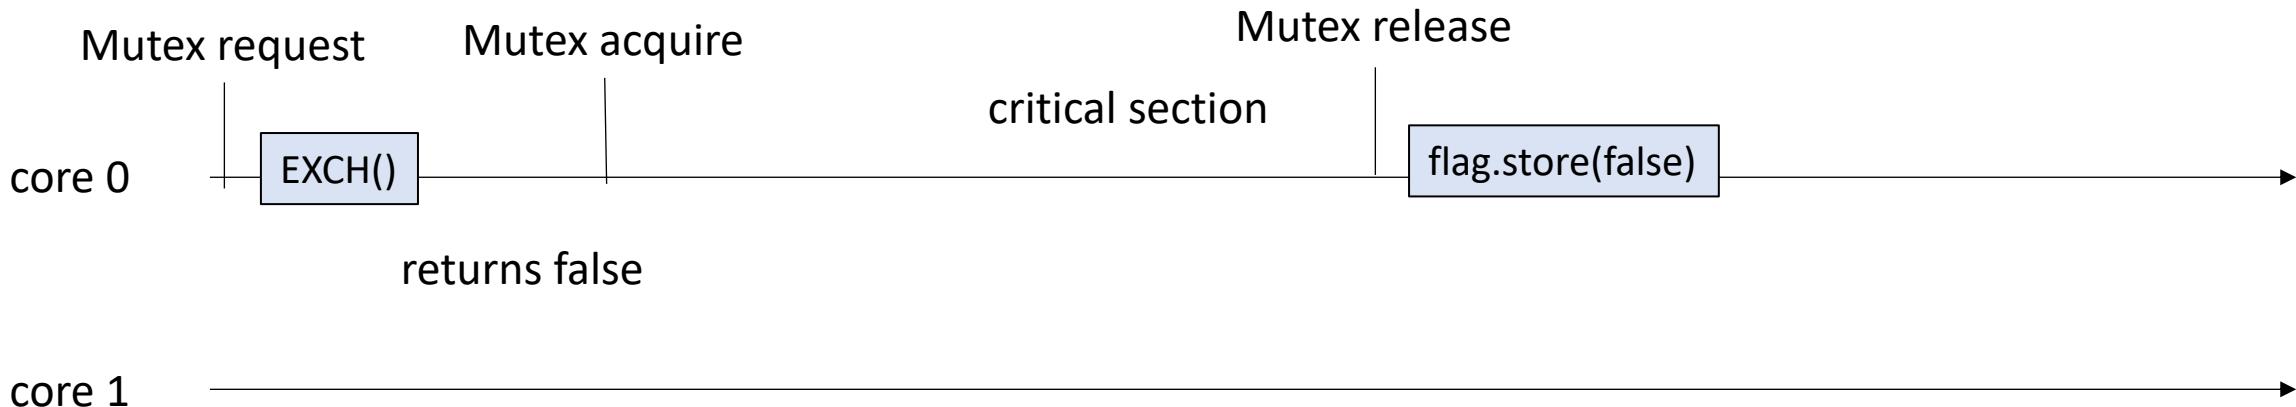
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



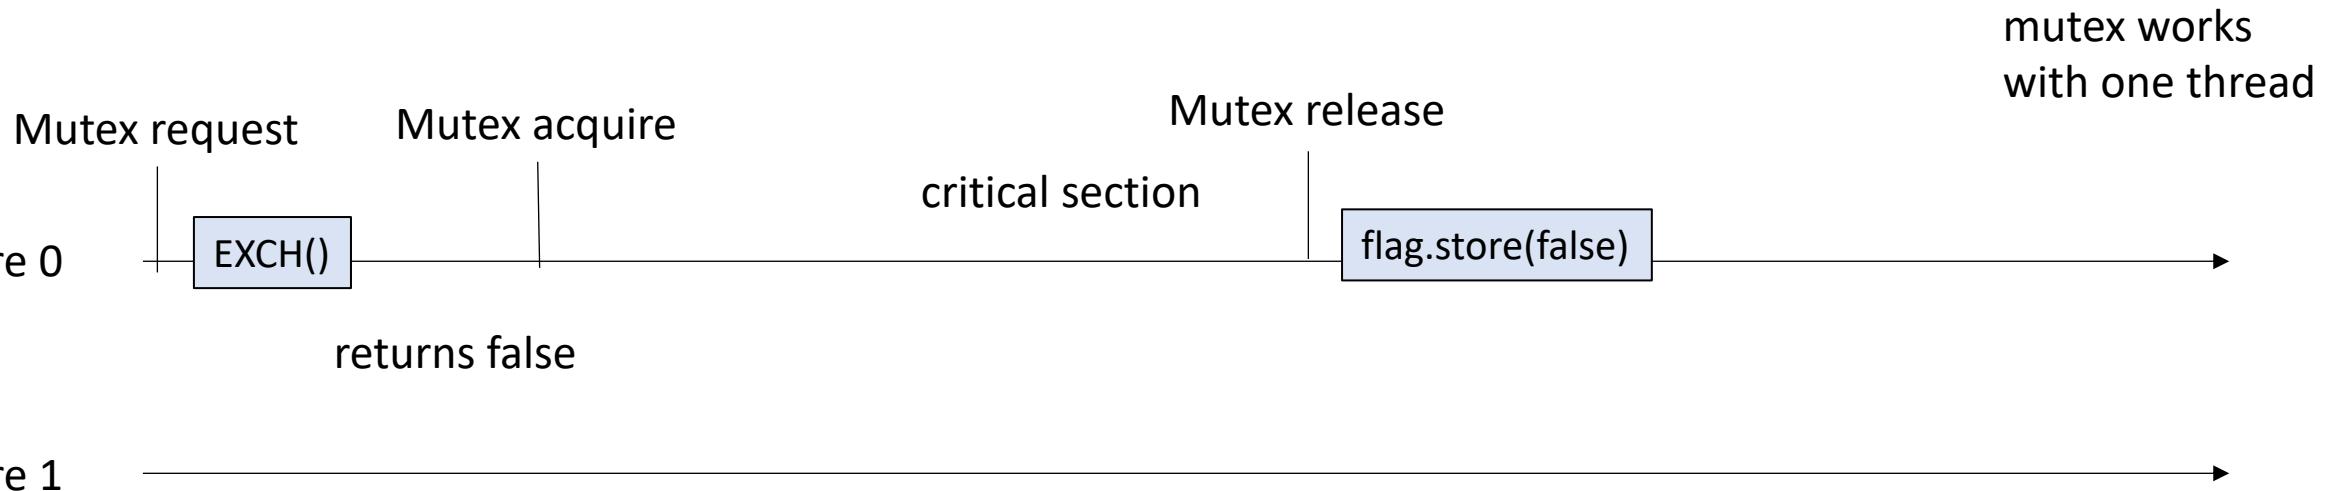
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



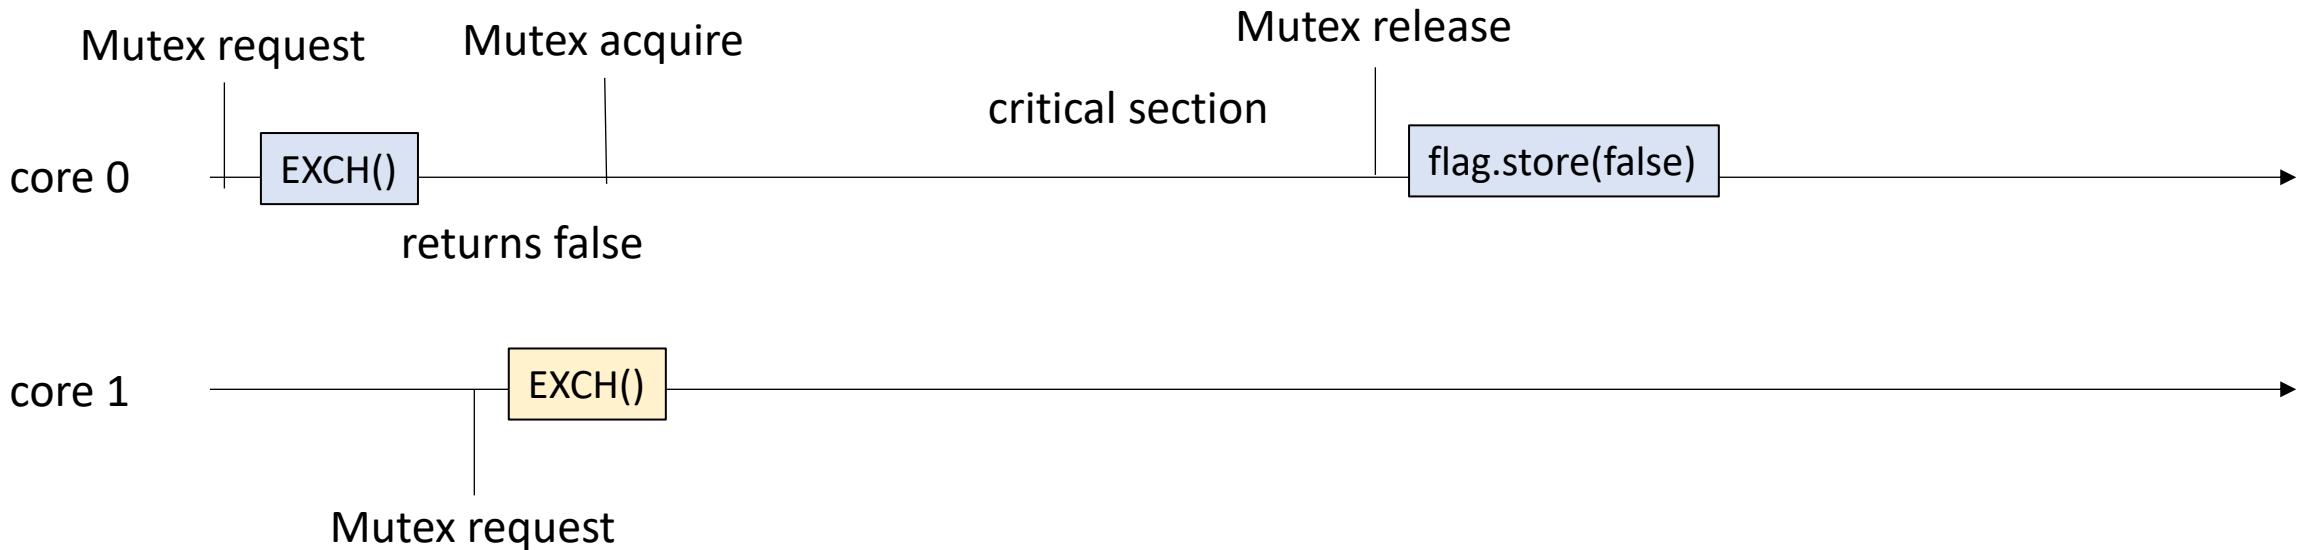
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



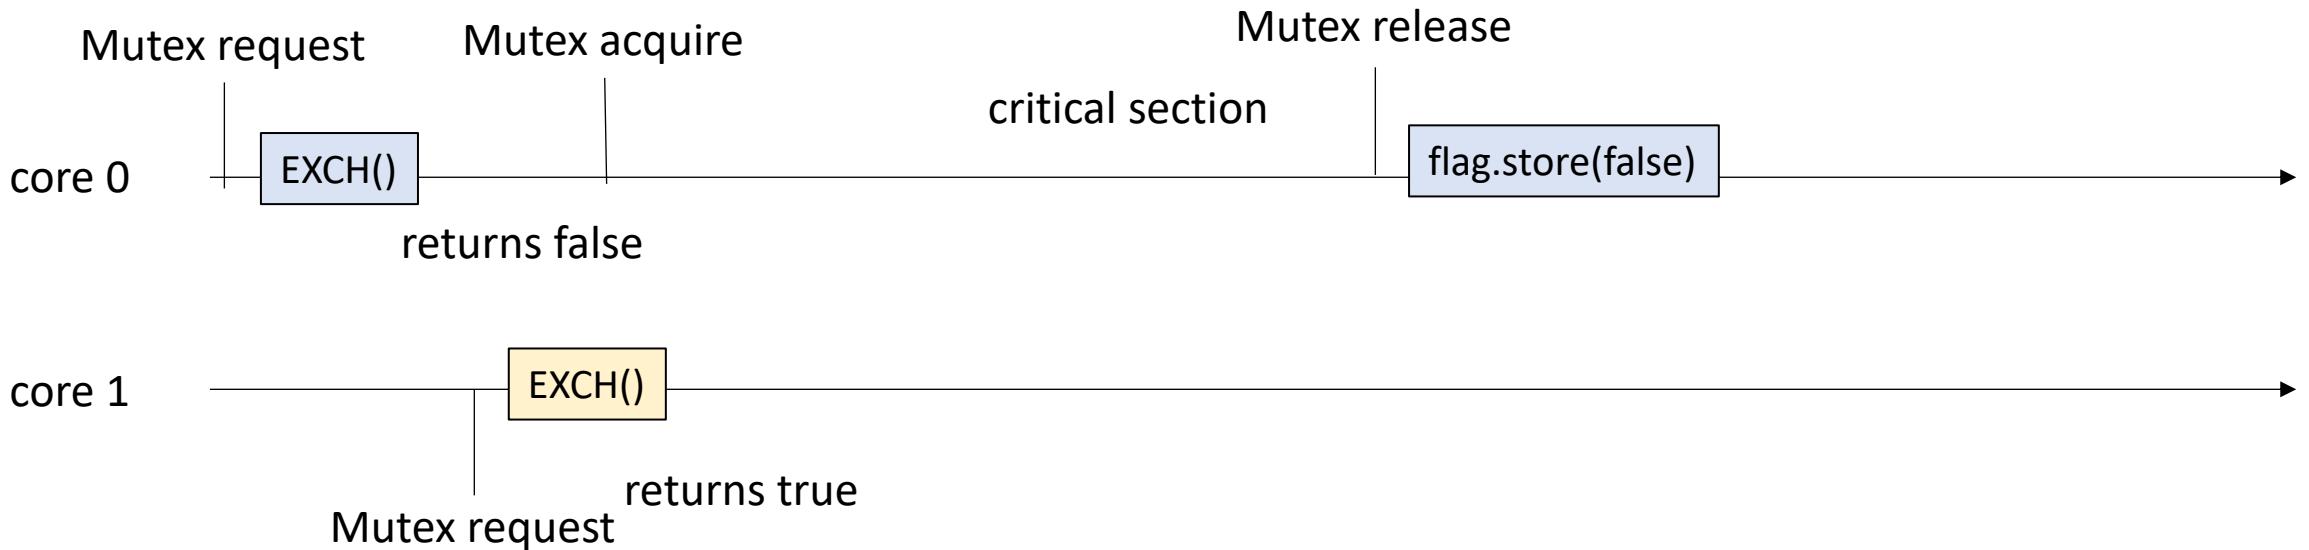
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



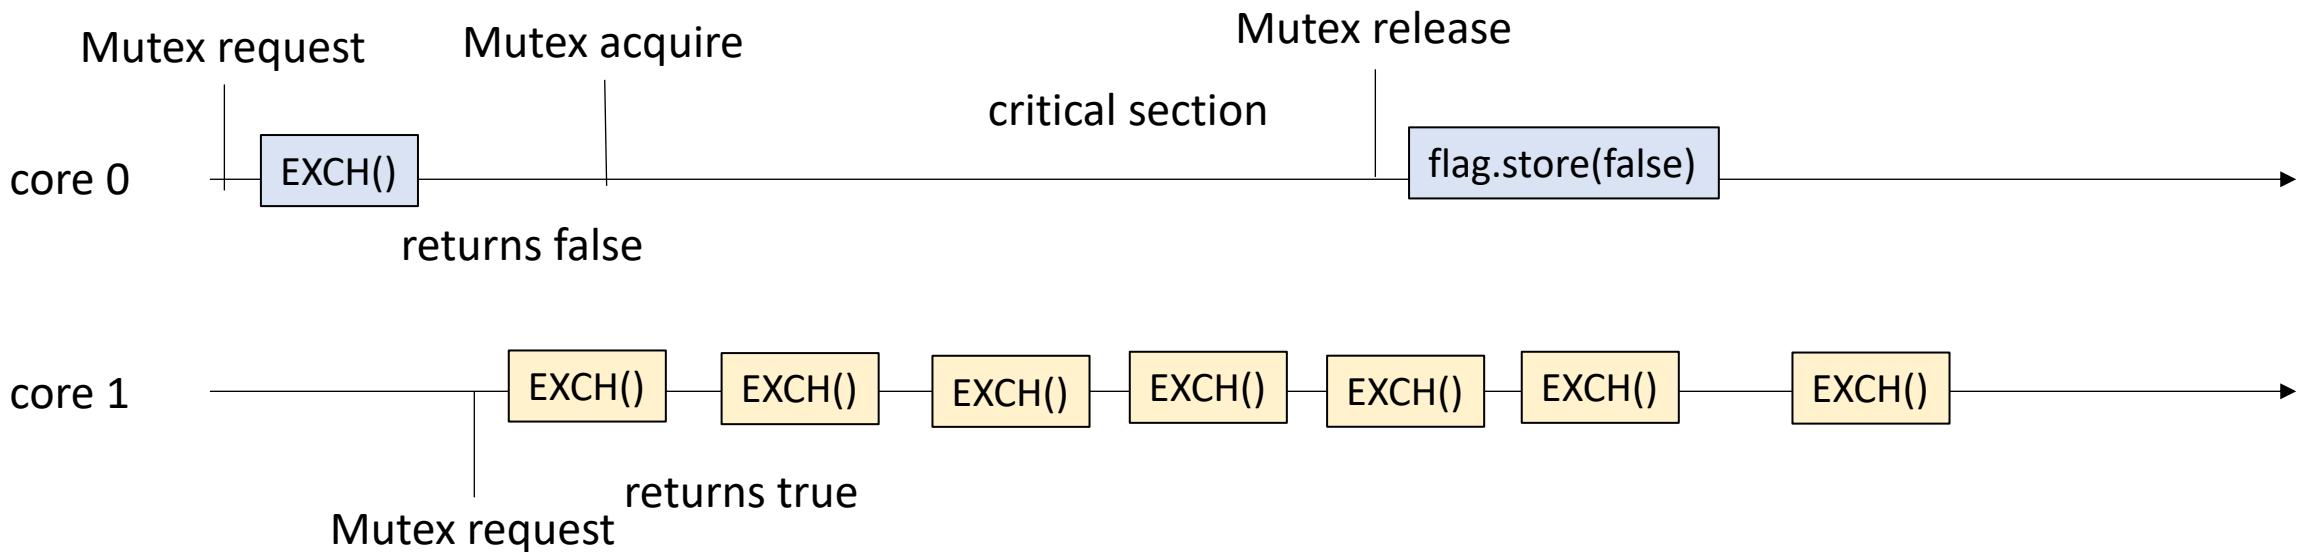
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



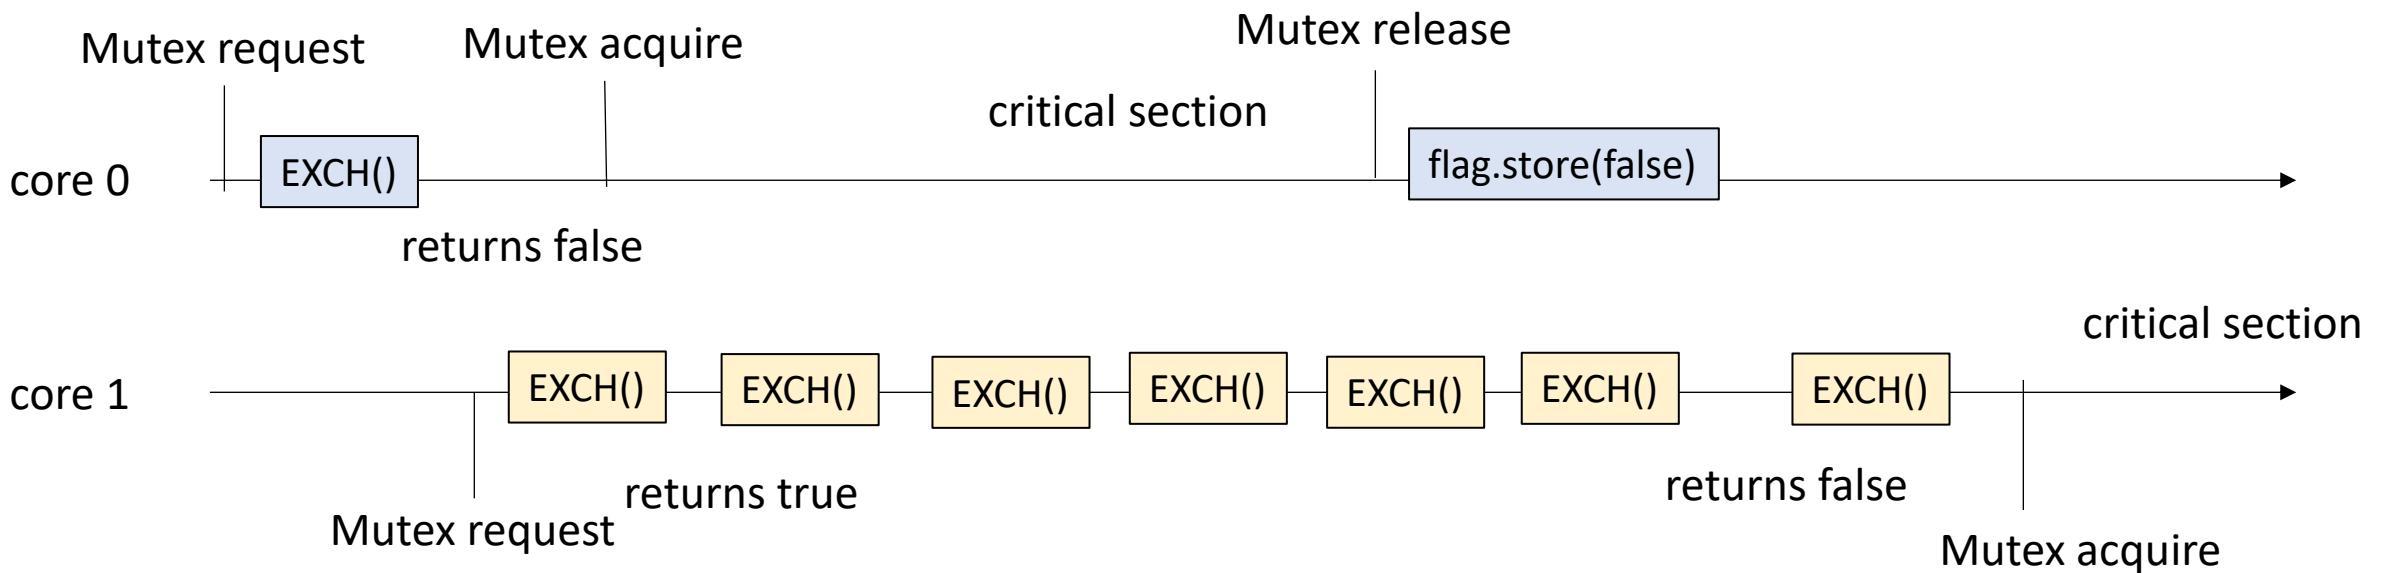
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



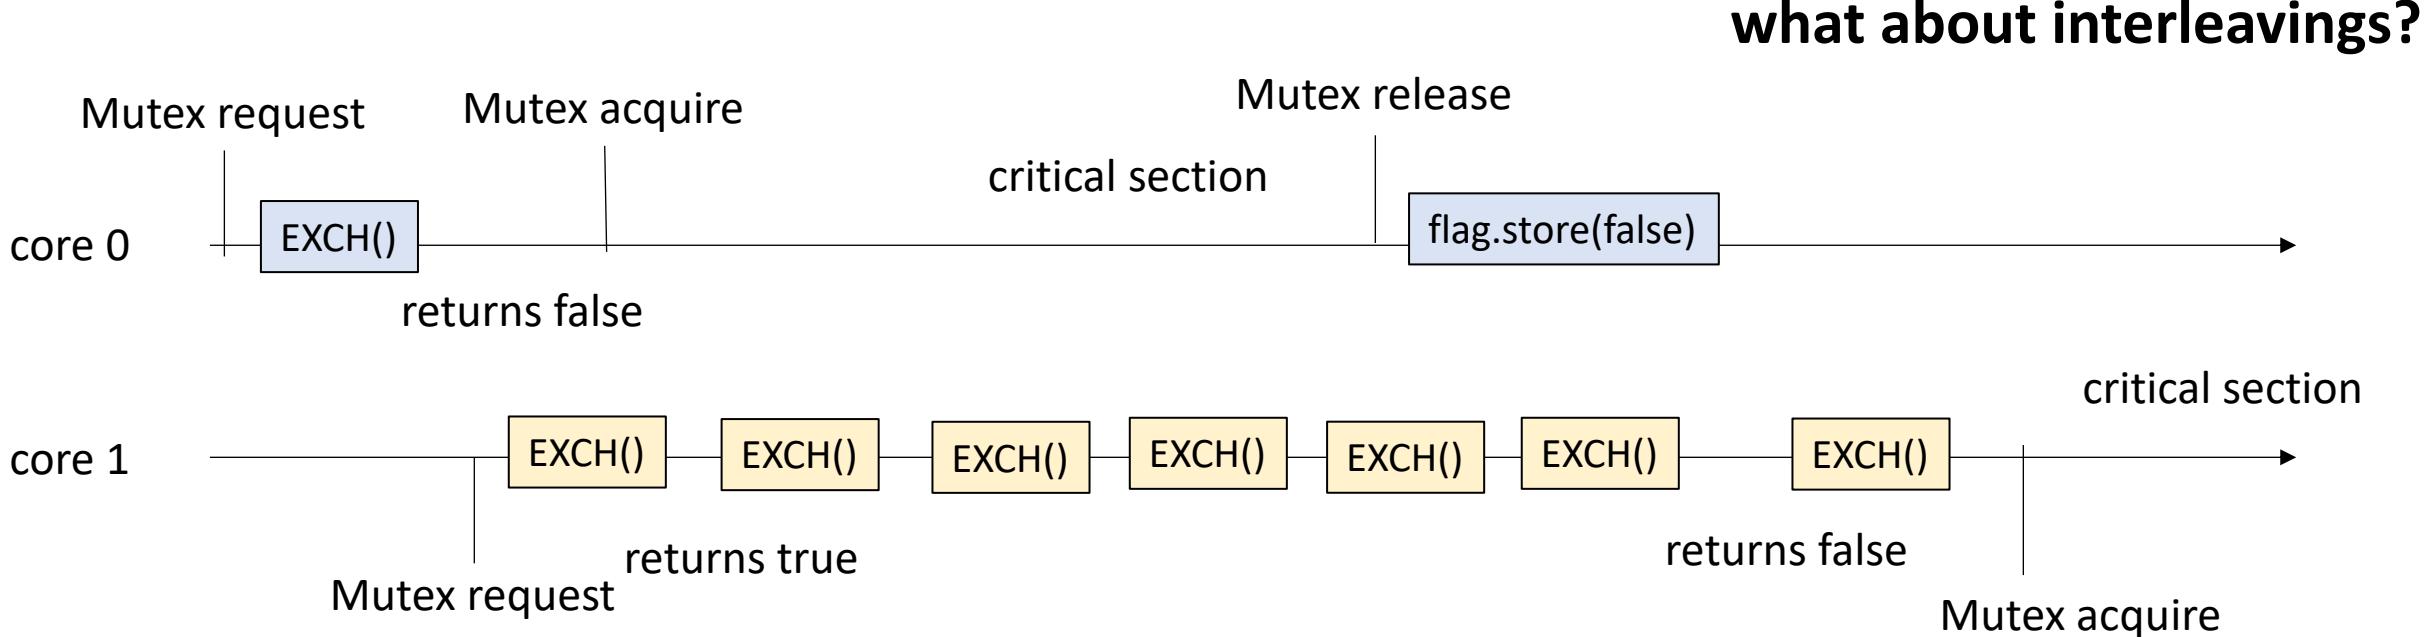
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```



# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```

**what about interleavings?**

Mutex request



core 1



Mutex request

# Analysis

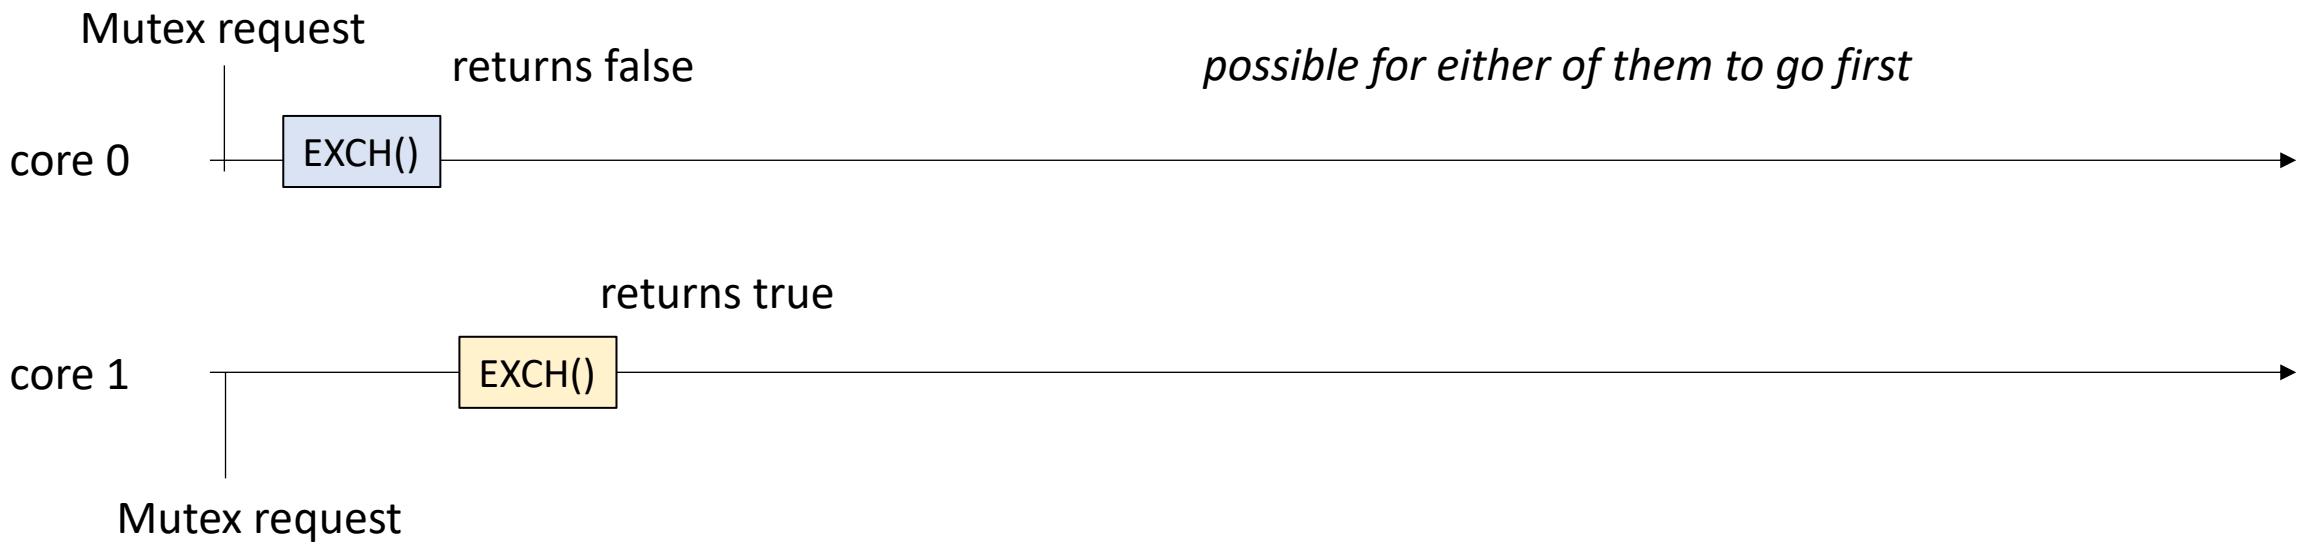
```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```

**recall RMWS can't overlap!**



# Analysis

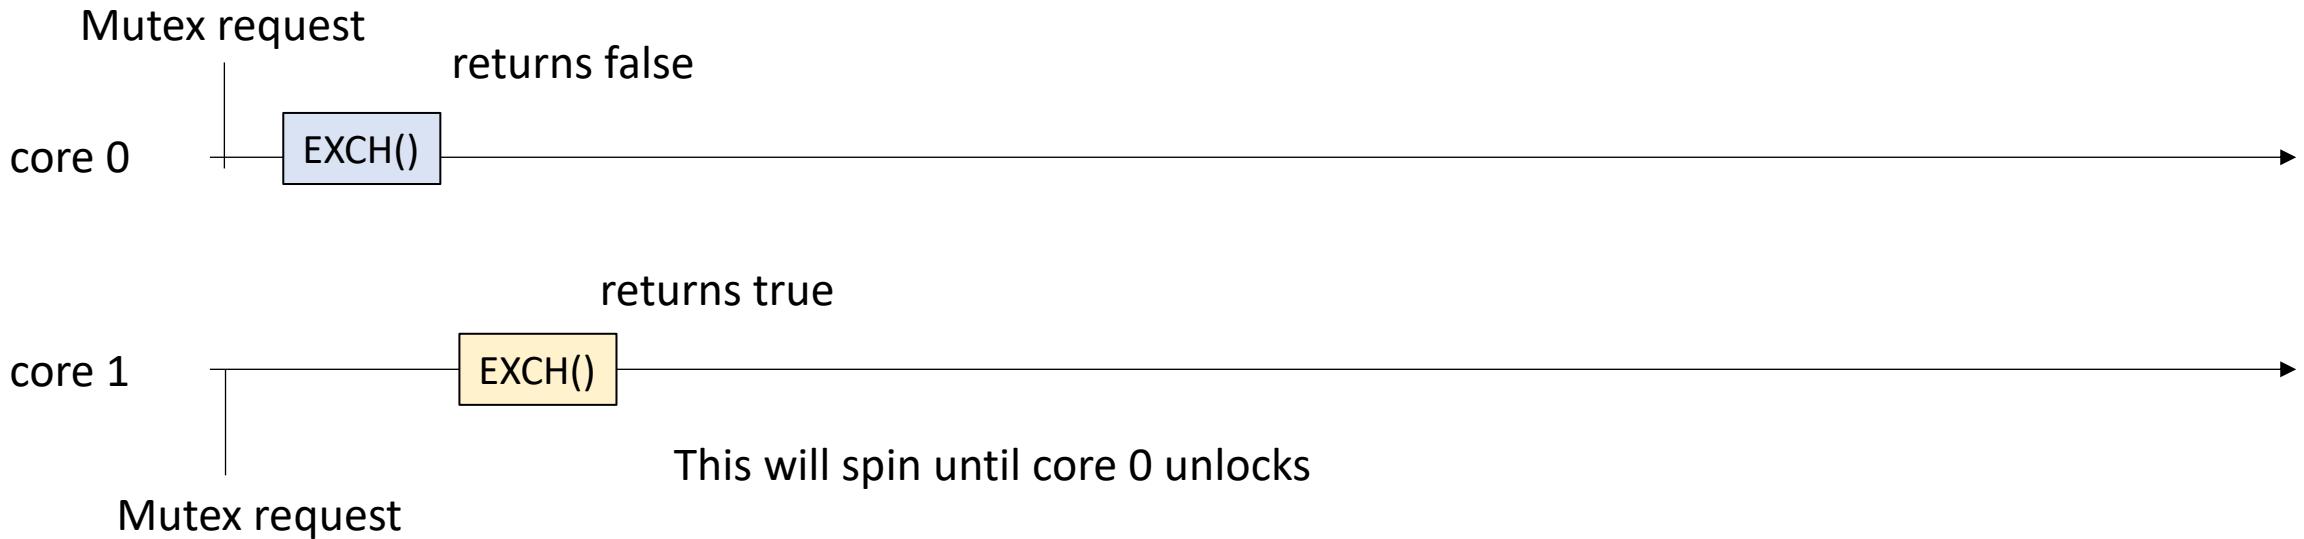
```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```

**recall RMWS can't overlap!**



# Analysis

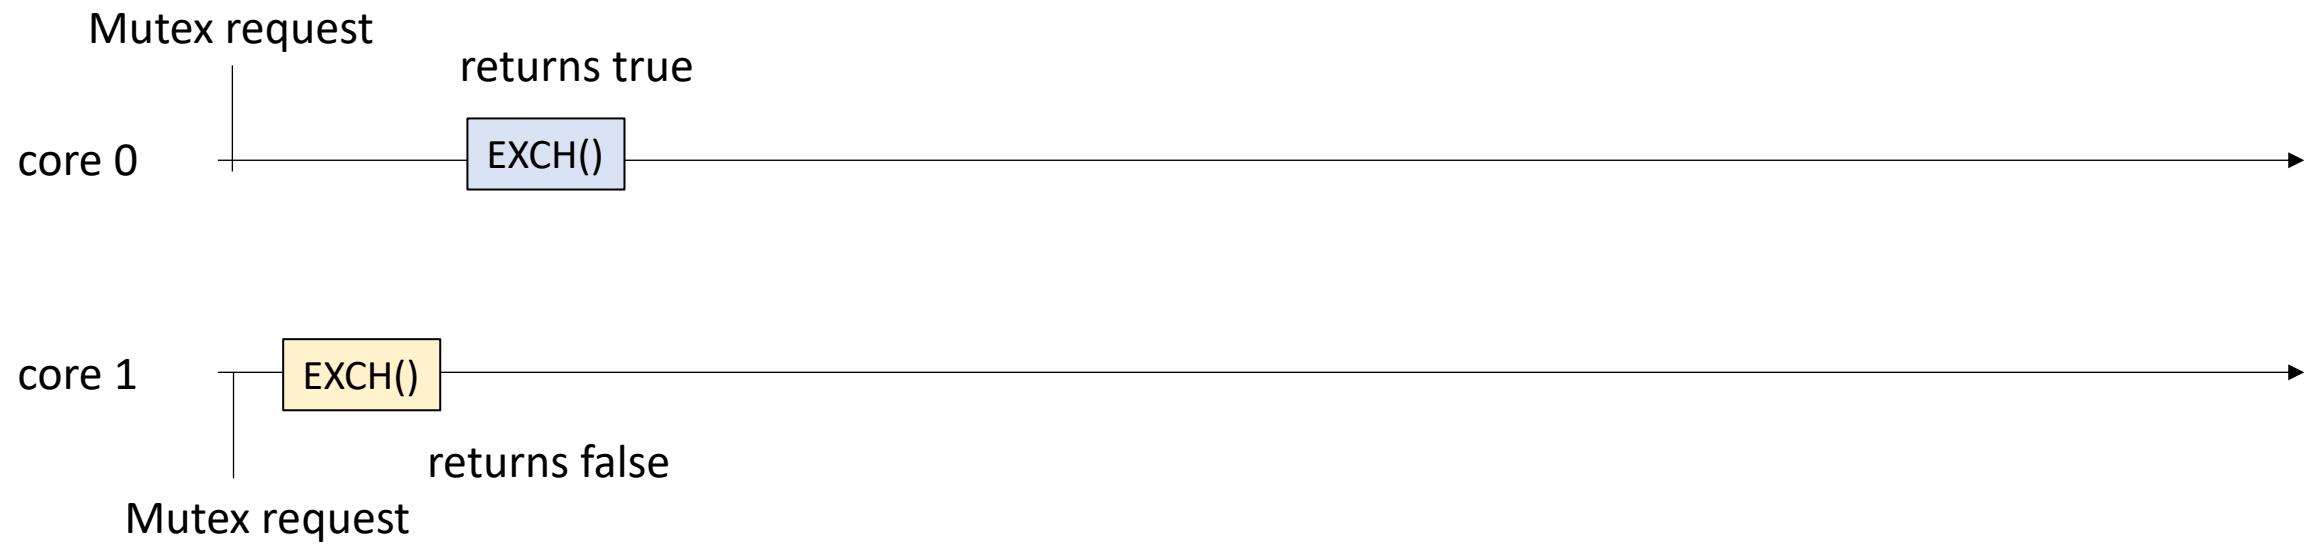
```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
m.lock();  
m.unlock();

Thread 1:  
m.lock();  
m.unlock();

```
void unlock() {  
    flag.store(false);  
}
```

recall RMWS can't overlap!



# Analysis

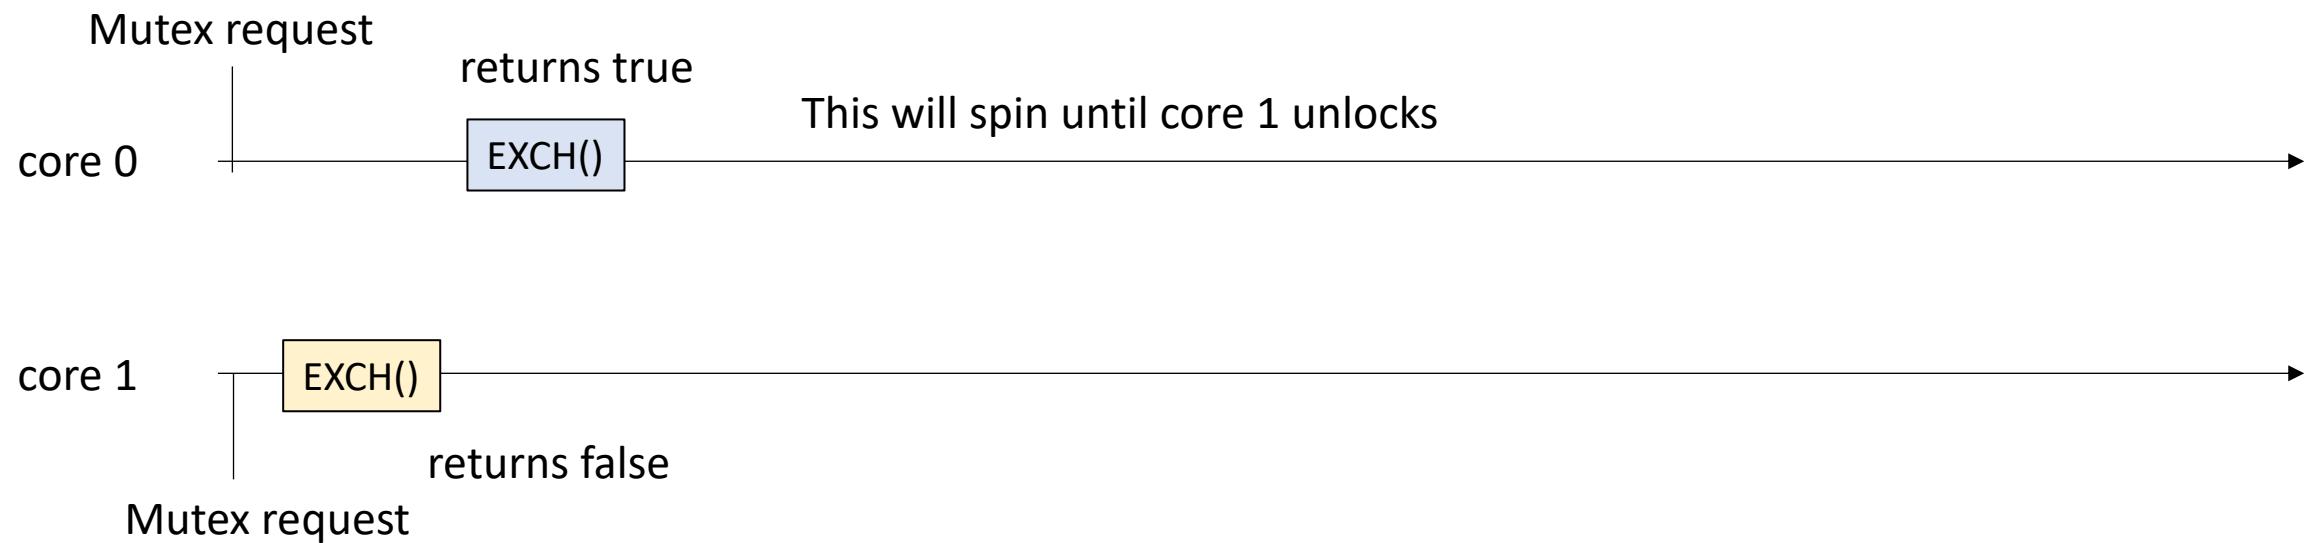
```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Thread 0:  
`m.lock();`  
`m.unlock();`

Thread 1:  
`m.lock();`  
`m.unlock();`

```
void unlock() {  
    flag.store(false);  
}
```

**recall RMWS can't overlap!**

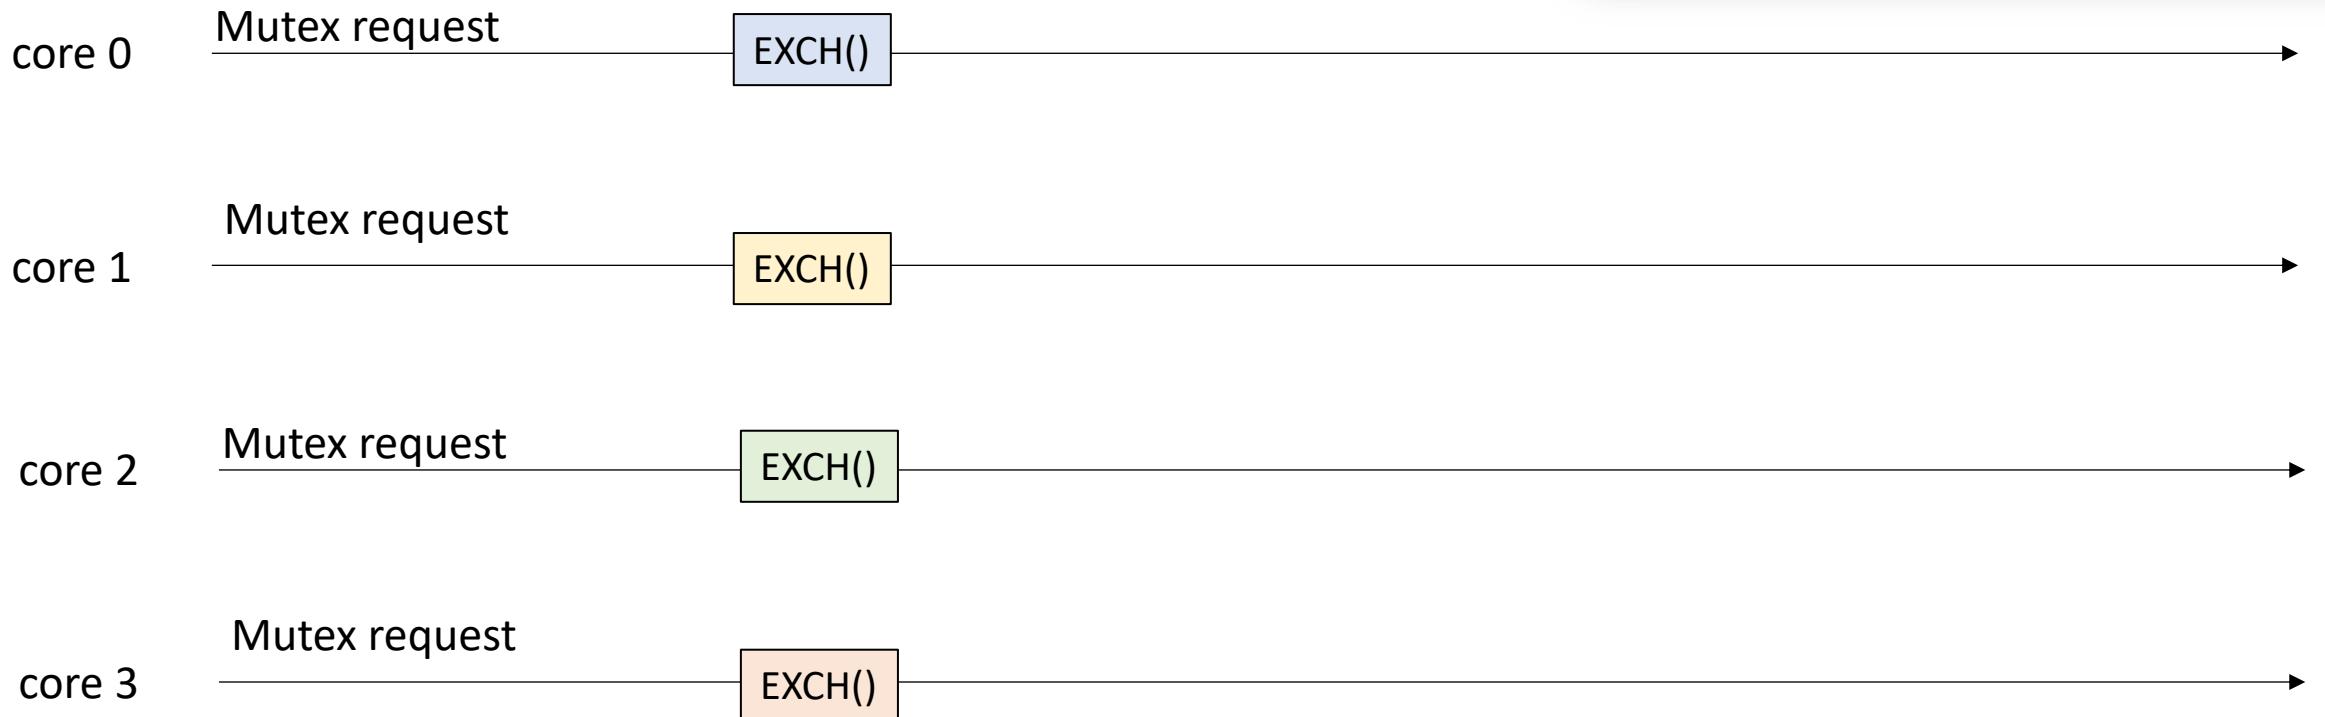


# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

```
void unlock() {  
    flag.store(false);  
}
```



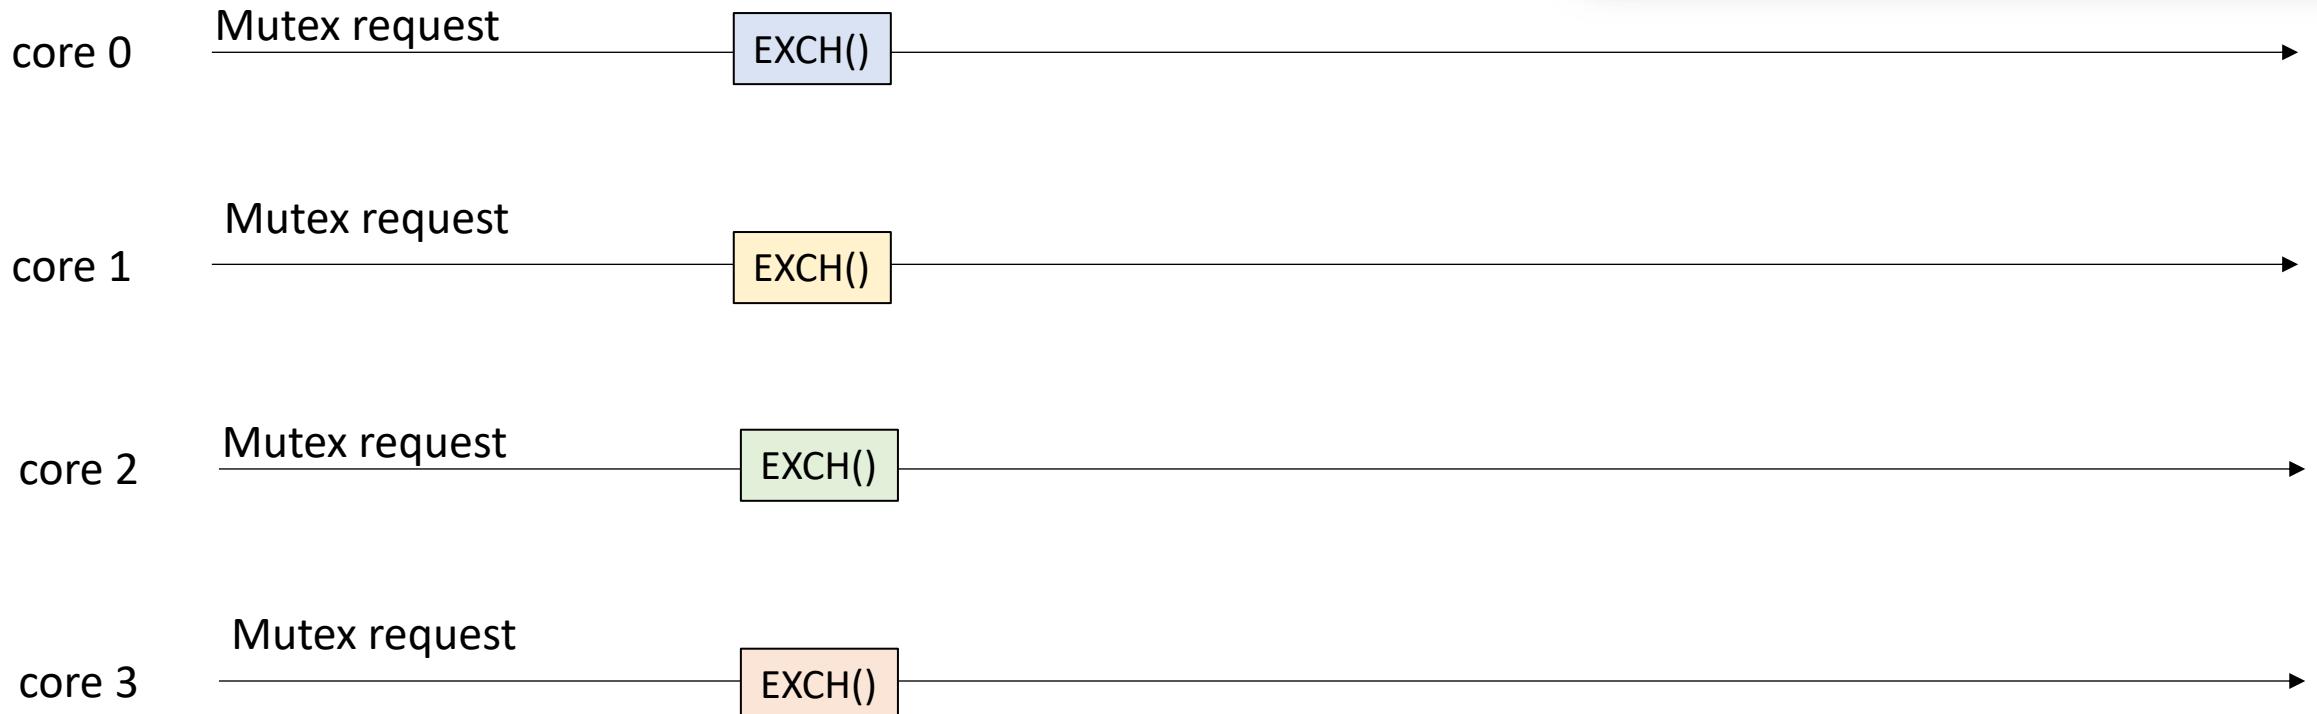
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



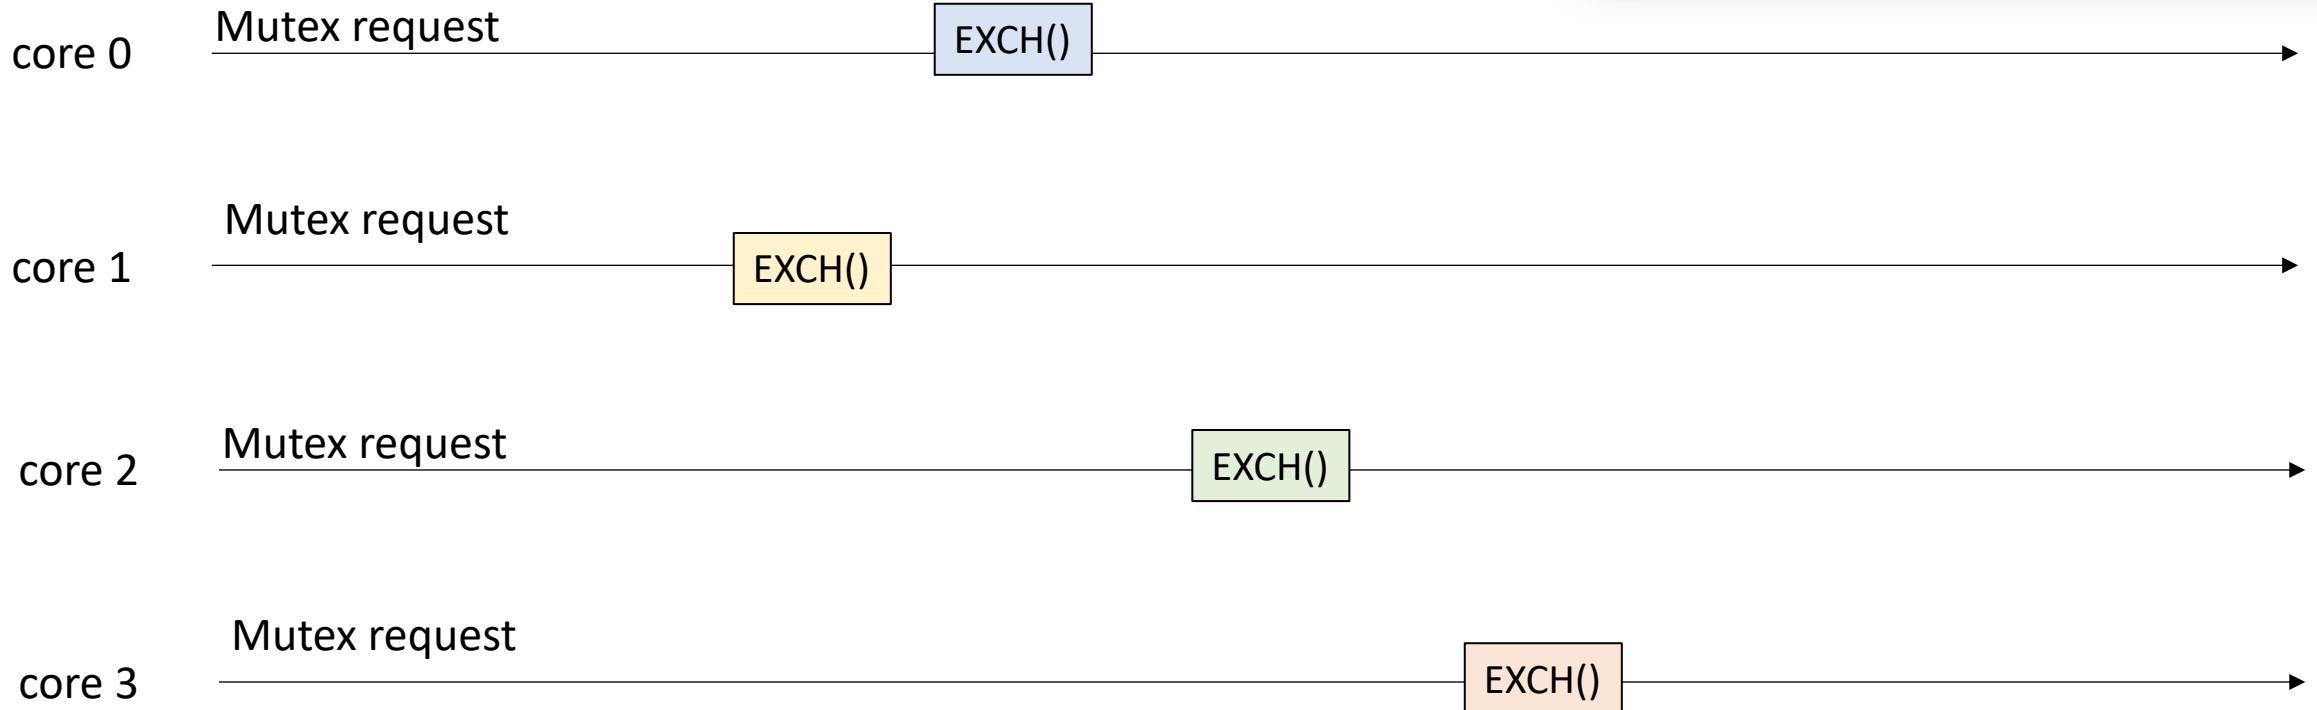
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



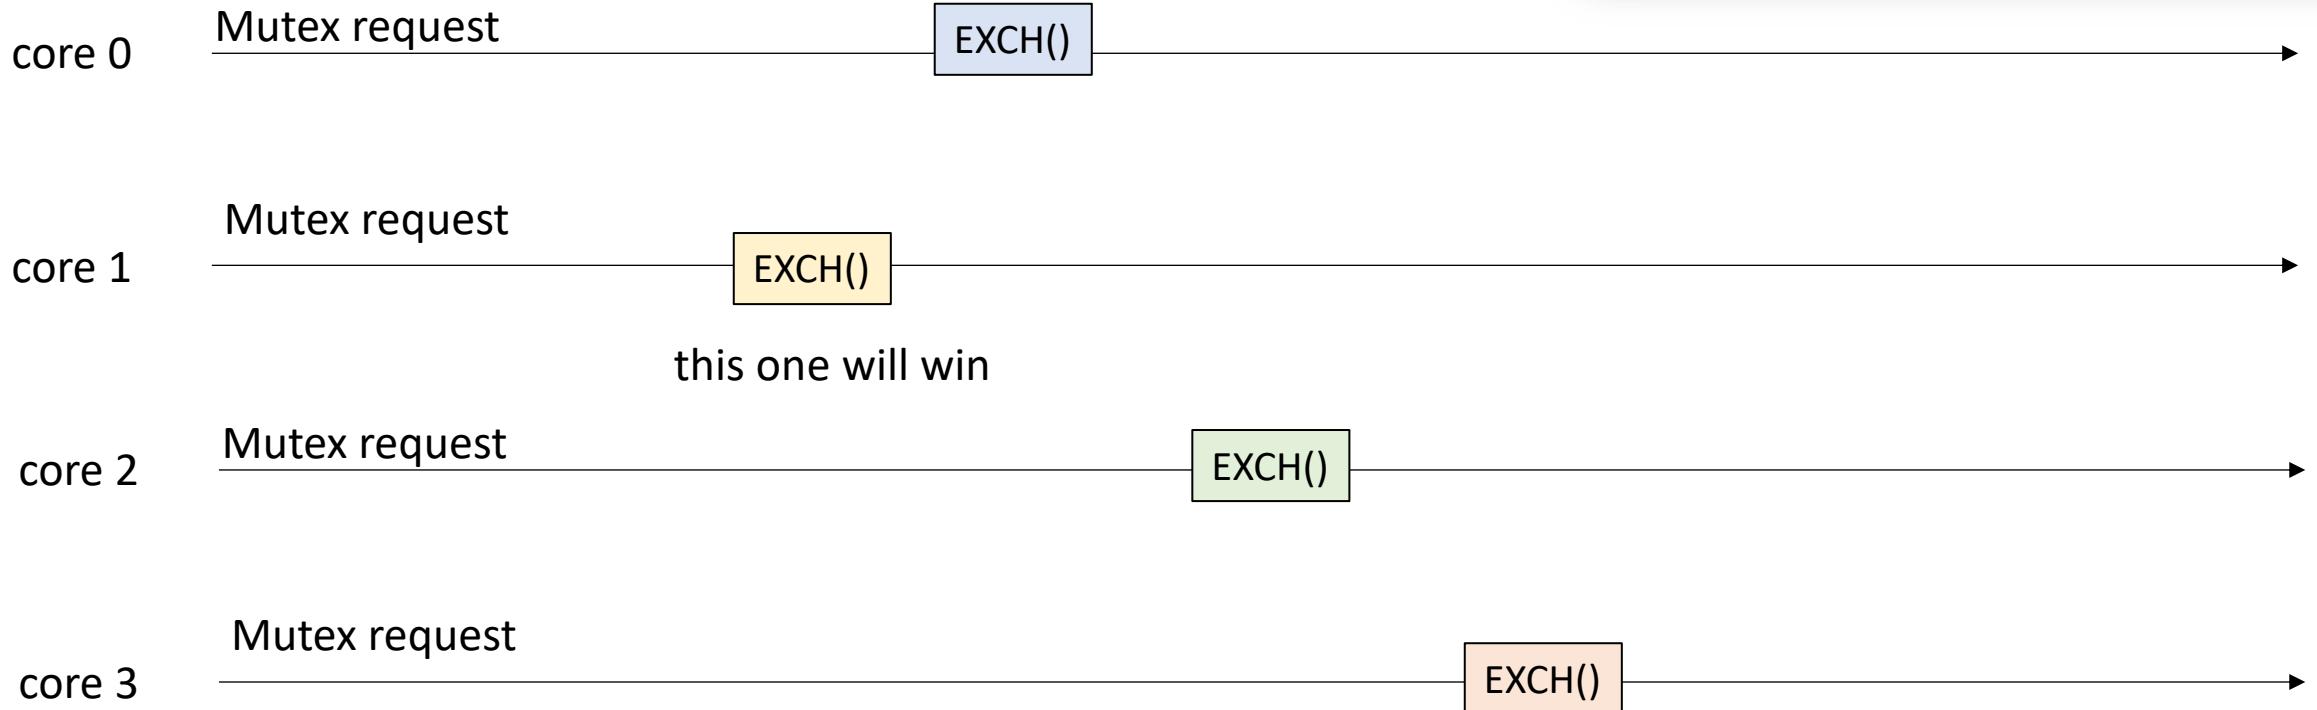
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



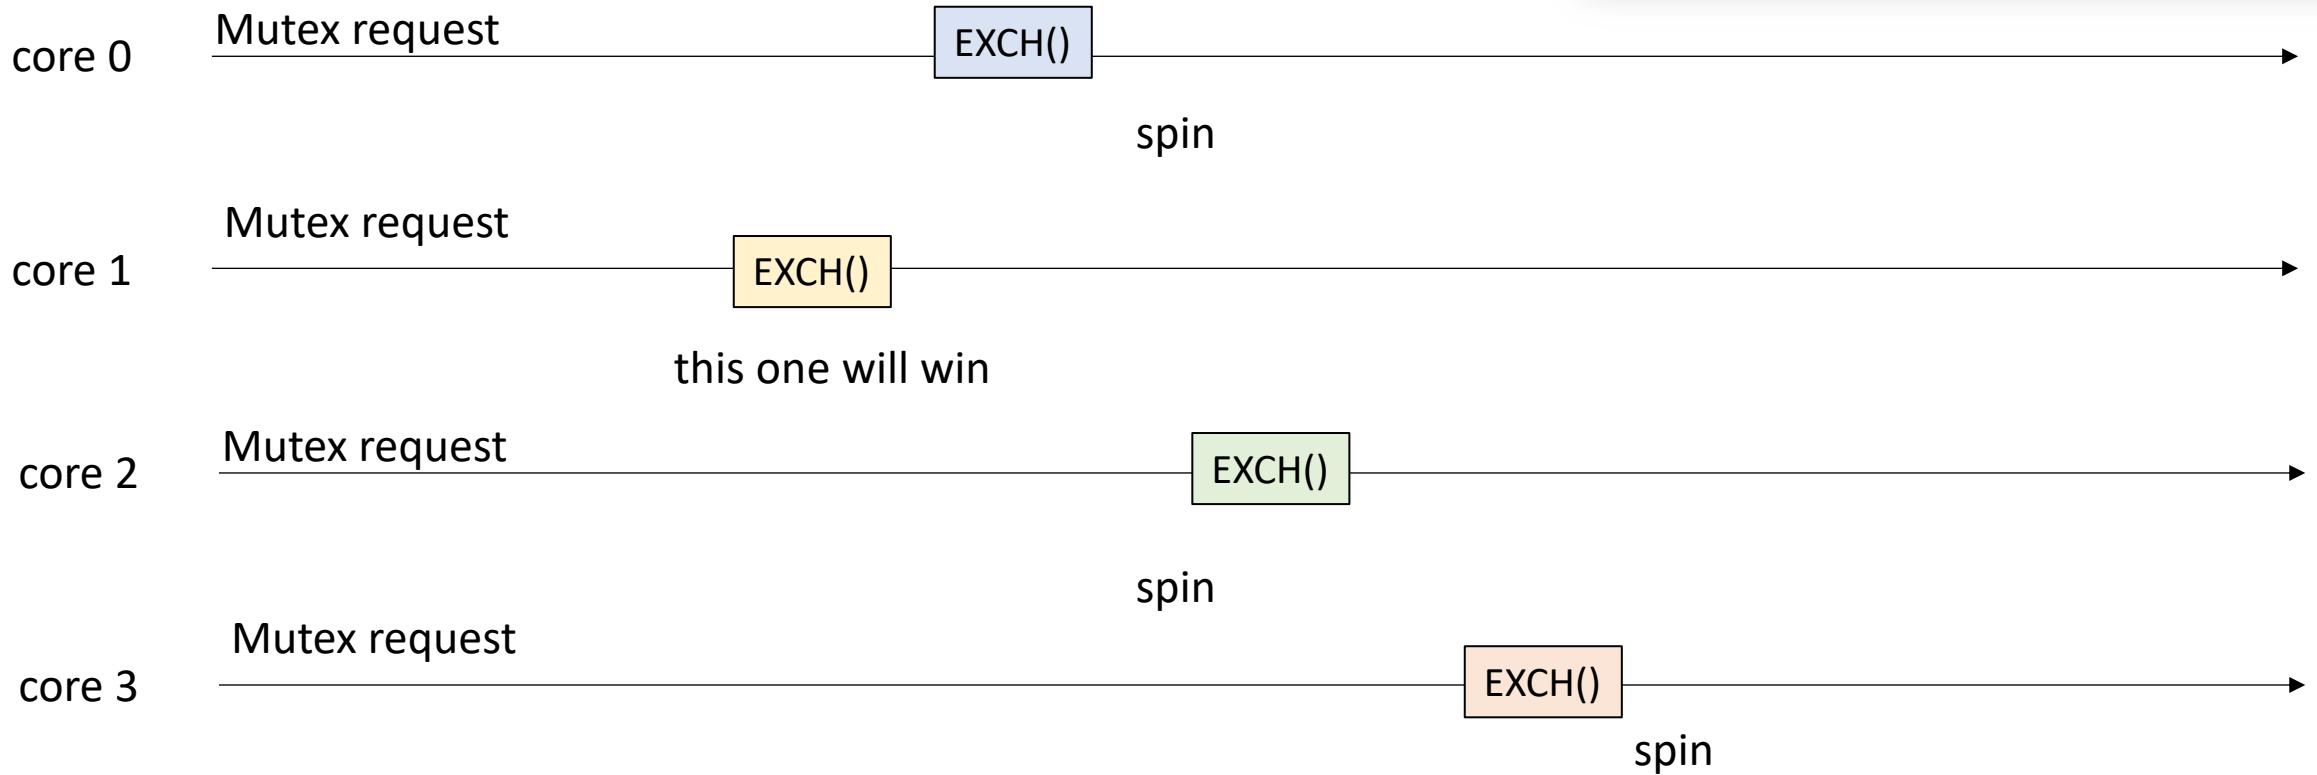
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



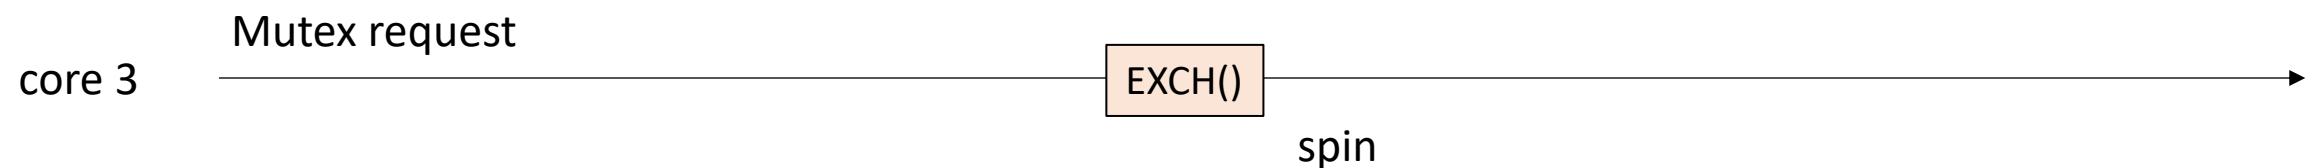
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



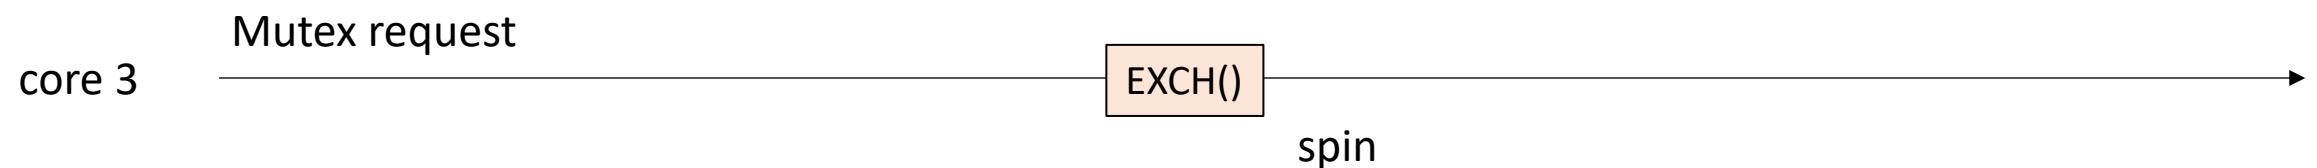
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



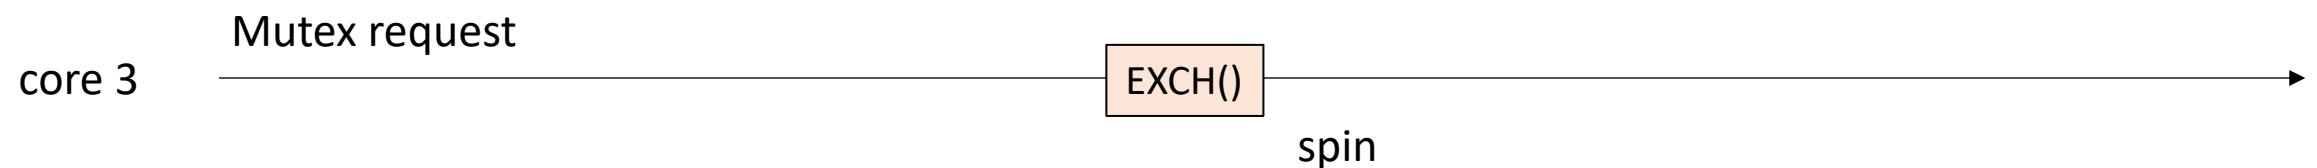
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



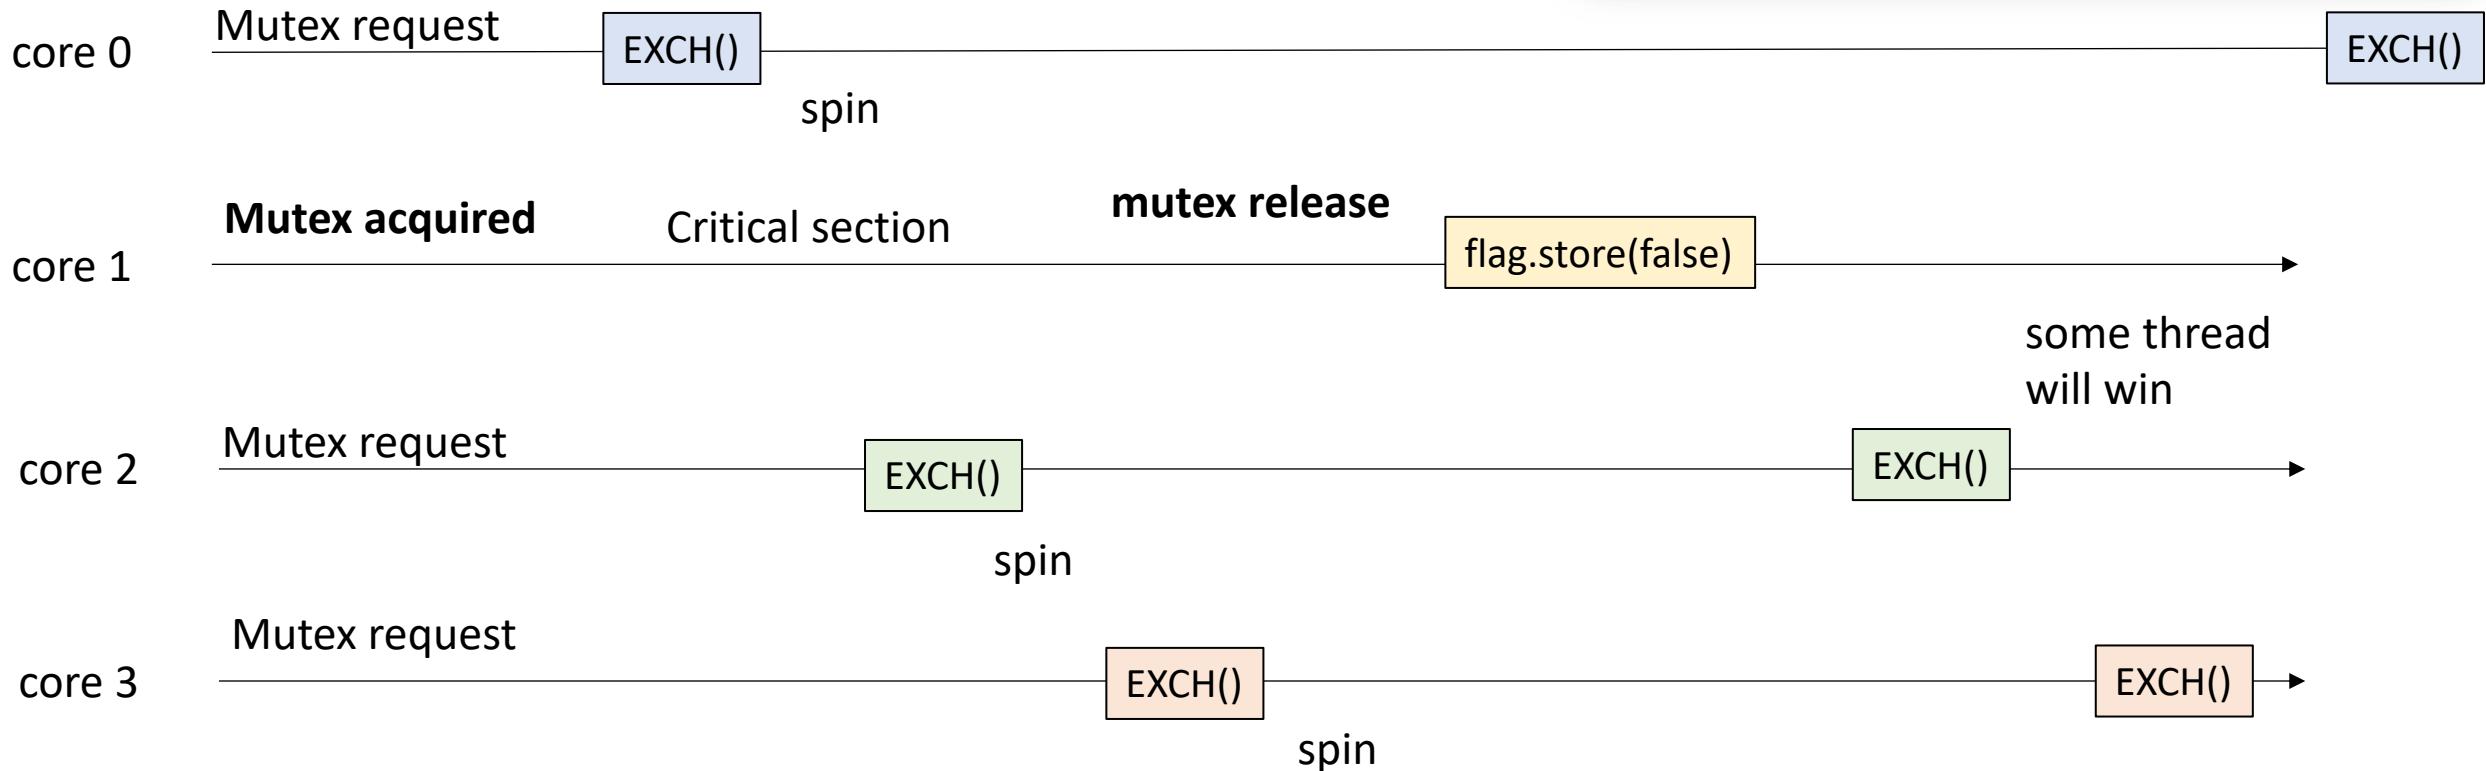
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



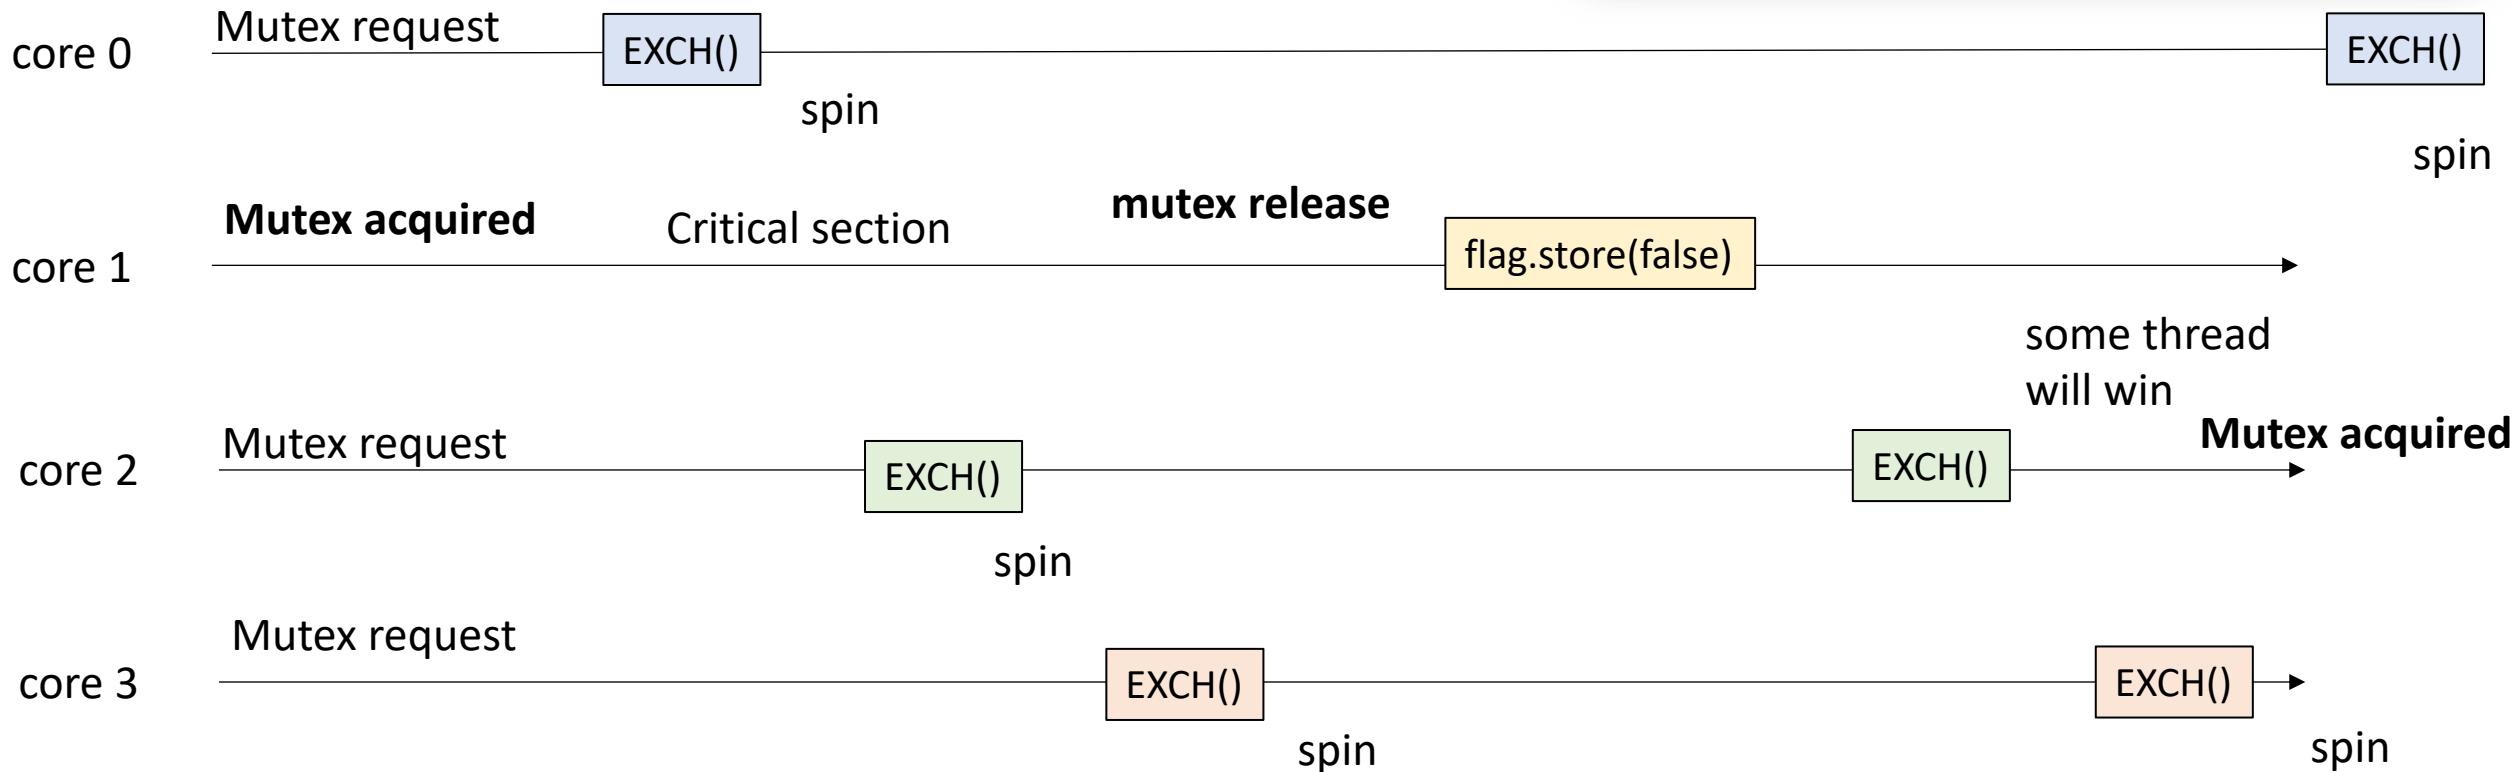
# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*what about 4 threads?*

atomic operations can't overlap

```
void unlock() {  
    flag.store(false);  
}
```



# First example: Exchange Mutex

- Questions?

# Lecture Schedule

- **Atomic RMW mutexes**

- Exchange
- CAS
- Ticket

- Optimizations

- Relaxed peeking
- Backoff

# Most versatile RMW: Compare-and-swap

- Exchange was the simplest RMW (no modify)
- Most versatile RMW: Compare-and-swap (CAS)

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace);
```

# Most versatile RMW: Compare-and-swap

- Exchange was the simplest RMW (no modify)
- Most versatile RMW: Compare-and-swap (CAS)

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace);
```

Checks if value at `a` is equal to the value at `expected`. If it is equal, swap with `replace`. Returns `True` if the values were equal. `False` otherwise.

# Most versatile RMW: Compare-and-swap

- Exchange was the simplest RMW (no modify)
- Most versatile RMW: Compare-and-swap (CAS)

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace);
```

Checks if value at `a` is equal to the value at `expected`. If it is equal, swap with `replace`.  
returns `True` if the values were equal. `False` otherwise.

`expected` is passed by reference: the previous value at `a` is returned

# Most versatile RMW: Compare-and-swap

- Exchange was the simplest RMW (no modify)
- Most versatile RMW: Compare-and-swap (CAS)

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

# Most versatile RMW: Compare-and-swap

- Exchange was the simplest RMW (no modify)
- Most versatile RMW: Compare-and-swap (CAS)

*we will discuss  
this soon!*

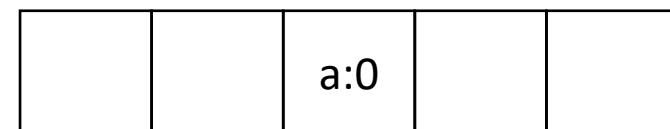
```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

# Most versatile RMW: Compare-and-swap

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

thread 0:

```
// some atomic int address a
int e = 0;
bool s = atomic_CAS(a,&e,6);
```

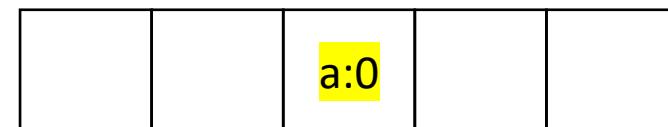


# Most versatile RMW: Compare-and-swap

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

thread 0:

```
// some atomic int address a
int e = 0;
bool s = atomic_CAS(a,&e,6);
```

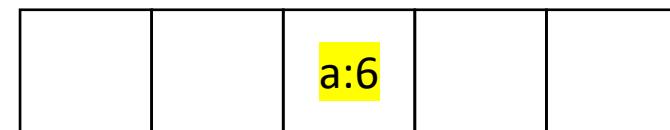


# Most versatile RMW: Compare-and-swap

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

thread 0:

```
// some atomic int address a
int e = 0;
bool s = atomic_CAS(a,&e,6);
```

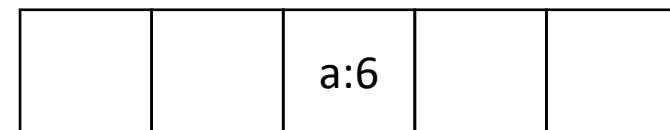


# Most versatile RMW: Compare-and-swap

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

thread 0:

```
// some atomic int address a
int e = 0;
bool s = atomic_CAS(a,&e,6);
true
```

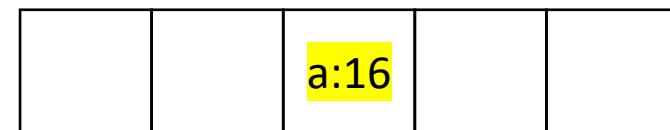


# Most versatile RMW: Compare-and-swap

```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

thread 0:

```
// some atomic int address a
int e = 0;
bool s = atomic_CAS(a,&e,6);
```



# Most versatile RMW: Compare-and-swap

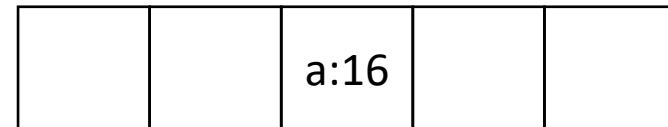
```
bool atomic_compare_exchange_strong(atomic *a, value *expected, value replace) {
    value tmp = a.load();
    if (tmp == *expected) {
        a.store(replace);
        return true;
    }
    *expected = tmp;
    return false;
}
```

thread 0:

```
// some atomic int address a
int e = 0;
bool s = atomic_CAS(a,&e,6);
```

false

16



# CAS lock

```
#include <atomic>
using namespace std;

class Mutex {
public:
    Mutex() {
        flag = false;
    }

    void lock();
    void unlock();

private:
    atomic_bool flag;
};
```

Pretty intuitive: only 1 bit required again:

# CAS lock

```
void lock() {  
    bool check = false;  
    while (atomic_compare_exchange_strong(&flag, &check, true) == false);  
}  
  
void unlock() {  
    flag.store(false);  
}
```

Check if the mutex is free, if so, take it.

compare the mutex to free (false), if so, replace it with taken (true). Spin while the thread isn't able to take the mutex.

# CAS: versatile

- Why do I say it is versatile?

# CAS: versatile

- Why do I say it is versatile?
- We can implement ANY other RMW using CAS!

# Implementing atomic\_fetch\_add

```
int atomic_fetch_add(atomic_int *a, value v) {
    // implement me using CAS
}
```

# Implementing atomic\_fetch\_add

```
int atomic_fetch_add(atomic_int *a, value v) {
    int old_val = a->load();
    int new_val = old_val + v;
    atomic_compare_exchange(a, &old_val, new_val);
}
```

# Implementing atomic\_fetch\_add

```
int atomic_fetch_add(atomic_int *a, value v) {
    do {
        int old_val = a->load();
        int new_val = old_val + v;
        bool success = atomic_compare_exchange(a, &old_val, new_val);
    } while (!success)
}
```

# Implementing atomic\_fetch\_add

could be any operation!

```
int atomic_fetch_add(atomic_int *a, value v) {
    do {
        int old_val = a->load();
        int new_val = old_val + v;
        bool success = atomic_compare_exchange(a, &old_val, new_val);
    } while (!success)
}
```

# Implementing RMWs with CAS

- Gives you access to a wide range of operations!
  - `atomic_fetch_add` for float (not often provided)
  - You have to be careful with bit casting
- Why might this be difficult to implement?
  - Not provided in C++
  - Not provided for GPUs either (generally)
  - But very useful, especially for reduction and flow algorithms

# How is CAS implemented?

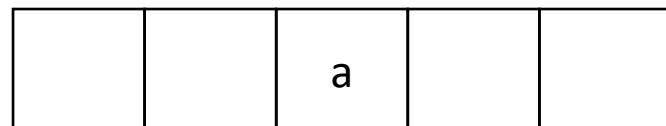
- X86 has an actual instruction
- ARM and POWER are load linked store conditional
- Show Godbolt example

# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start

thread 0:

```
atomic_CAS(a, ...);
```

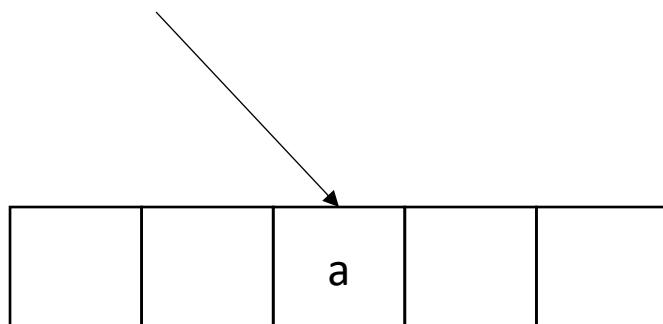


# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start

thread 0:

```
atomic_CAS(a, ...);
```



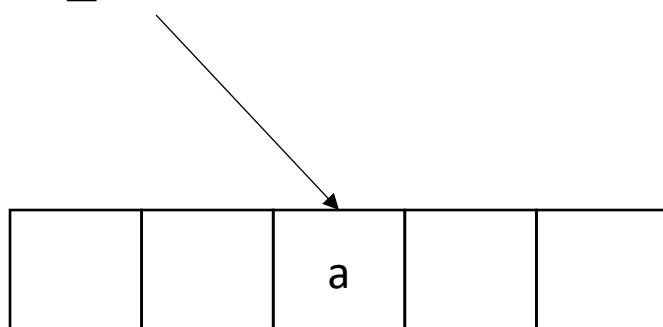
no other thread can access

# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start

thread 0:  
`atomic_CAS(a, ...);`

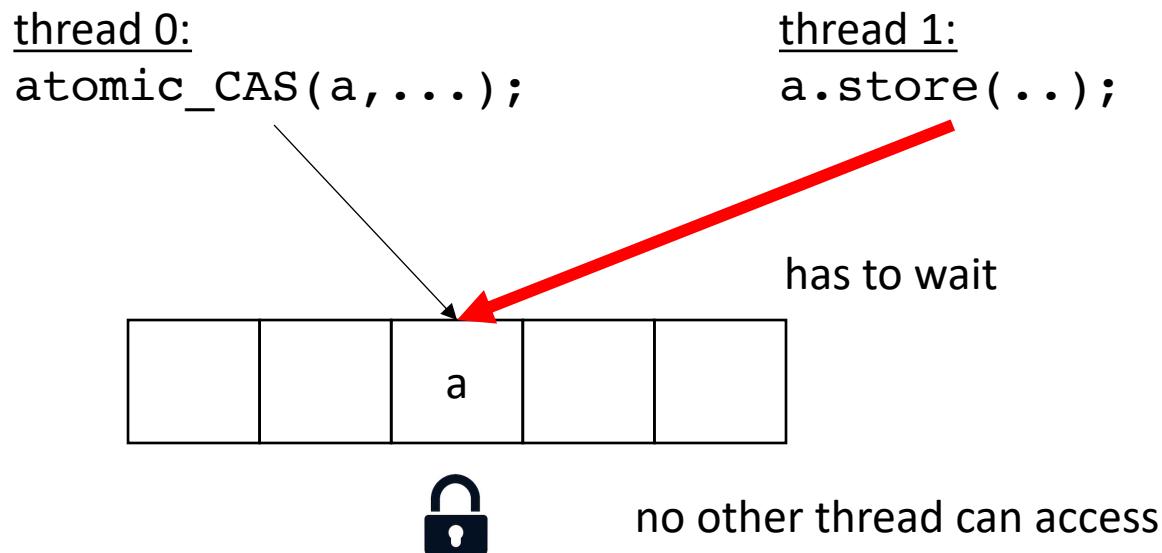
thread 1:  
`a.store(..);`



no other thread can access

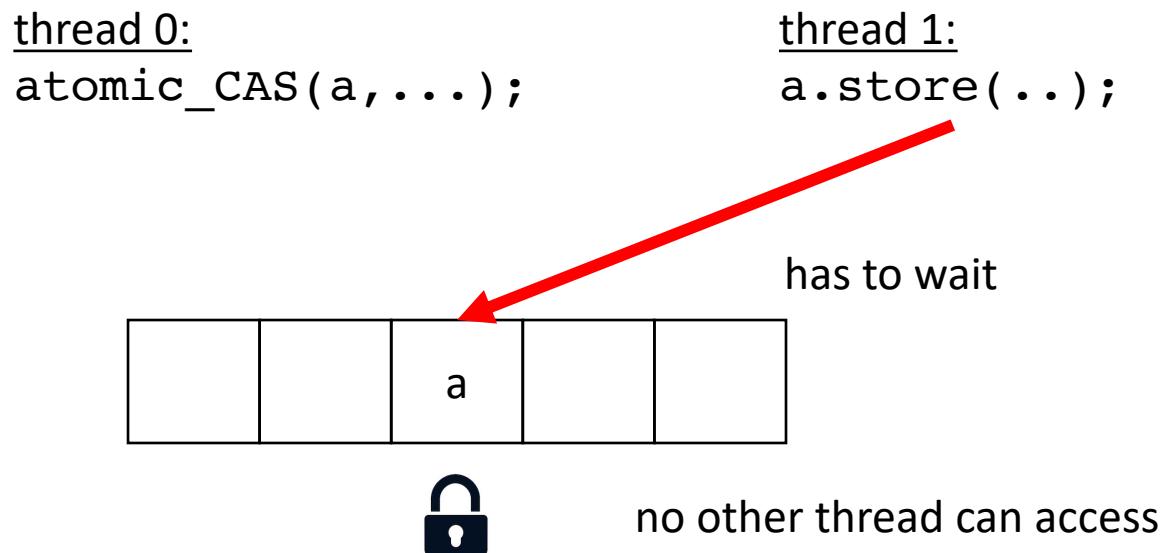
# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start



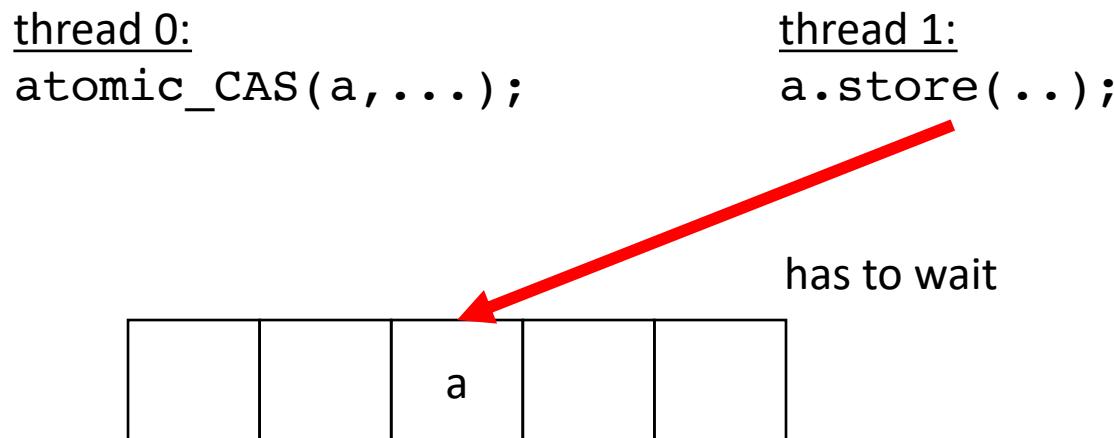
# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start



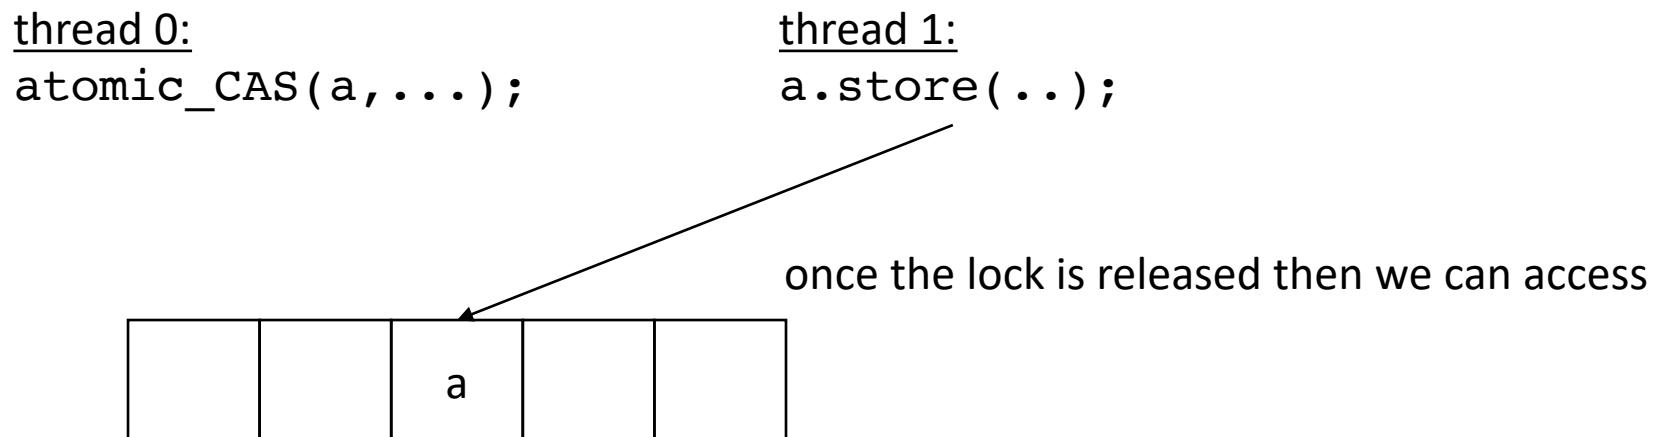
# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start



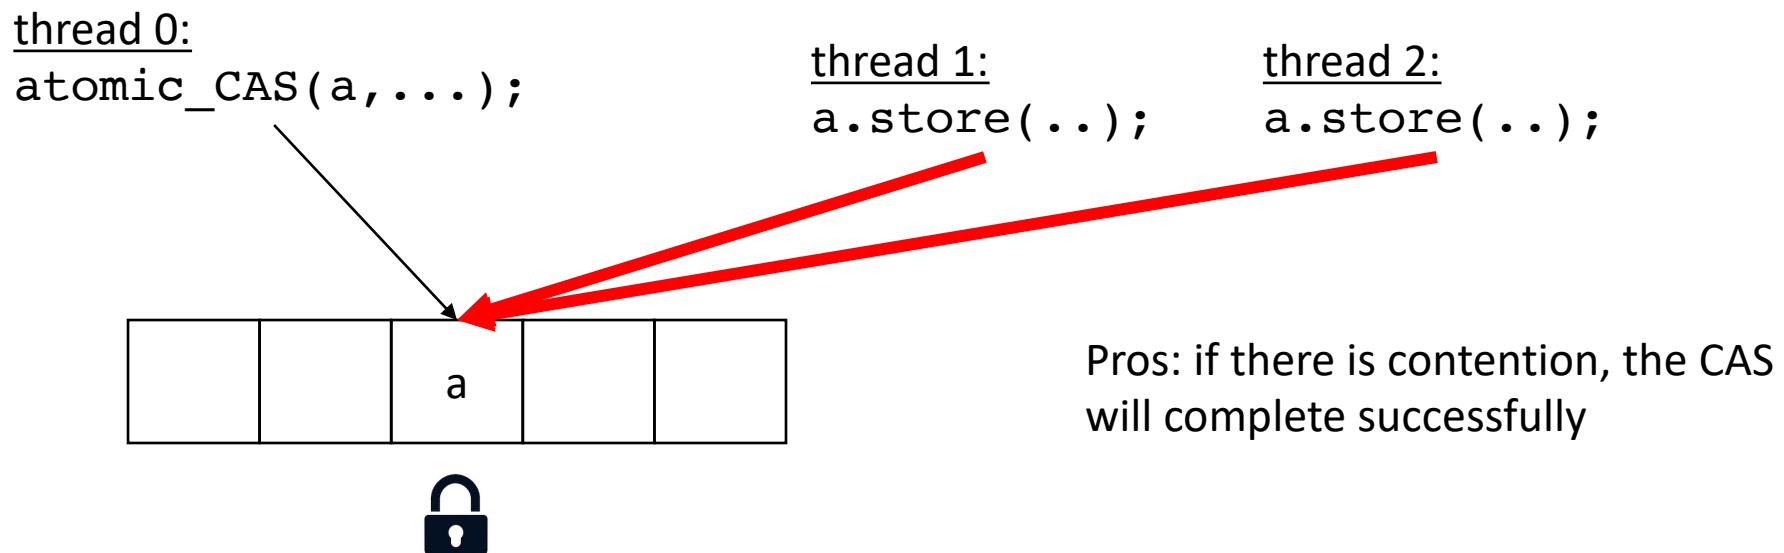
# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start



# Pessimistic Concurrency

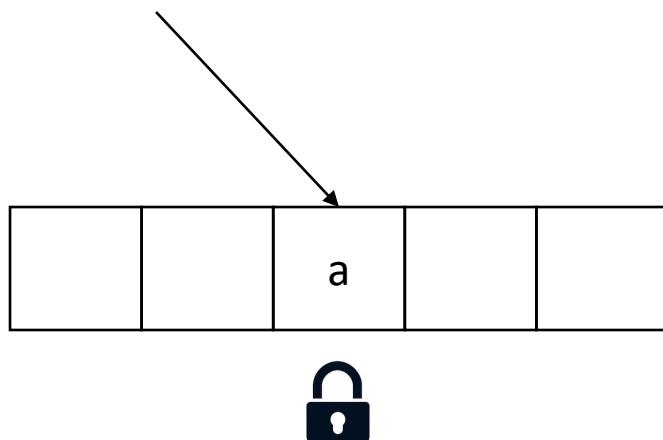
- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start



# Pessimistic Concurrency

- X86 has an actual instruction: lock the memory location
- Known as **Pessimistic Concurrency**
- Assume conflicts will happen and defend against them from the start

thread 0:  
`atomic_CAS(a, ...);`



Cons: if no other threads are contending, lock overhead is high

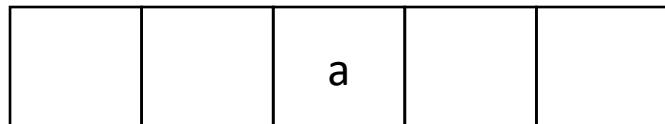
# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume **no** conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```

*For this example consider an atomic increment*

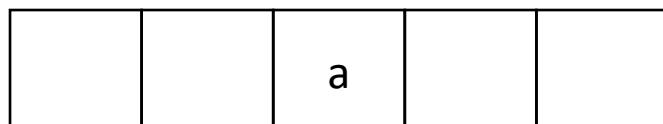


# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```

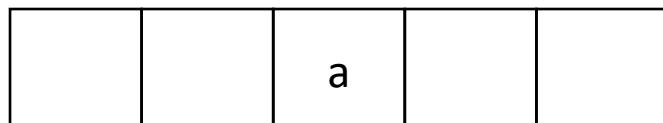


T0\_exclusive = 1

# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:  
tmp = load\_exclusive(a, ...);  
**tmp += 1;**  
store\_exclusive(a, tmp);



T0\_exclusive = 1

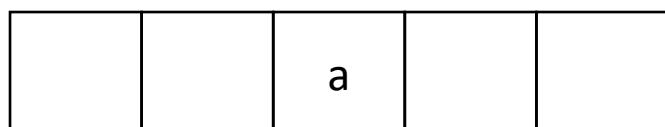
# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```

before we store, we have to check if there was a conflict.



**T0\_exclusive = 1**

# Optimistic Concurrency

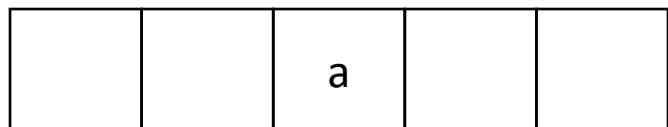
- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```

thread 1:

```
a.store(...)
```



# Optimistic Concurrency

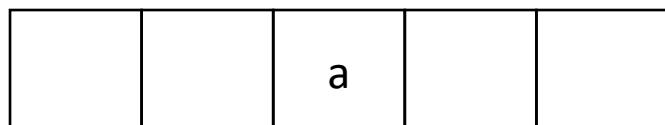
- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```

thread 1:

```
a.store(...)
```



T0\_exclusive = 1

# Optimistic Concurrency

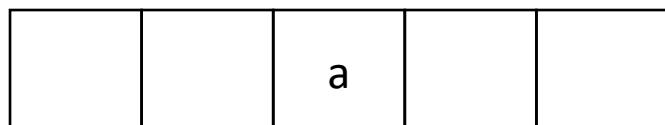
- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```

thread 1:

```
a.store(...)
```



T0\_exclusive = 1

# Optimistic Concurrency

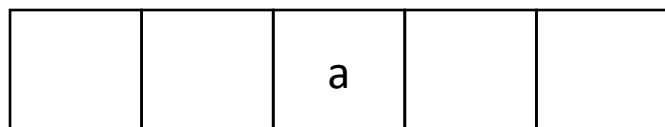
- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```

thread 1:

```
a.store(...)
```



T0\_exclusive = 0

# Optimistic Concurrency

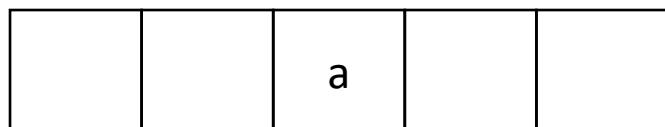
- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a,...);
tmp += 1;
store_exclusive(a, tmp);
```

thread 1:

```
a.store(...)
```



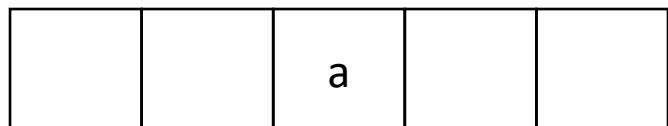
T0\_exclusive = 0

# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume **no** conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a,...);  
tmp += 1;  
store_exclusive(a, tmp);
```



T0\_exclusive = 0

thread 1:

```
a.store(...)
```

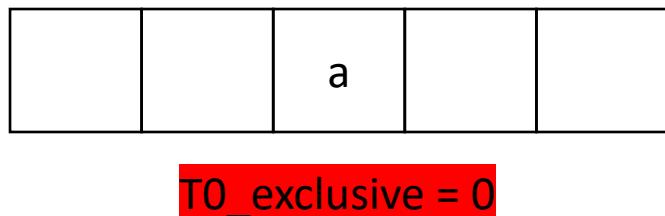
*can't store because our exclusive bit was changed, i.e. there was a conflict!*

# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume **no** conflicts will happen. Detects and reacts to them.

thread 0:

```
tmp = load_exclusive(a, ...);  
tmp += 1;  
store_exclusive(a, tmp);
```



thread 1:

```
a.store(...)
```

*can't store because our exclusive bit was changed, i.e. there was a conflict!*

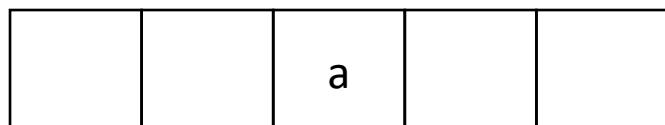
*solution: loop until success:*

# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
do {  
    tmp = load_exclusive(a,...);  
    tmp += 1;  
} while(!store_exclusive(a, tmp));
```



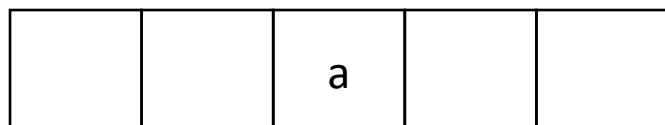
T0\_exclusive = 0

# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
do {  
    tmp = load_exclusive(a,...);  
    tmp += 1;  
} while(!store_exclusive(a, tmp));
```



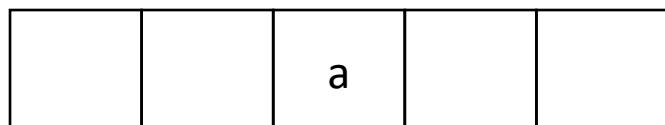
T0\_exclusive = 0

# Optimistic Concurrency

- ARM has load/store exclusive
- Known as **Optimistic Concurrency**
- Assume *no* conflicts will happen. Detects and reacts to them.

thread 0:

```
do {  
    tmp = load_exclusive(a,...);  
    tmp += 1;  
} while(!store_exclusive(a, tmp));
```



T0\_exclusive = 0

Pros: very efficient when there is no conflicts!

Cons: conflicts are very expensive!

Spinning thread might starve (but not indefinitely) if other threads are constantly writing.

# Back to mutexes...

- Speaking of starvation:
- Are the Exchange lock or Spin lock starvation free?

# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*Is this mutex starvation Free?*

```
void unlock() {  
    flag.store(false);  
}
```

mutex  
request

core 0

mutex  
request

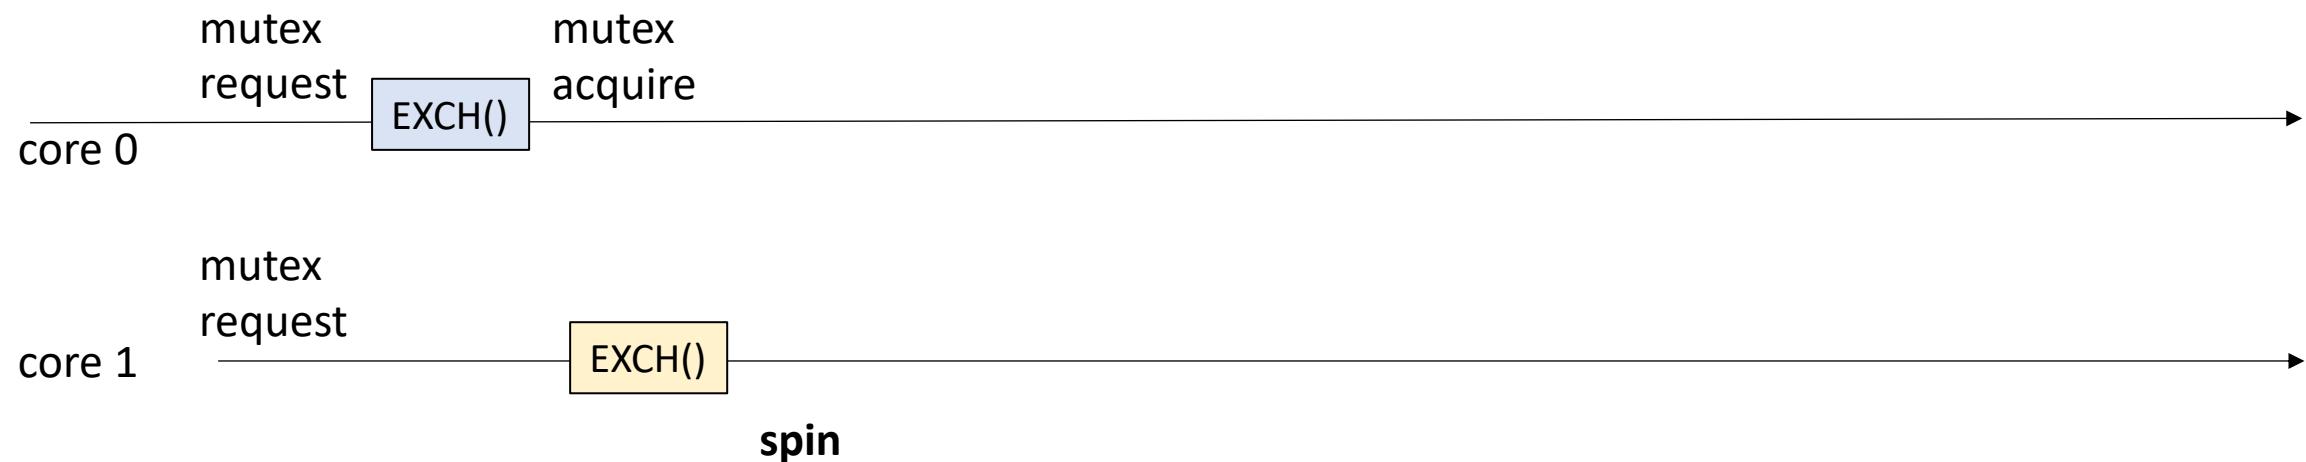
core 1

# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*Is this mutex starvation Free?*

```
void unlock() {  
    flag.store(false);  
}
```

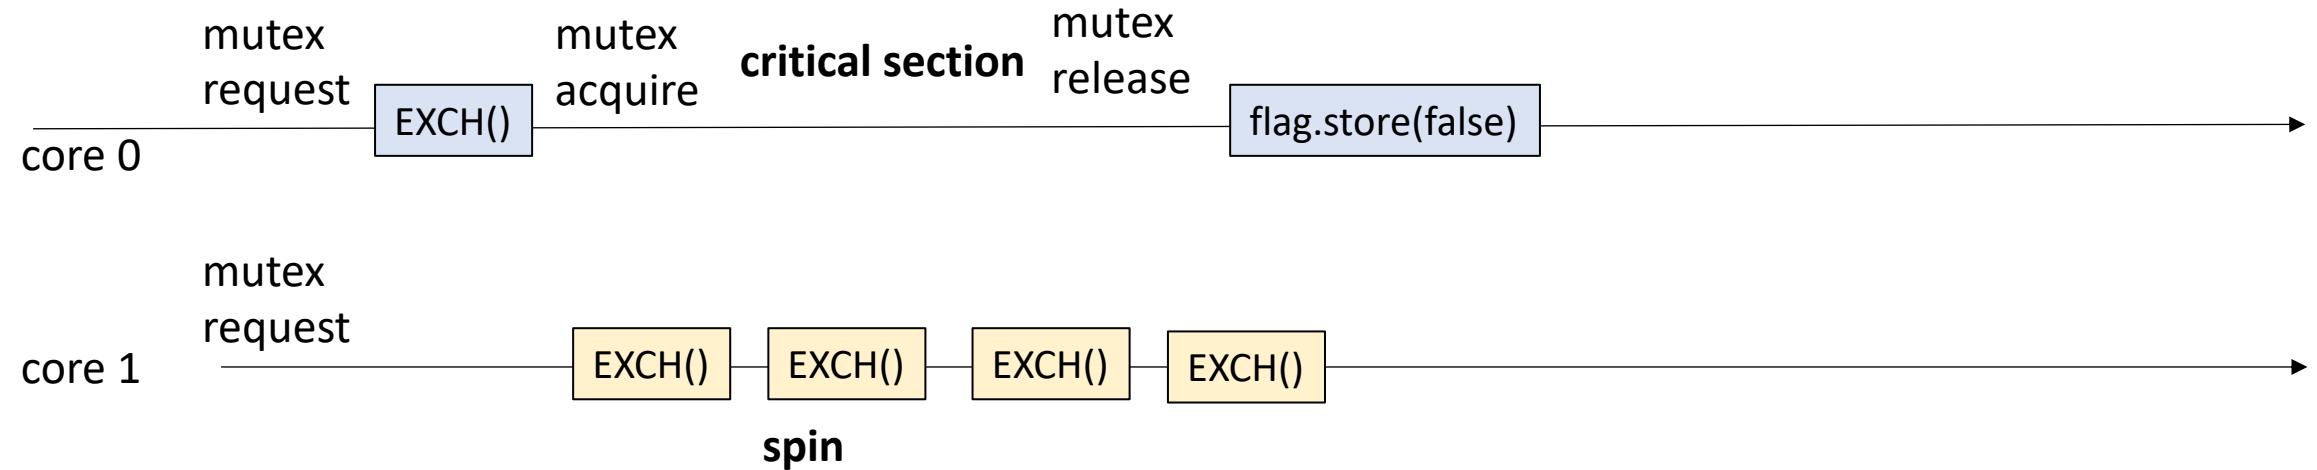


# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*Is this mutex starvation Free?*

```
void unlock() {  
    flag.store(false);  
}
```

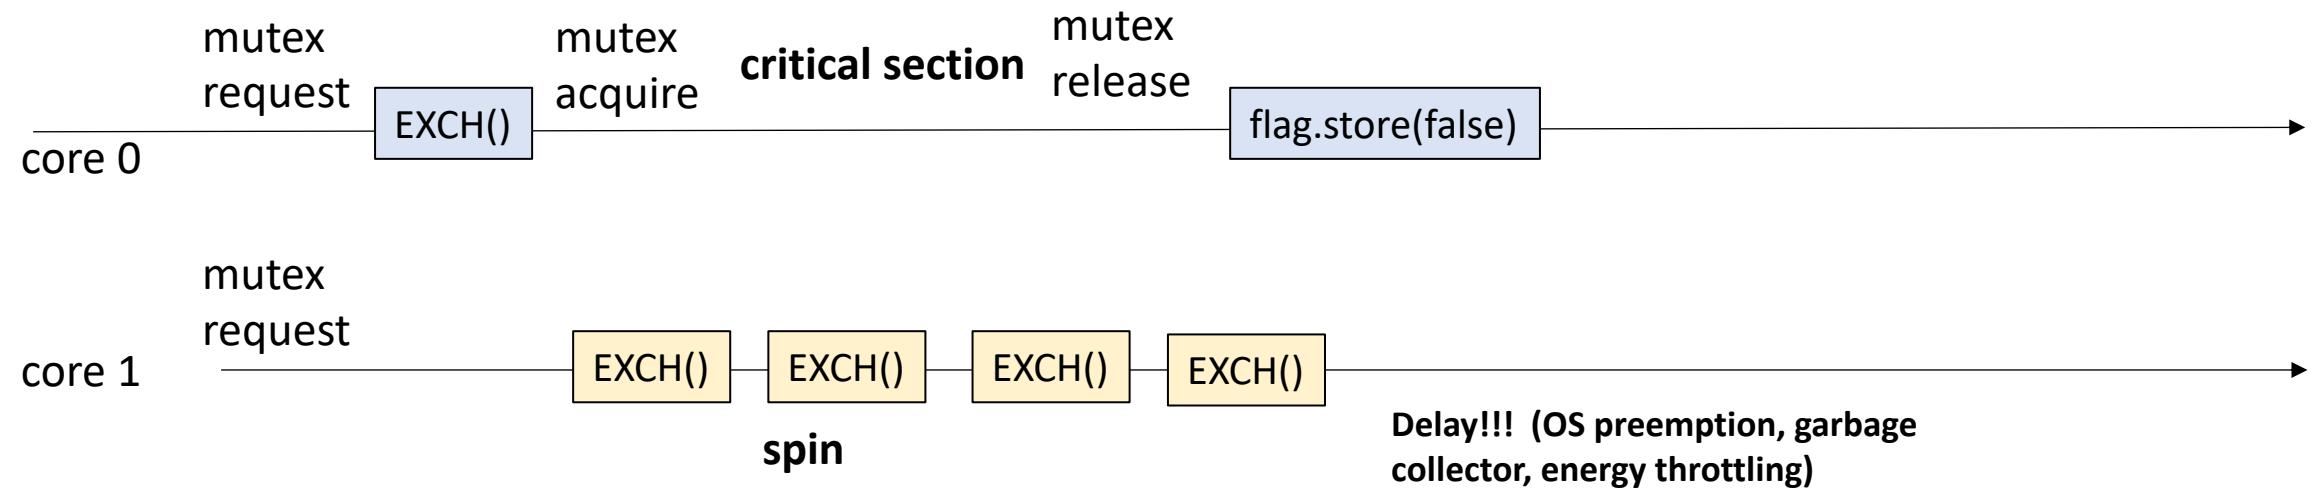


# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

*Is this mutex starvation Free?*

```
void unlock() {  
    flag.store(false);  
}
```

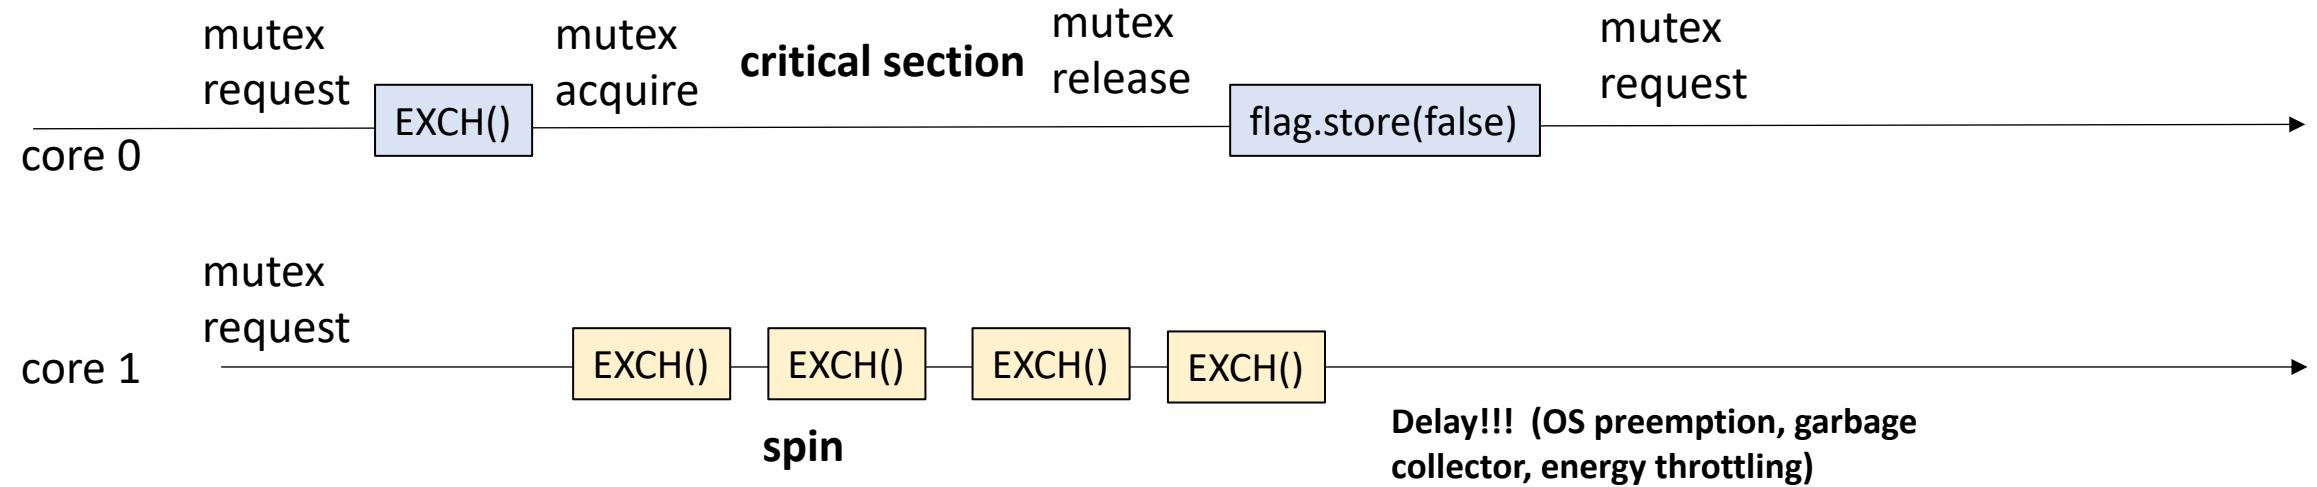


# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Is this mutex starvation Free?

```
void unlock() {  
    flag.store(false);  
}
```

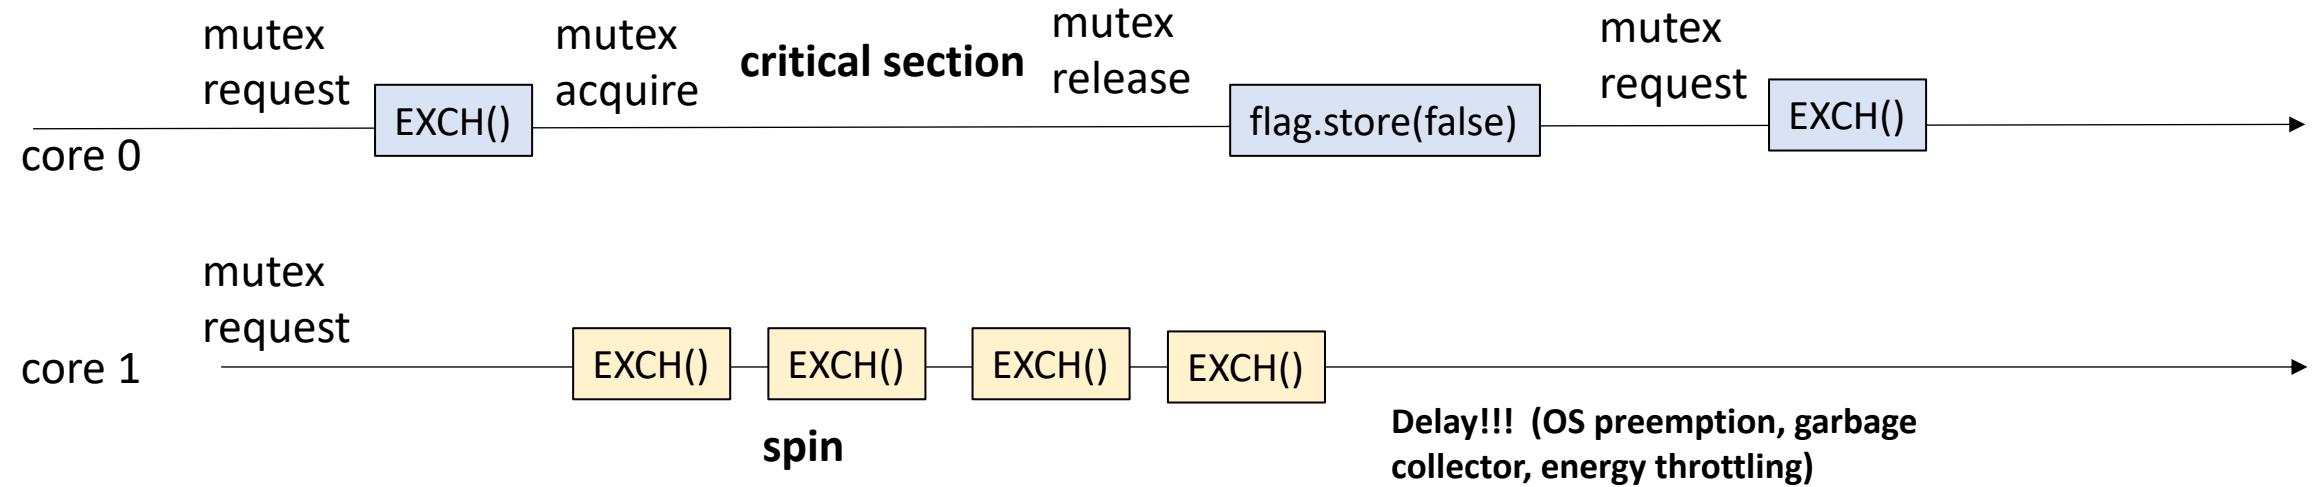


# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Is this mutex starvation Free?

```
void unlock() {  
    flag.store(false);  
}
```

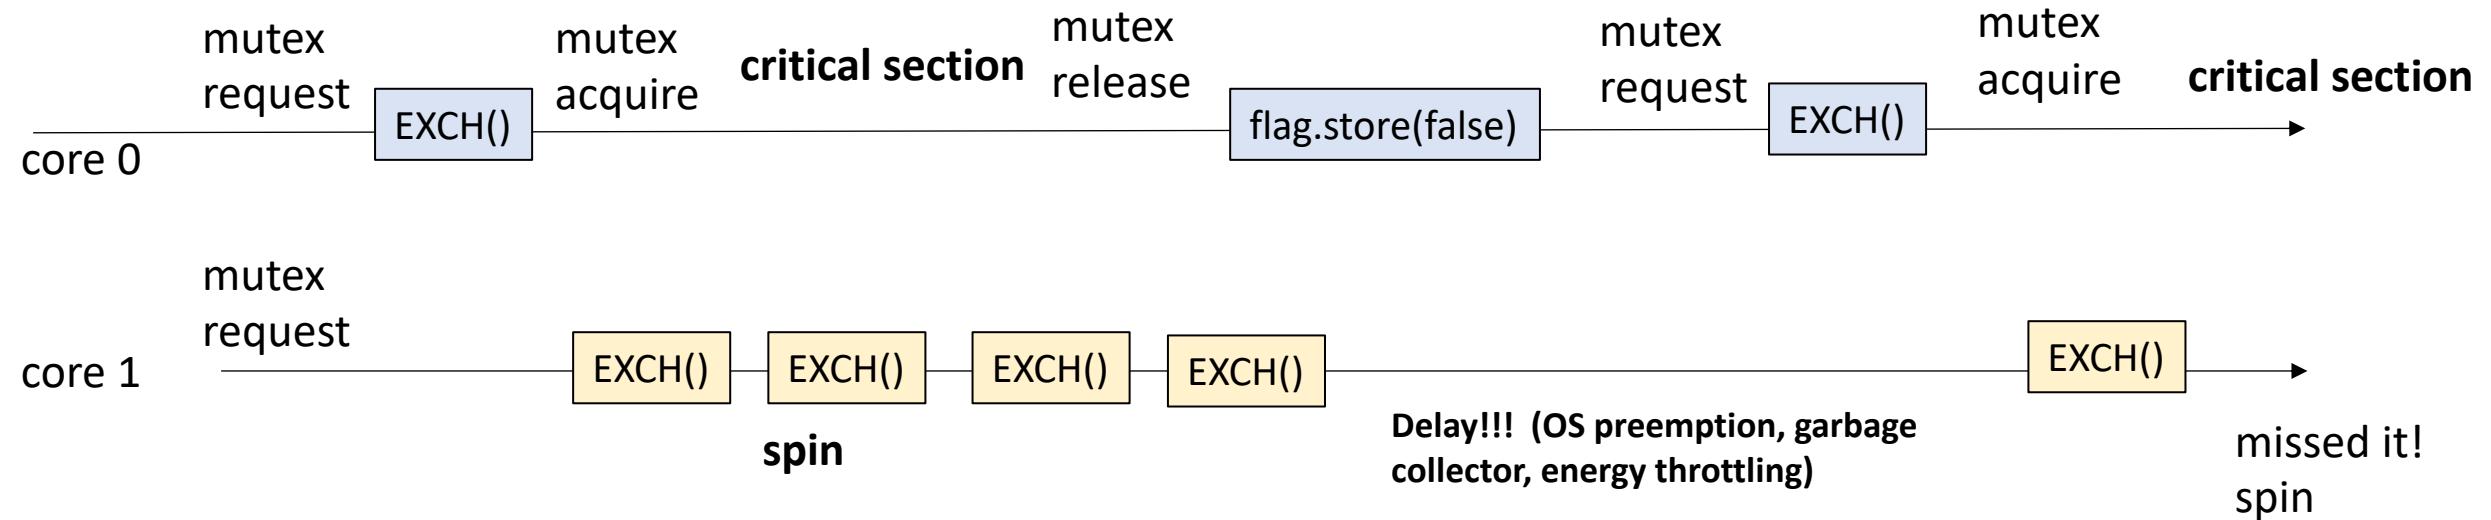


# Analysis

```
void lock() {  
    while (atomic_exchange(&flag, true) == true);  
}
```

Is this mutex starvation Free?

```
void unlock() {  
    flag.store(false);  
}
```



# How can we make this more fair?

- Use a different atomic instruction:
  - `int atomic_fetch_add(atomic_int *a, int v);`

*We've seen this one before!*

# How can we make this more fair?

- Use a different atomic instruction:
  - `int atomic_fetch_add(atomic_int *a, int v);`

*We've seen this one before!*  
*intuition: take a ticket*



like at Zoccoli's!



# Lecture Schedule

- **Atomic RMW mutexes**

- Exchange
- CAS
- Ticket

- Optimizations

- Relaxed peeking
- Backoff

# Ticket lock

```
class Mutex {
public:
    Mutex() {
        counter = 0;
        currently_serving = 0;
    }

    void lock() {
        int my_number = atomic_fetch_add(&counter, 1);
        while (currently_serving.load() != my_number);
    }

    void unlock() {
        int tmp = currently_serving.load();
        tmp += 1;
        currently_serving.store(tmp);
    }

private:
    atomic_int counter;
    atomic_int currently_serving;
};
```

- Ticket lock: instead of 1 bit, we need an integer for the counter.
- The mutex also needs to track of which ticket is currently being served

# Ticket lock

```
class Mutex {
public:
    Mutex() {
        counter = 0;
        currently_serving = 0;
    }

    void lock() {
        int my_number = atomic_fetch_add(&counter, 1);
        while (currently_serving.load() != my_number);
    }

    void unlock() {
        int tmp = currently_serving.load();
        tmp += 1;
        currently_serving.store(tmp);
    }

private:
    atomic_int counter;
    atomic_int currently_serving;
};
```

- Ticket lock: instead of 1 bit, we need an integer for the counter.
- The mutex also needs to track of which ticket is currently being served

*Get a unique number*

*Spin while your number isn't being served*

*To release, increment the number that's currently being served.*

# Analysis

*Is this mutex starvation Free?*

```
void lock() {
    int my_number = atomic_fetch_add(&counter, 1);
    while (currently_serving.load() != my_number);
}

void unlock() {
    int tmp = currently_serving.load();
    tmp += 1;
    currently_serving.store(tmp);
}
```

mutex  
request

core 0

mutex  
request

core 1

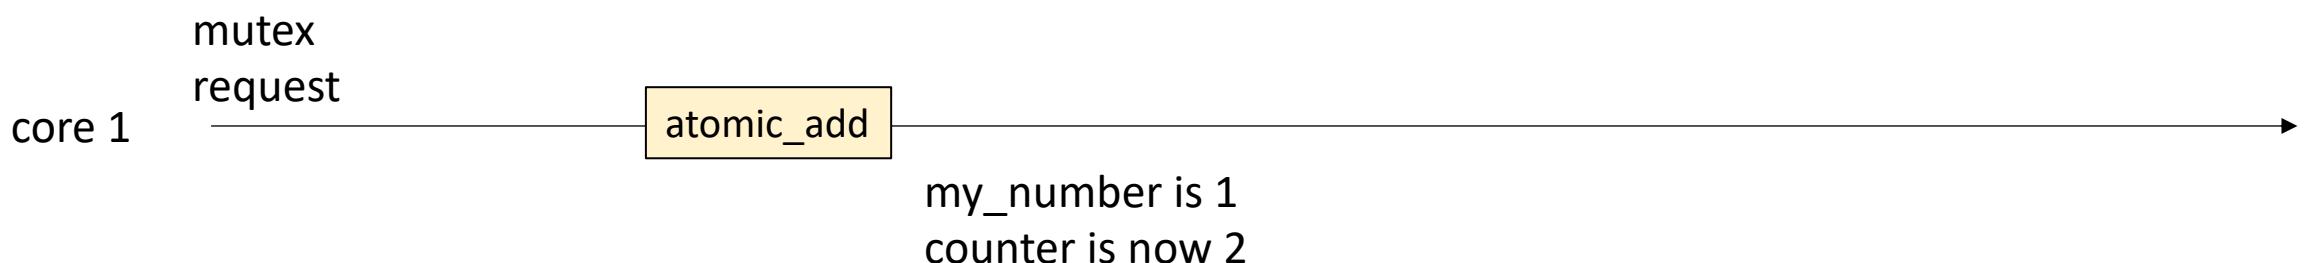
# Analysis

*Is this mutex starvation Free?*

```
void lock() {
    int my_number = atomic_fetch_add(&counter, 1);
    while (currently_serving.load() != my_number);
}

void unlock() {
    int tmp = currently_serving.load();
    tmp += 1;
    currently_serving.store(tmp);
}
```

currently\_serving is 0



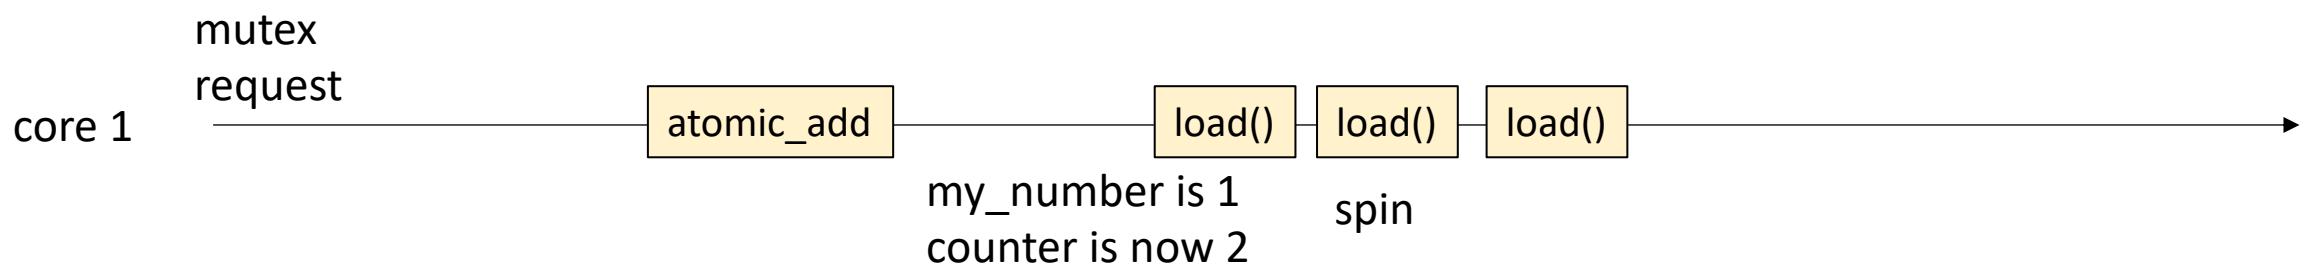
# Analysis

Is this mutex starvation Free?

```
void lock() {
    int my_number = atomic_fetch_add(&counter, 1);
    while (currently_serving.load() != my_number);
}

void unlock() {
    int tmp = currently_serving.load();
    tmp += 1;
    currently_serving.store(tmp);
}
```

currently\_serving is 0

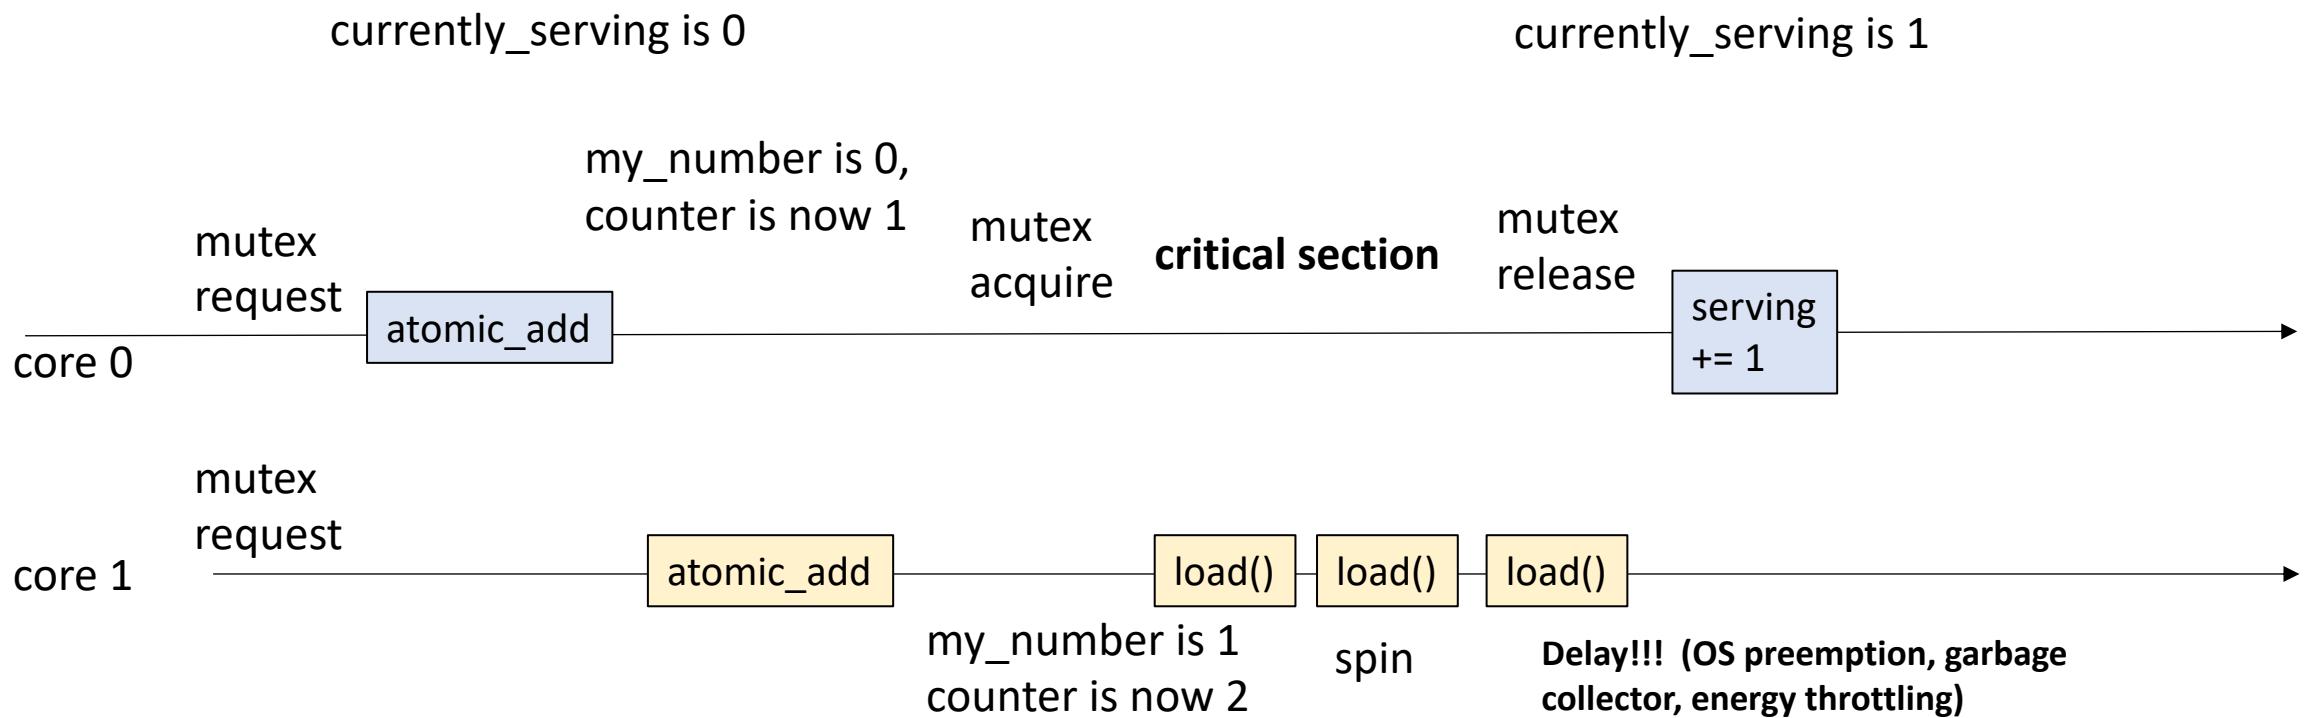


# Analysis

Is this mutex starvation Free?

```
void lock() {
    int my_number = atomic_fetch_add(&counter, 1);
    while (currently_serving.load() != my_number);
}

void unlock() {
    int tmp = currently_serving.load();
    tmp += 1;
    currently_serving.store(tmp);
}
```

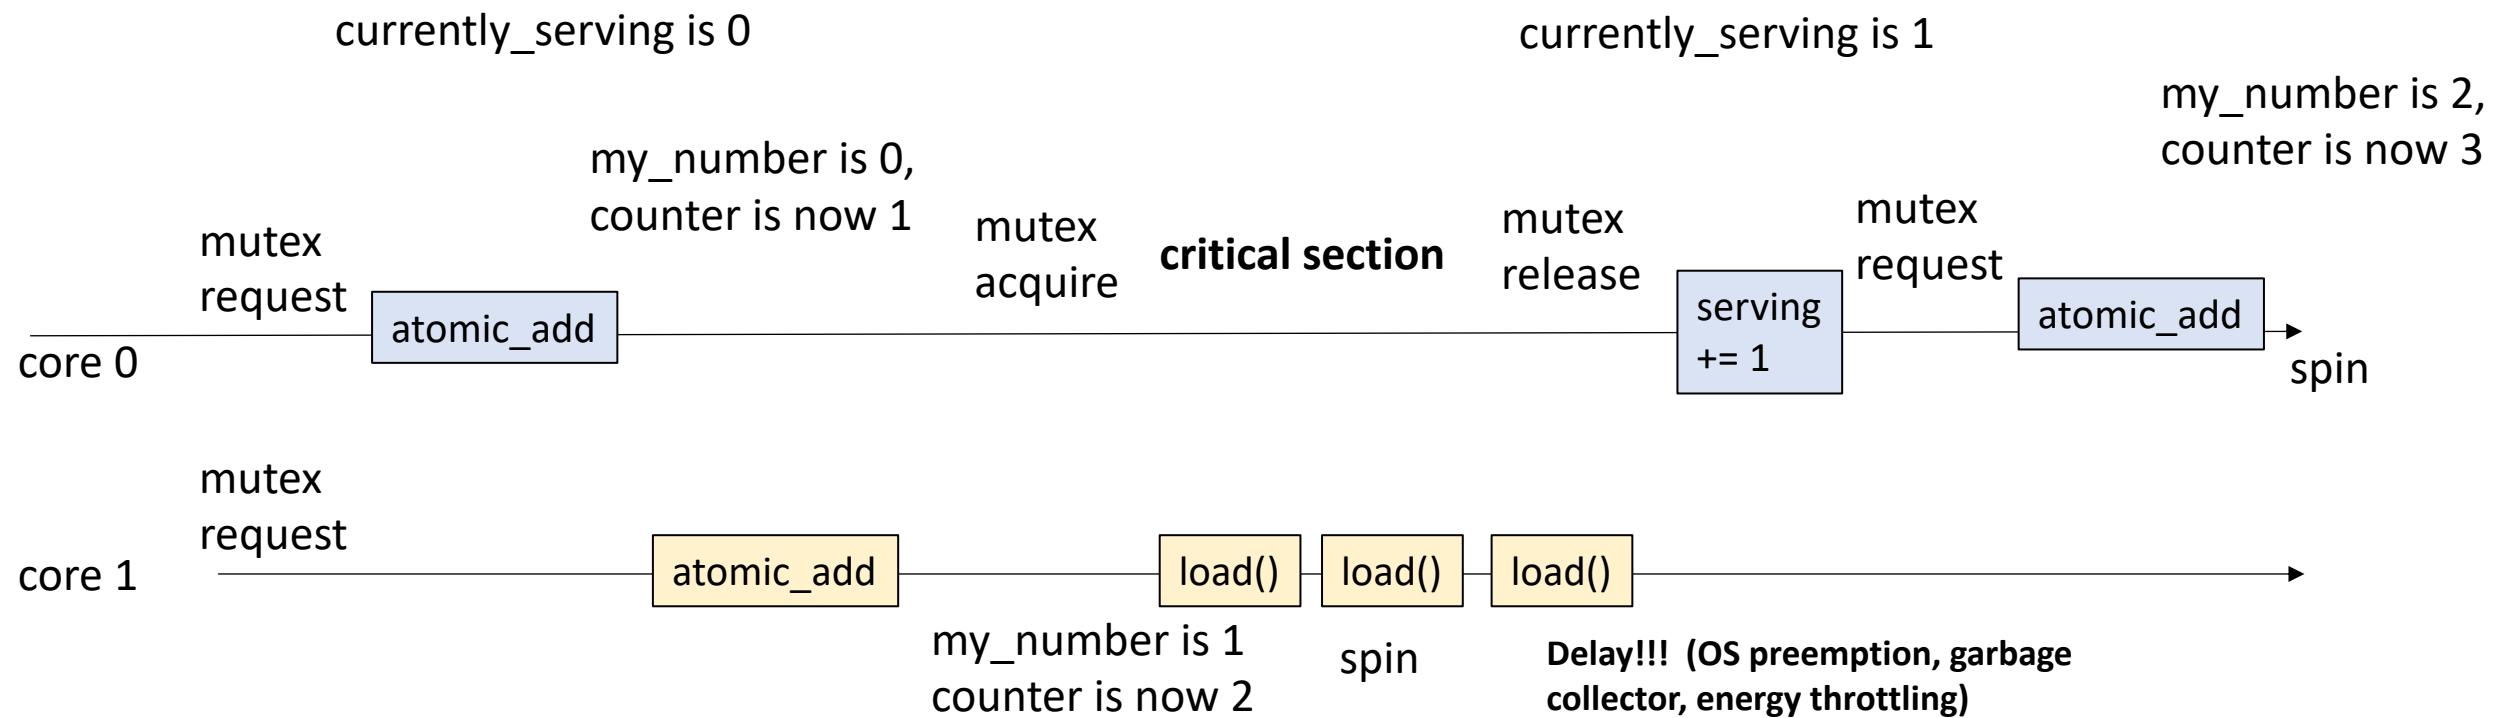


# Analysis

Is this mutex starvation Free?

```
void lock() {
    int my_number = atomic_fetch_add(&counter, 1);
    while (currently_serving.load() != my_number);
}

void unlock() {
    int tmp = currently_serving.load();
    tmp += 1;
    currently_serving.store(tmp);
}
```

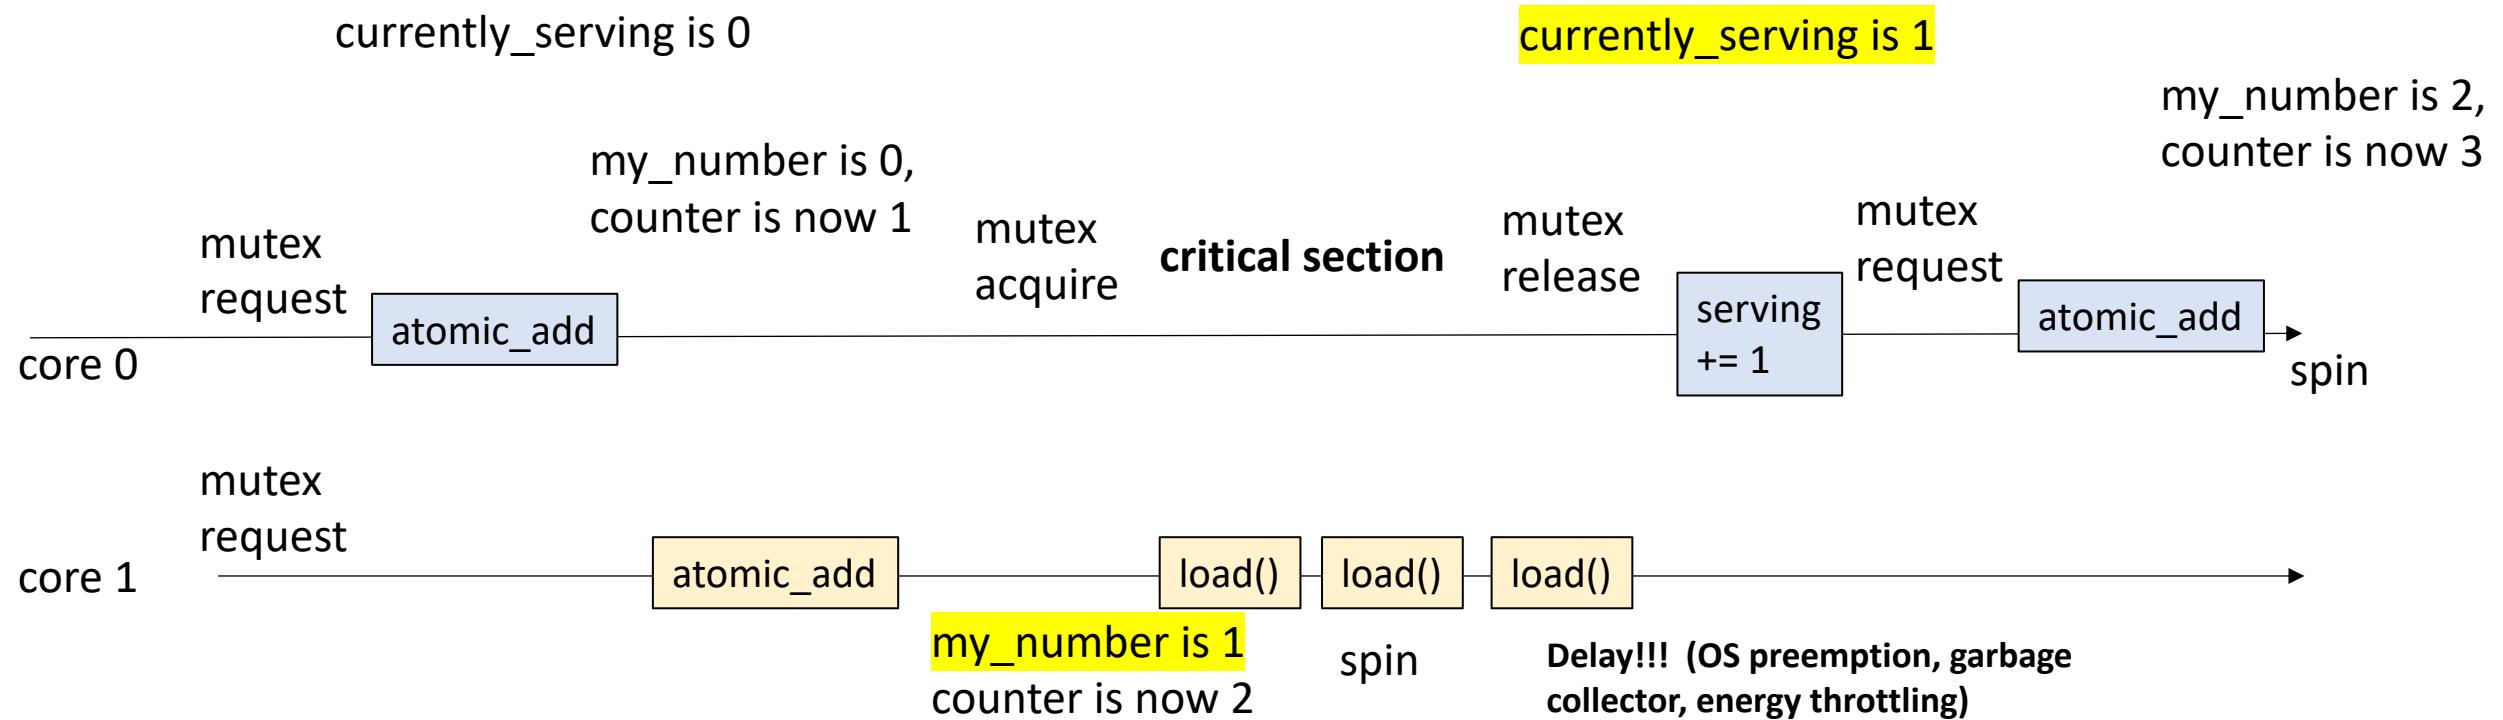


# Analysis

Is this mutex starvation Free?

```
void lock() {
    int my_number = atomic_fetch_add(&counter, 1);
    while (currently_serving.load() != my_number);
}

void unlock() {
    int tmp = currently_serving.load();
    tmp += 1;
    currently_serving.store(tmp);
}
```

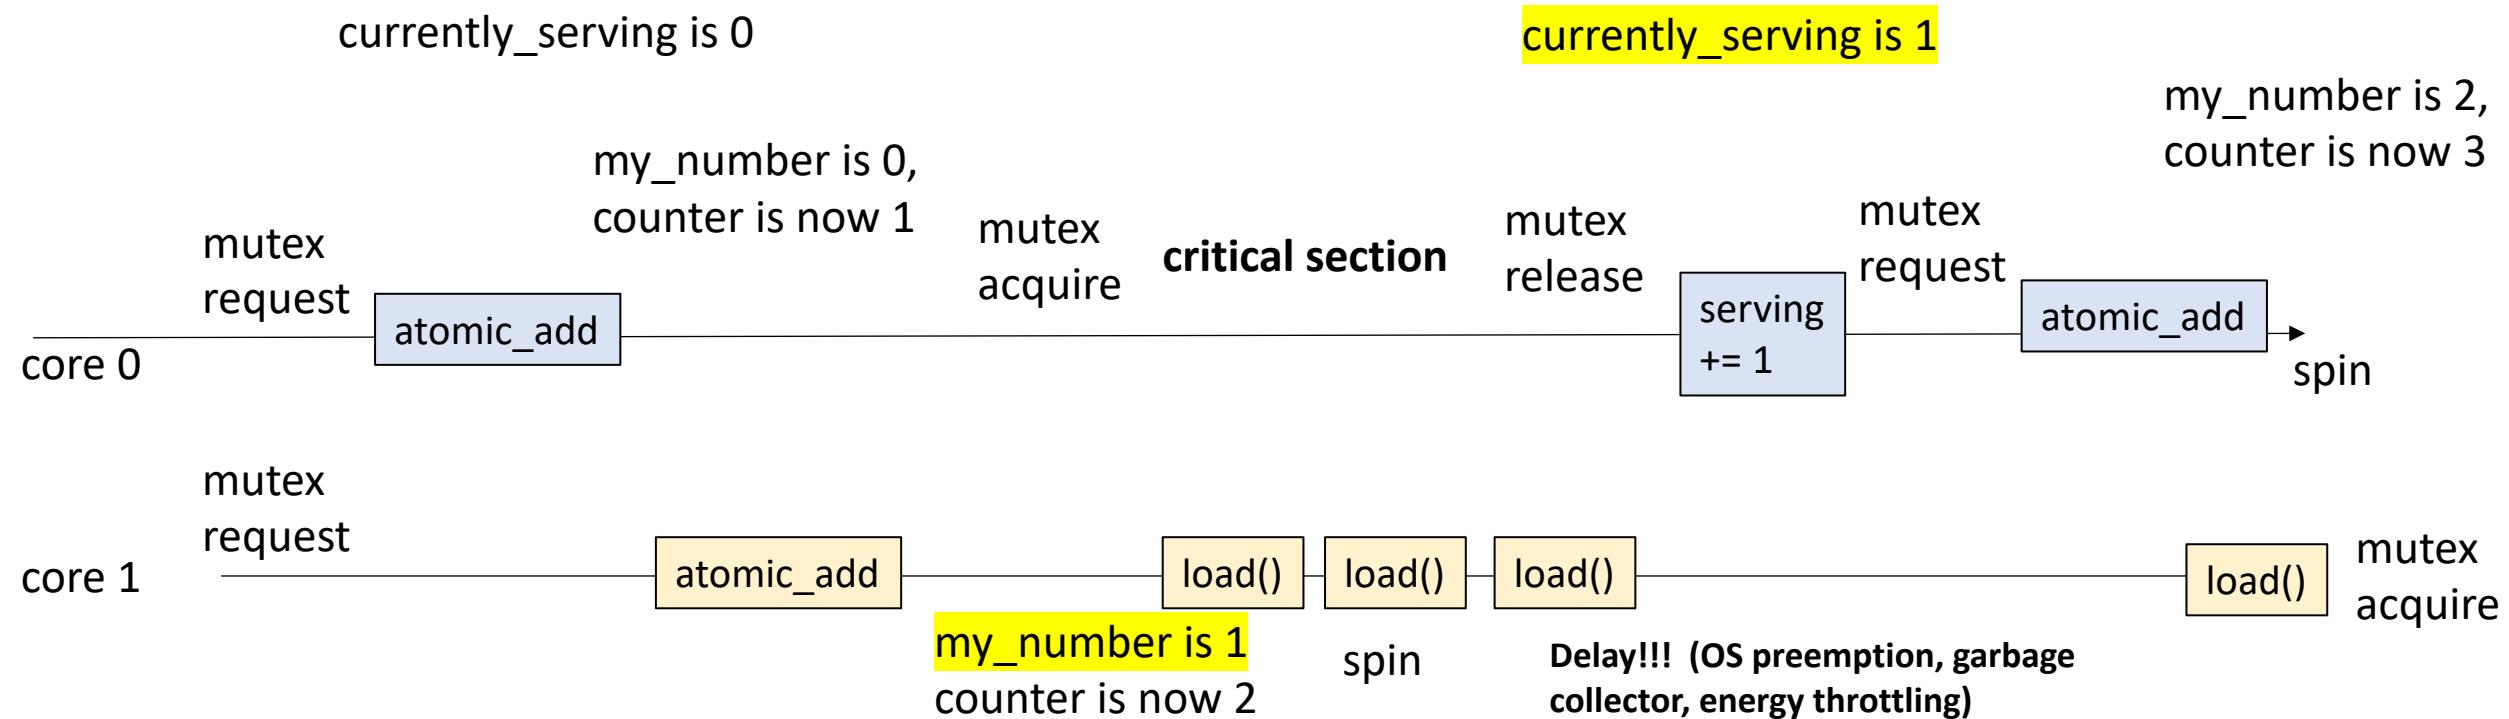


# Analysis

Is this mutex starvation Free?

```
void lock() {
    int my_number = atomic_fetch_add(&counter, 1);
    while (currently_serving.load() != my_number);
}

void unlock() {
    int tmp = currently_serving.load();
    tmp += 1;
    currently_serving.store(tmp);
}
```



# Fair but at what cost?

- Example

# Lecture Schedule

- Atomic RMW mutexes

- Exchange
- CAS
- Ticket

- Optimizations

- Relaxed peeking
- Backoff

# Lecture Schedule

- Atomic RMW mutexes

- Exchange
- CAS
- Ticket

- **Optimizations**

- **Relaxed peeking**
- Backoff

# Optimizations: relaxed peeking

- Relaxed Peeking
  - the Writes in RMWs cost extra; rather than always modify, we can do a simple check first

```
void lock() {
    bool check = false;
    while (atomic_compare_exchange_strong(&flag, &check, true) == false);
}

void unlock() {
    flag.store(false);
}
```

# Optimizations: relaxed peeking

- Relaxed Peeking
  - the Writes in RMWs cost extra; rather than always modify, we can do a simple check first

```
void lock() {
    bool acquired = false;
    bool expected = false;
    while (!acquired) {
        while(flag.load() == false);
        acquired = atomic_compare_exchange_strong(&flag, &expected, true);
    }
}
```

# Optimizations: relaxed peeking

- What about the load in the loop? Remember the memory fence? Do we need to flush our caches every time we peek?
- We only need to flush when we actually acquire the mutex

```
void lock() {
    bool acquired = false;
    bool expected = false;
    while (!acquired) {
        while(flag.load() == false);
        acquired = atomic_compare_exchange_strong(&flag, &expected, true);
    }
}
```

# Optimizations: relaxed peeking

- What about the load in the loop? Remember the memory fence? Do we need to flush our caches every time we peek?
- We only need to flush when we actually acquire the mutex

```
void lock() {
    bool acquired = false;
    bool expected = false;
    while (!acquired) {
        while(flag.load_explicit(memory_order_relaxed) == false);
        acquired = atomic_compare_exchange_strong(&flag, &expected, true);
    }
}
```

```
void lock() {  
    bool acquired = false;  
    bool expected = false;  
    while (!acquired) {  
        while(flag.load_explicit(memory_order_relaxed) == false);  
        acquired = atomic_compare_exchange_strong(&flag, &expected, true);  
    }  
}
```

C0 memory operations are performed and flushed

core 0



core 1



C1 memory operations have **not** yet been performed and cache is invalidated

# Relaxed atomics

- Enter expert mode!
  - explicit atomics with relaxed semantics
  - Beware! they do not provide a memory fence!
  - Only use when a memory fence is issued later before leaving your mutex implementation. Good for “peeking” before you actually execute your RMW.

# Lecture Schedule

- Atomic RMW mutexes

- Exchange
- CAS
- Ticket

- **Optimizations**

- **Relaxed peeking**
- **Backoff**

# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - In non-parallel systems, concurrent threads can get in the way of progress

# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

*Say threads 0 and 1 are executing concurrently*

core 0



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

*Say threads 0 and 1 are executing concurrently*

core 0



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

*Say threads 0 and 1 are executing concurrently*

Thread 0 in critical  
section!

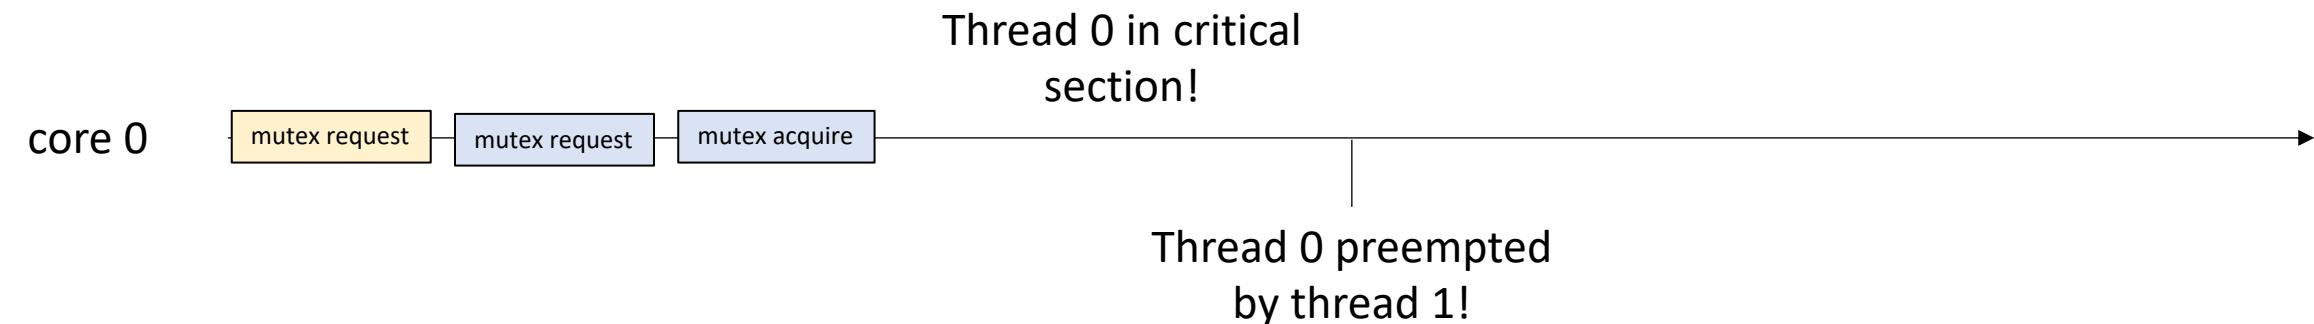
core 0



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

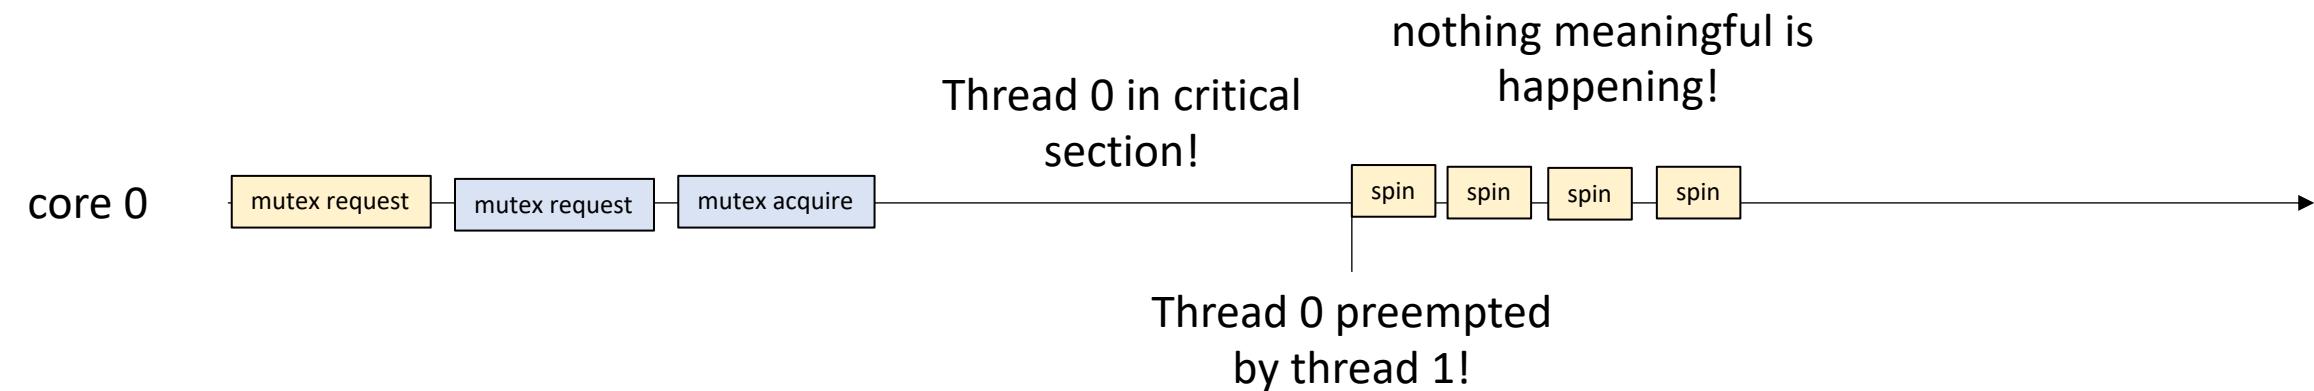
*Say threads 0 and 1 are executing concurrently*



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

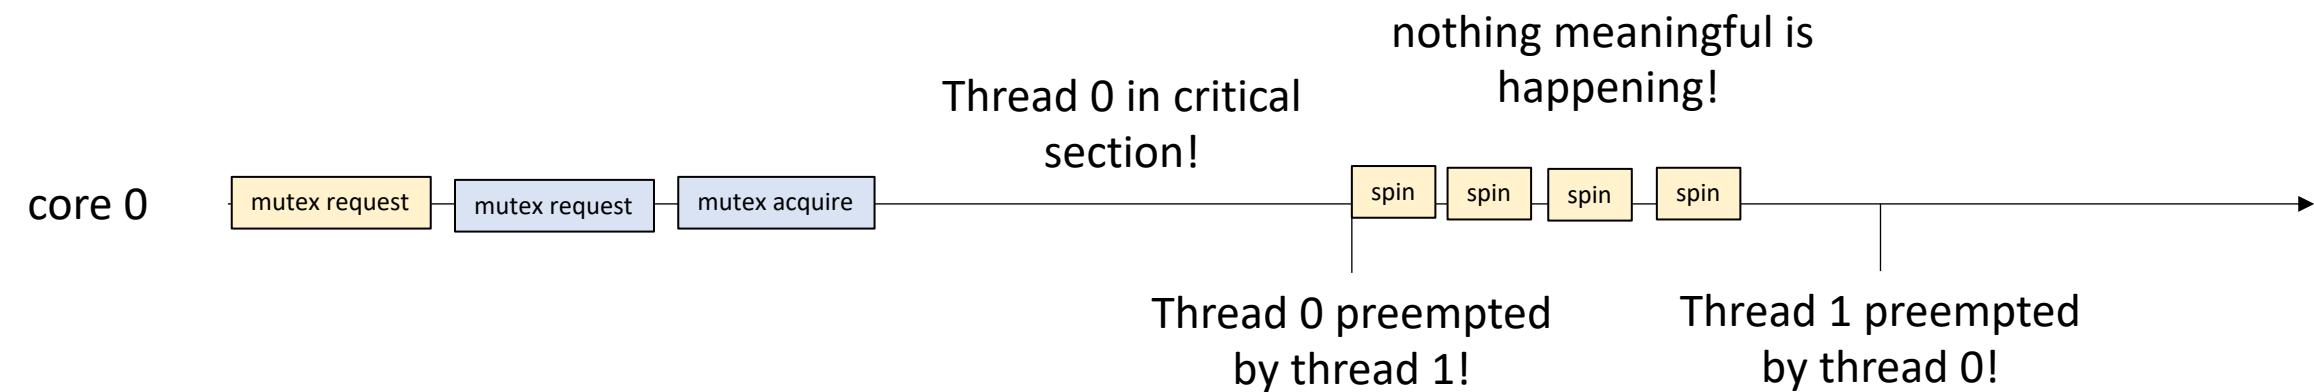
*Say threads 0 and 1 are executing concurrently*



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if it's not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

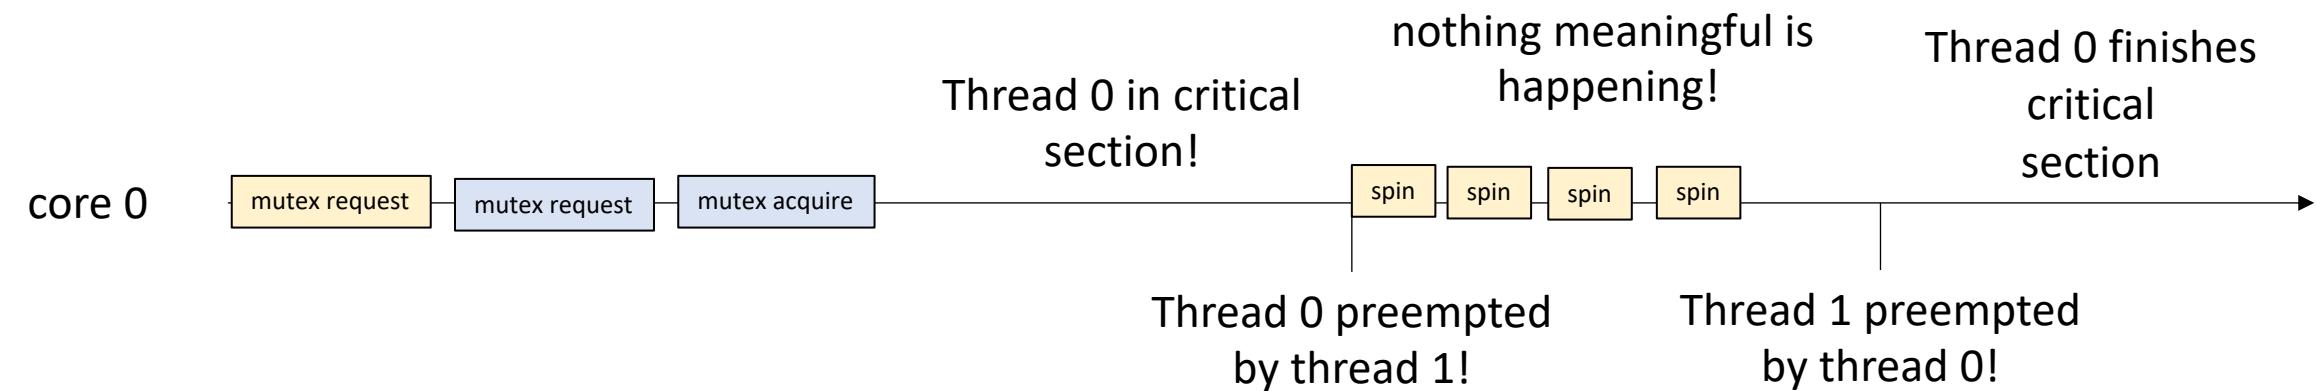
*Say threads 0 and 1 are executing concurrently*



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

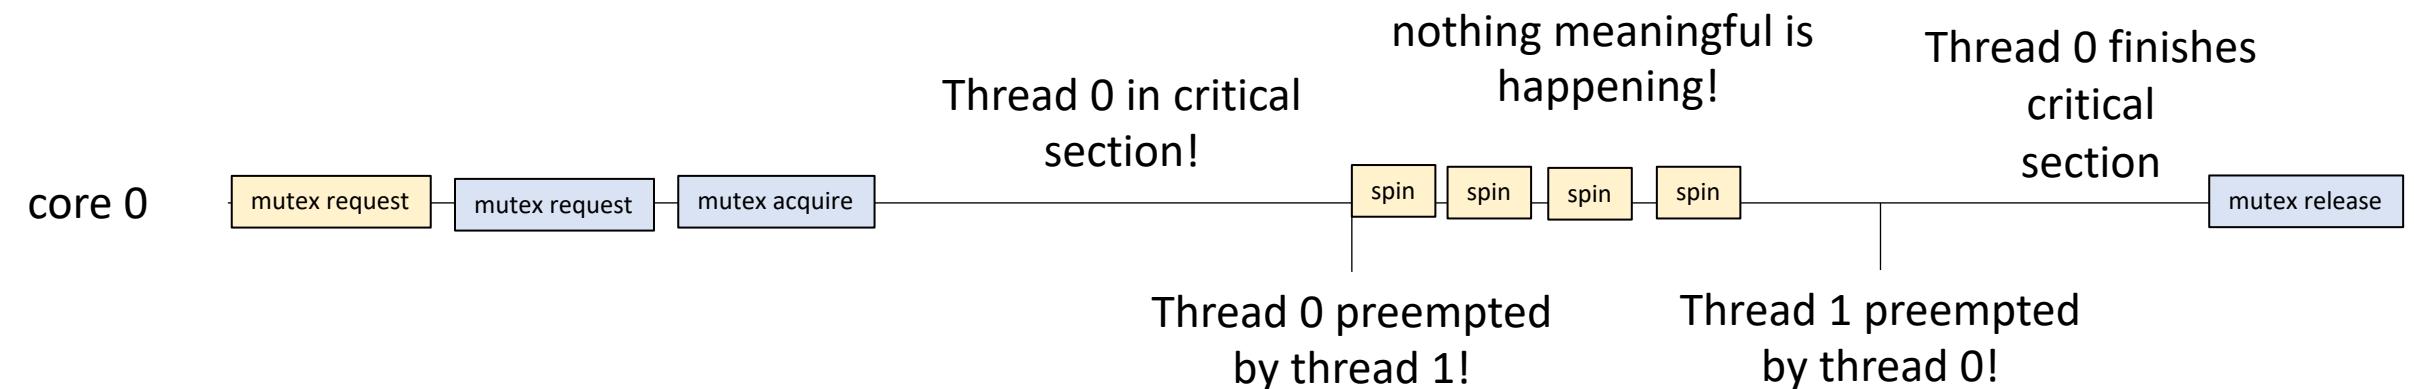
*Say threads 0 and 1 are executing concurrently*



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

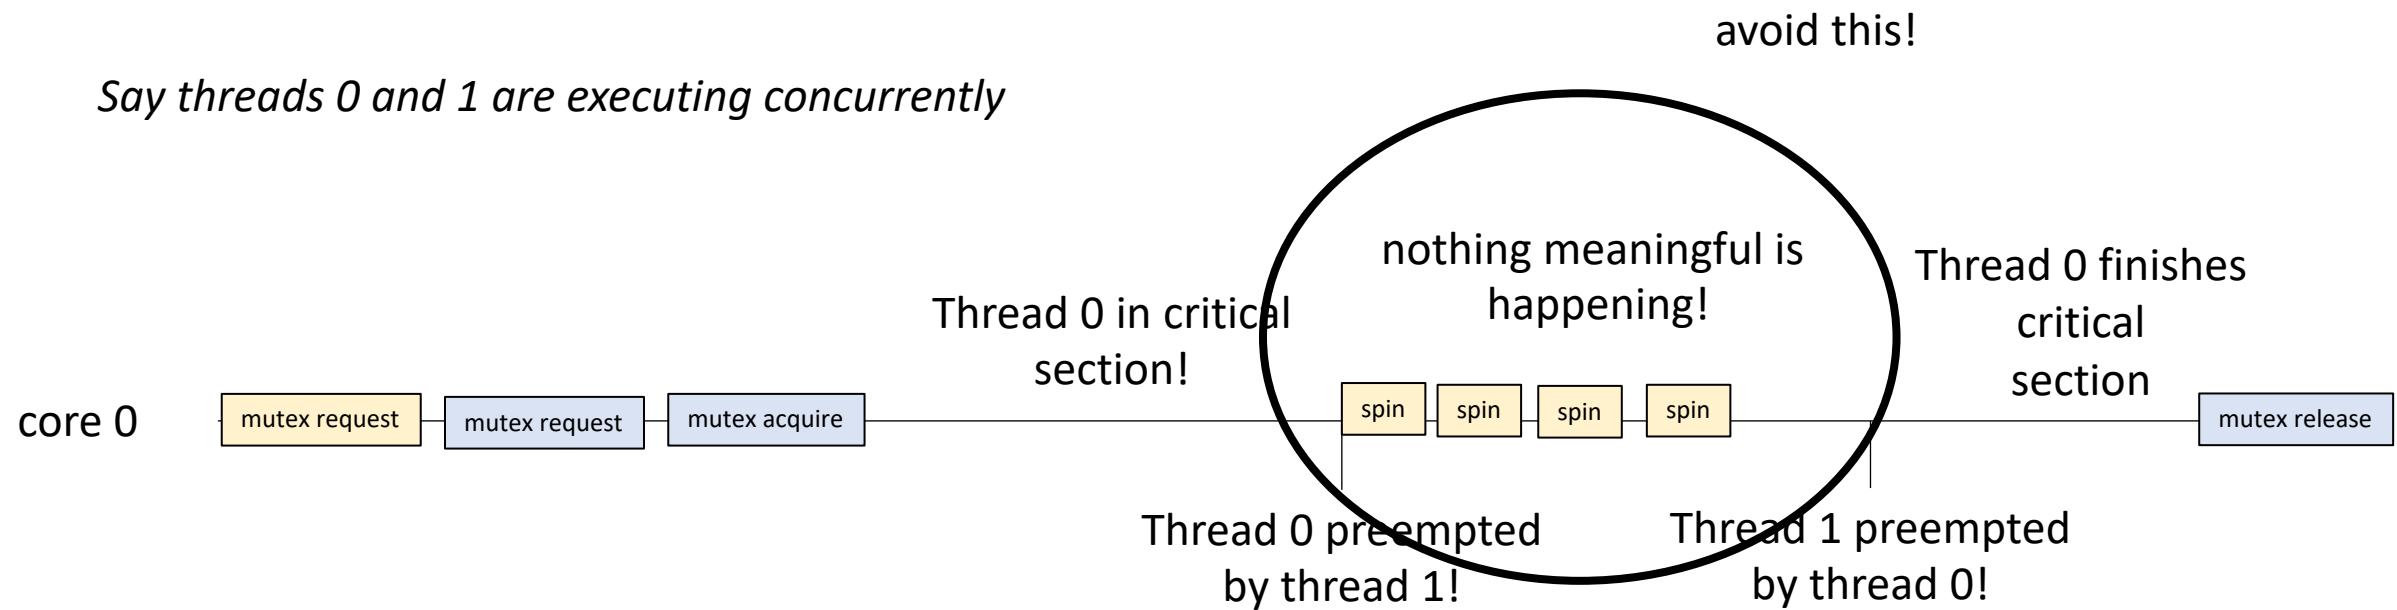
*Say threads 0 and 1 are executing concurrently*



# Optimizations: backoff

- Two issues remain:
  - Loads still cause bus traffic (even if its not as bad as RMWs)
  - **In non-parallel systems, concurrent threads can get in the way of progress**

Say threads 0 and 1 are executing concurrently



# Optimizations: backoff

- C++
  - `this_thread::yield();`
- Hints to the operating system that we should take a break while other threads (potentially the threads that have the mutex) get scheduled.

# Optimizations: backoff

where do we put it?

- C++
  - `this_thread::yield();`
- Hints to the operating system that we should take a break while other threads (potentially the threads that have the mutex) get scheduled.

```
void lock() {
    bool acquired = false;
    bool expected = false;
    while (!acquired) {
        while(flag.load_explicit(memory_order_relaxed) == false);
        acquired = atomic_compare_exchange_strong(&flag, &expected, true);
    }
}
```

# Optimizations: backoff

```
void lock() {
    bool acquired = false;
    bool expected = false;
    while (!acquired) {
        while(flag.load_explicit(memory_order_relaxed) == false) {
            this_thread::yield();
        }
        acquired = atomic_compare_exchange_strong(&flag, &expected, true);
    }
}
```

# Optimizations: backoff

- Other backoff strategies:
  - sleep: more control over how long the thread is suspended.
- Some mutexes try to learn how long to sleep for:
  - Keep track of a sleep time.
  - Every time you spin, increase the sleep time (remember for next spin)
  - If you acquire, reduce the sleep time

# Optimizations: when to use them

- Spinning is useful for short waits on non-oversubscribed systems
- Sleeping/yielding is useful for oversubscribed systems, repeated and regular tasks (e.g. updating UIs)

# End of spin-lock optimizations

- When to use what optimization?
  - Start with C++ mutex, then
  - microbenchmark
  - profile
- Can we optimize around existing mutexes? e.g. C++ mutex?
  - Yes!

# try\_lock

- another common mutex API call: try\_lock()
- one-shot mutex attempt (implementation defined)
- You can then implement your own sleep/yield strategy around this

```
void lock() {
    while (atomic_exchange(&flag, true) == true);
}

void try_lock() {
    return atomic_exchange(&flag, true));
}
```

# Next week

- More locks
  - Queue locks
  - Reentrant locks
  - Reader-writer locks
- Tools to help you with checking your code for data conflicts
- Good luck with homework!